

# About me

## My journey

- Joined the Seoul Bitcoin Meetup in 2017  
([meetup.com/seoulbitcoin](https://www.meetup.com/seoulbitcoin))

# About me

## My journey

- Joined the Seoul Bitcoin Meetup in 2017  
([meetup.com/seoulbitcoin](https://www.meetup.com/seoulbitcoin))
- Started contributing to OSS in 2019

# About me

## My journey

- Joined the Seoul Bitcoin Meetup in 2017  
([meetup.com/seoulbitcoin](https://www.meetup.com/seoulbitcoin))
- Started contributing to OSS in 2019
- Started getting funded for my work in 2020

*"The most important book about technology today,  
with implications that go far beyond programming."*  
—Guy Kawasaki

Revised & Expanded

# THE CATHEDRAL & THE BAZAAR

MUSINGS ON LINUX AND OPEN SOURCE  
BY AN ACCIDENTAL REVOLUTIONARY



ERIC S. RAYMOND

WITH A FOREWORD BY BOB YOUNG, CHAIRMAN & CEO OF RED HAT, INC.

- There isn't an official roadmap

- There isn't an official roadmap
- Everyone looks at scalability differently

- There isn't an official roadmap
- Everyone looks at scalability differently
- These are my thoughts

- Anthony Towns - [www.erisian.com.au/wordpress/  
2023/06/21/putting-the-b-in-btc](http://www.erisian.com.au/wordpress/2023/06/21/putting-the-b-in-btc)
- James O'Berine - [jameso.be/2023/08/14/scaling-  
bitcoin.html](http://jameso.be/2023/08/14/scaling-bitcoin.html)

# **On on-chain scalability**

**Current bottlenecks and my proposed  
solutions**

<https://github.com/kcalvinalvin/2024-05-01-bobspaces>



**“We very, very much need such a system, but  
the way I understand your proposal, it does  
not seem to scale to the required size.”**

<https://www.mail-archive.com/cryptography@metzdowd.com/msg09963.html>

# **What makes Bitcoin special**

**Something that  
cannot be negotiated**

# Trustlessness

**What is it?**

Only trusting the  
software I run

**“Of Bitcoin’s many properties, trustlessness, or the ability to use Bitcoin without trusting anything but the open-source software you run, is, by far, king”**

**<https://bluematt.bitcoin.ninja/2017/02/28/bitcoin-trustlessness/>**

**“Of Bitcoin’s many properties, **trustlessness**, or the ability to use Bitcoin without trusting anything but the open-source software you run, is, by far, king”**

<https://bluematt.bitcoin.ninja/2017/02/28/bitcoin-trustlessness/>

- Decentralization
- P2P
- Self custody

We've rejected  
anything that hurts  
trustlessness

**How do you use Bitcoin  
trustlessly?**

# Running a full node

## Achieving Trustlessness

- Full node = fully validating



# Running a full node

## Achieving Trustlessness

- Full node = fully validating
- Self-validating blocks/txs



# **The Validation Process**

**What all full nodes go through**

- Download ~600GB worth of blocks

# The Validation Process

What all full nodes go through

- Download ~600GB worth of blocks
- Validate that the transaction is spending Bitcoins that exist

# The Validation Process

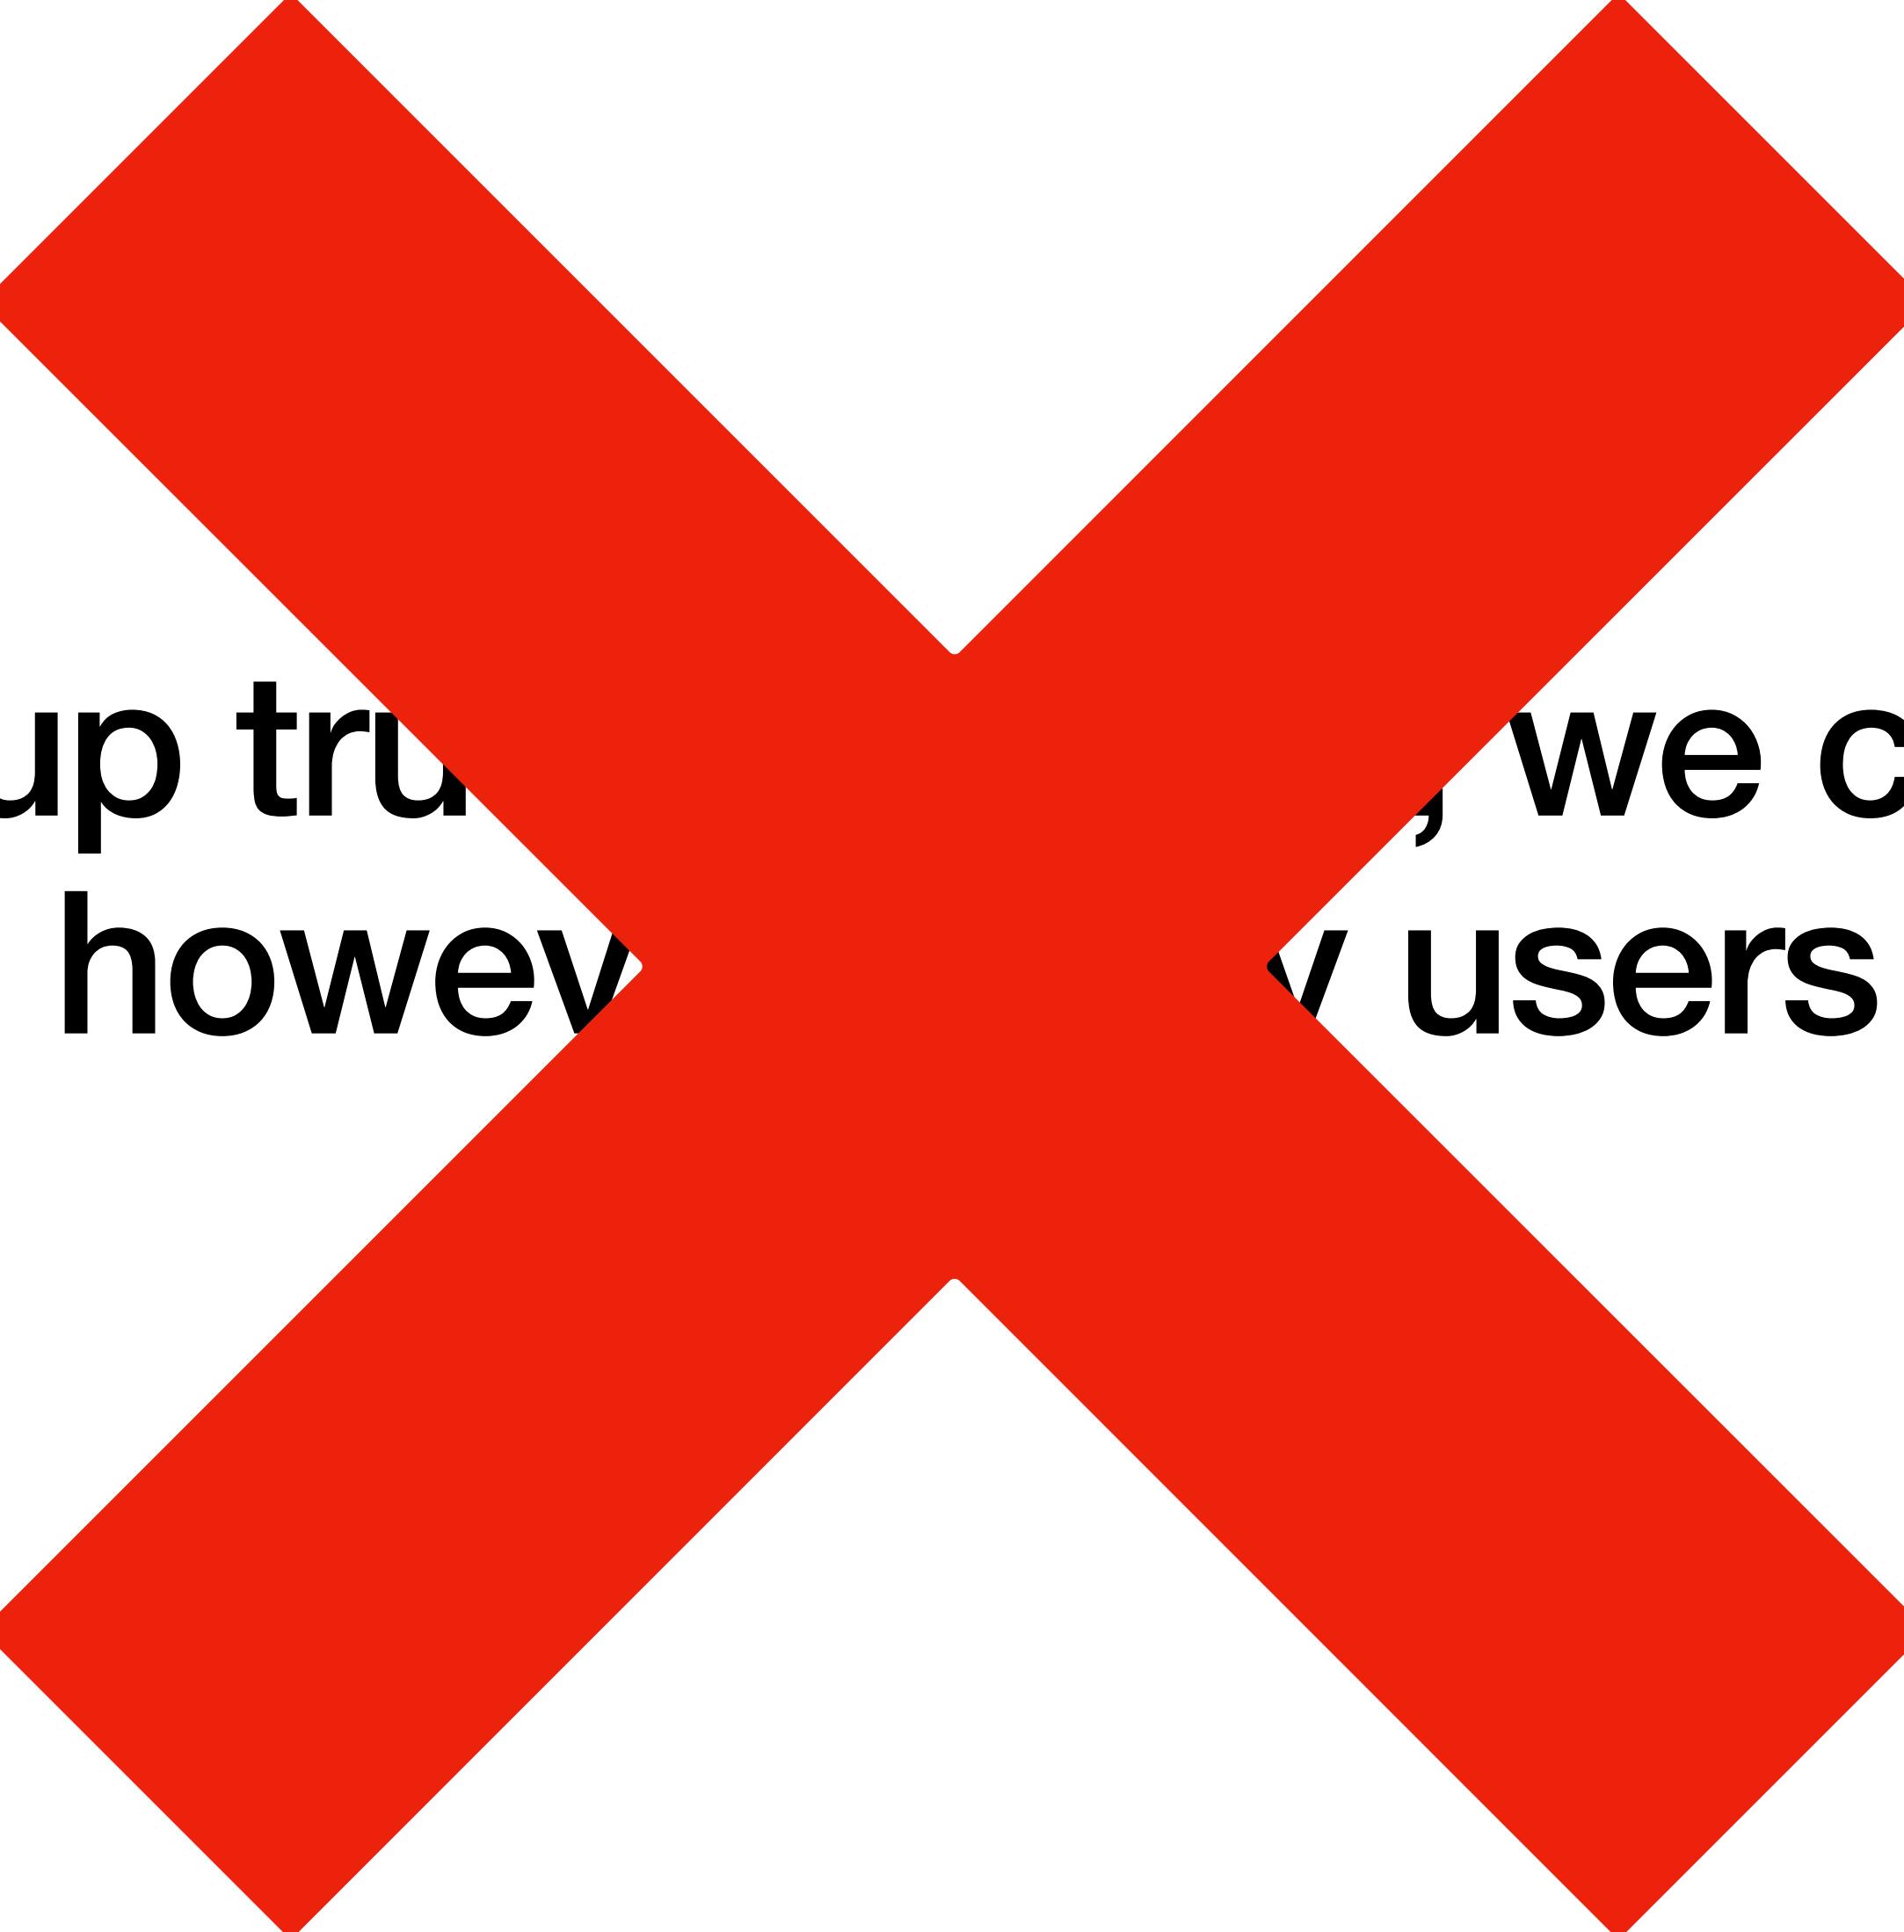
What all full nodes go through

- Download ~600GB worth of blocks
- Validate that the transaction is spending Bitcoins that exist
- Validate that the transaction is able to spend that Bitcoin

# Things limiting scalability

# **Trustlessness!**

**If we give up trustlessness, we can scale  
Bitcoin to however many users we'd like.**



If we give up true  
Bitcoin to however  
we can scale  
to however many users we'd like.

# **Giving up on trustlessness**

**Letting everyone use Bitcoin but very few being able to validate**

- Increase the block size to 1 gigabyte

# **Giving up on trustlessness**

**Letting everyone use Bitcoin but very few being able to validate**

- Increase the block size to 1 gigabyte
- Everyone uses SPV and blindly trusts miners

# **Giving up on trustlessness**

**Letting everyone use Bitcoin but very few being able to validate**

- Increase the block size to 1 gigabyte
- Everyone uses SPV and blindly trusts miners
- Miners do whatever they want

**If we give up on the amount of users,  
everyone can use Bitcoin trustlessly**



**If we give up on  
everyone can use it  
of users,  
n trustlessly**

# Giving up on user count

Letting a few amount of people use Bitcoin trustlessly

- Limit the block size to 1KB

# Giving up on user count

Letting a few amount of people use Bitcoin trustlessly

- Limit the block size to 1KB
- The entire chain is only 1GB up to block 1 million. Validation takes less than a minute on a normal phone.

# **Giving up on user count**

**Letting a few amount of people use Bitcoin trustlessly**

- Limit the block size to 1KB
- The entire chain is only 1GB up to block 1 million. Validation takes less than a minute on a normal phone
- Fees are crazy. I can't afford them

**There's a  
balance**

**SMALLER BLOCK BETTER**



**BIGGER BLOCK BETTER**



# What I strive for

The most amount of  
users while maintaining  
trustlessness

# **Methods for getting more bang for the buck**

## **While maintaining trustlessness**

- Use offchain protocols

# **Off chain protocols**

**What they're really all about**

**Achieve multiple transfers with  
a single on-chain transaction**

# Liquid Network

## Offchain protocols

- Trading Bitcoin for IOUs on liquid sidechain

# Liquid Network

## Offchain protocols

- Trading Bitcoin for IOUs on liquid sidechain
- Trading off trustlessness for scalability

# Liquid Network

## Offchain protocols

- Trading Bitcoin for IOUs on liquid sidechain
- Trading off trustlessness for scalability
- Opt-in

# Liquid Network

## Offchain protocols

- Trading Bitcoin for IOUs on liquid sidechain
- Trading off trustlessness for scalability
- Opt-in
- `liquid.net`

# **Statechains**

## **Offchain protocols**

- Moving a whole UTXO to a statechain ran by an entity

# **Statechains**

## **Offchain protocols**

- Moving a whole UTXO to a statechain ran by an entity
- Need to put trust in the Statechain entity

# Statechains

## Offchain protocols

- Moving a whole UTXO to a statechain ran by an entity
- Need to put trust in the Statechain entity
- [medium.com/@RubenSomsen/statechains-non-custodial-off-chain-bitcoin-transfer-1ae4845a4a39](https://medium.com/@RubenSomsen/statechains-non-custodial-off-chain-bitcoin-transfer-1ae4845a4a39)

# **Lightning Network**

## **Offchain protocols**

- Create channels between each other

# Lightning Network

## Offchain protocols

- Create channels between each other
- Transfer as long as the channel is balanced and is open

# **Lightning Network**

## **Offchain protocols**

- Create channels between each other
- Transfer as long as the channel is balanced and is open
- Security tradeoff. Passive -> active

**Ultimately these protocols require  
an on-chain tx somewhere**

# **Increase the block size**

**Needed to support more users.**

**Currently can't due to it hurting trustlessness**

# Why can't we increase the block size?

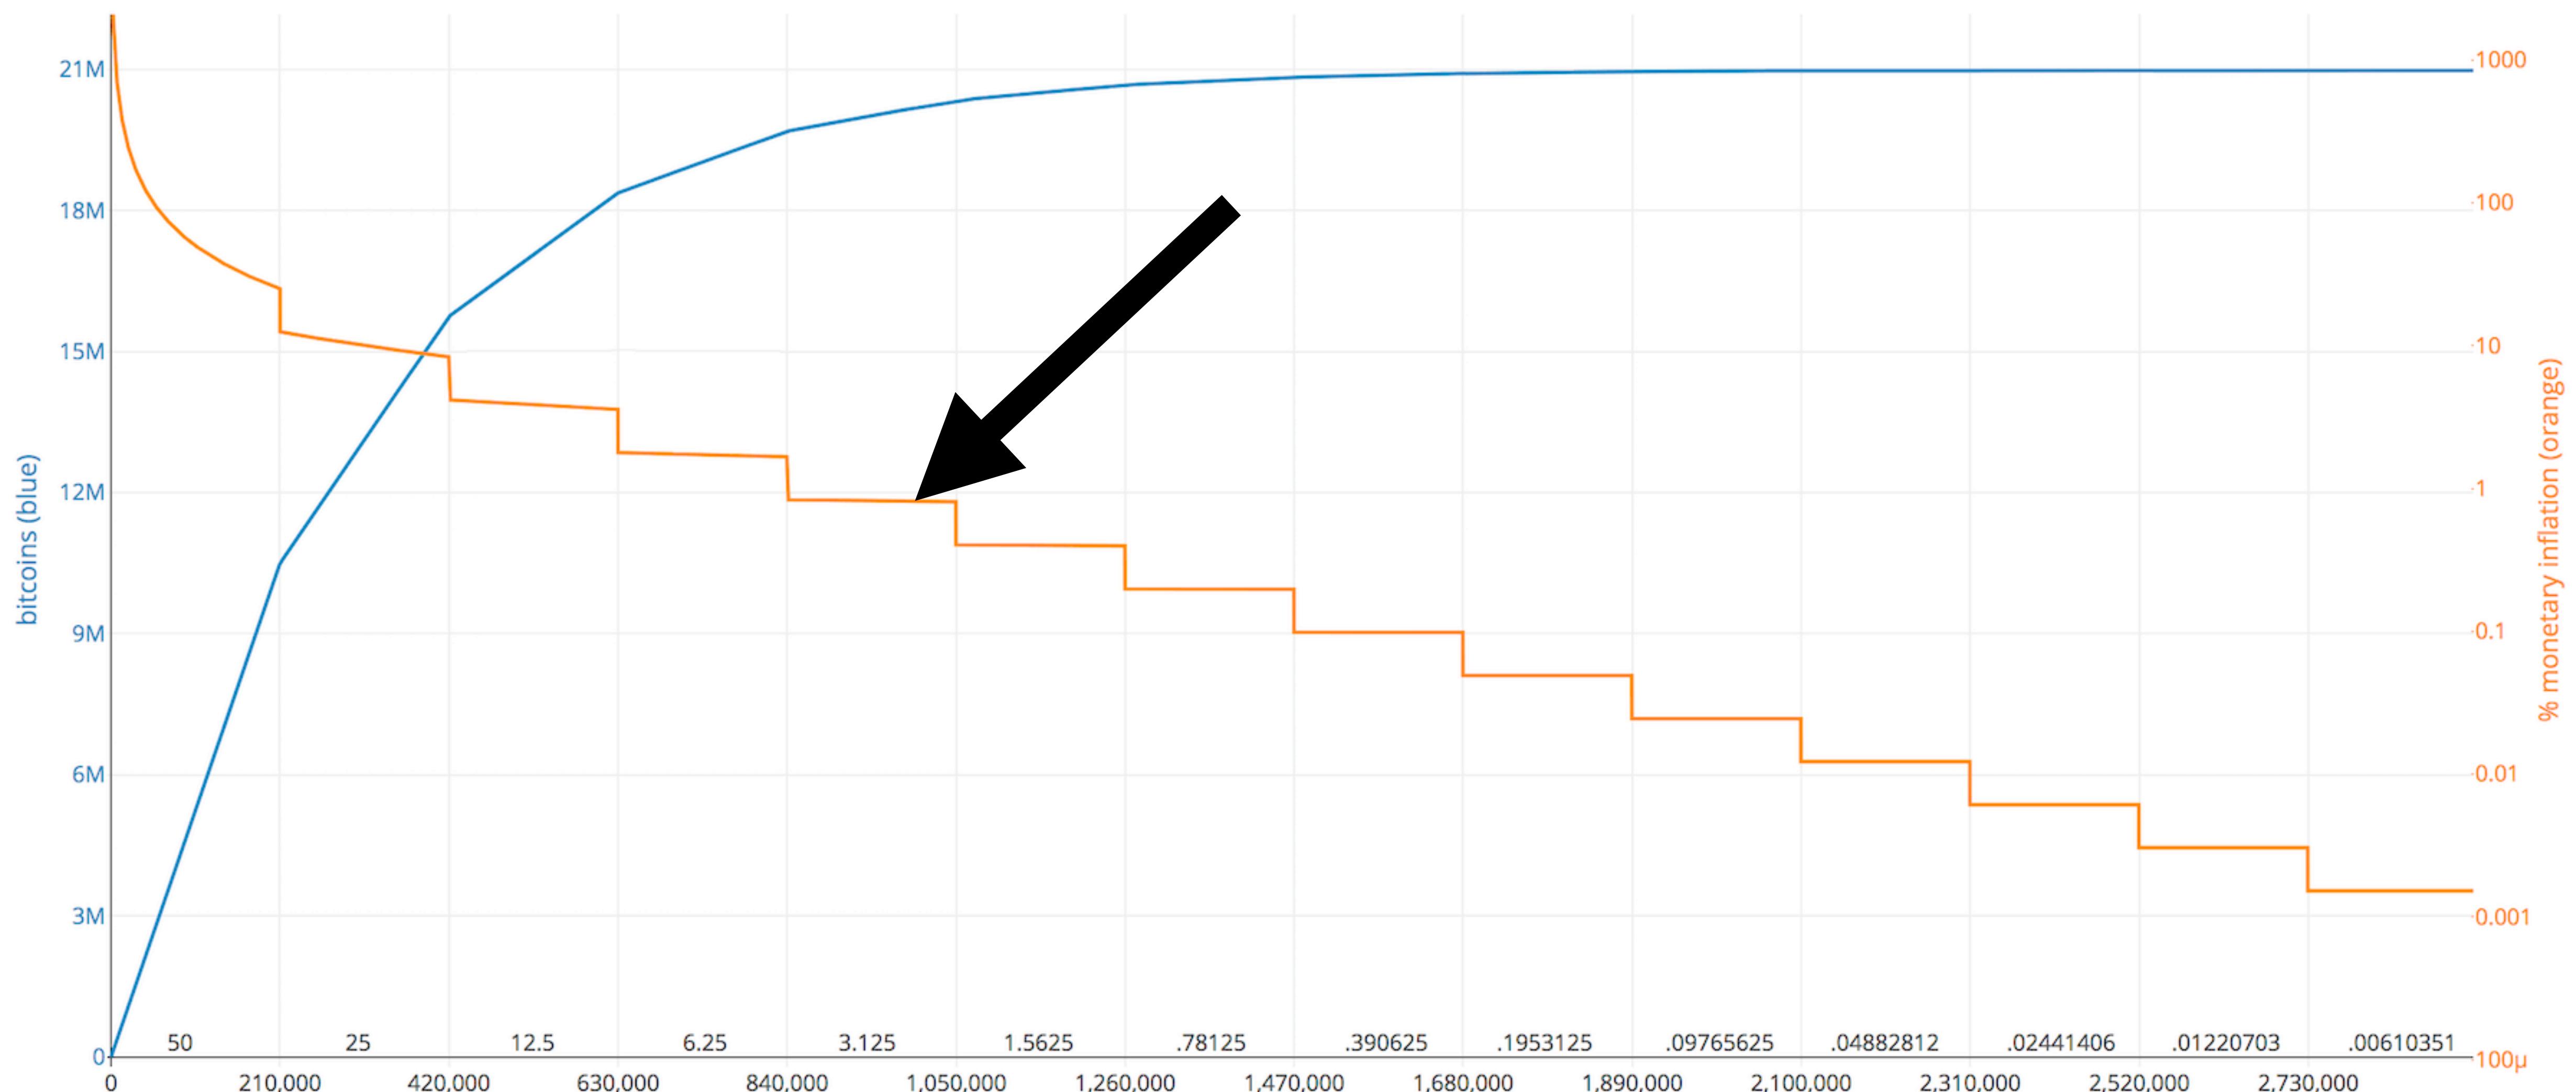
2 reasons

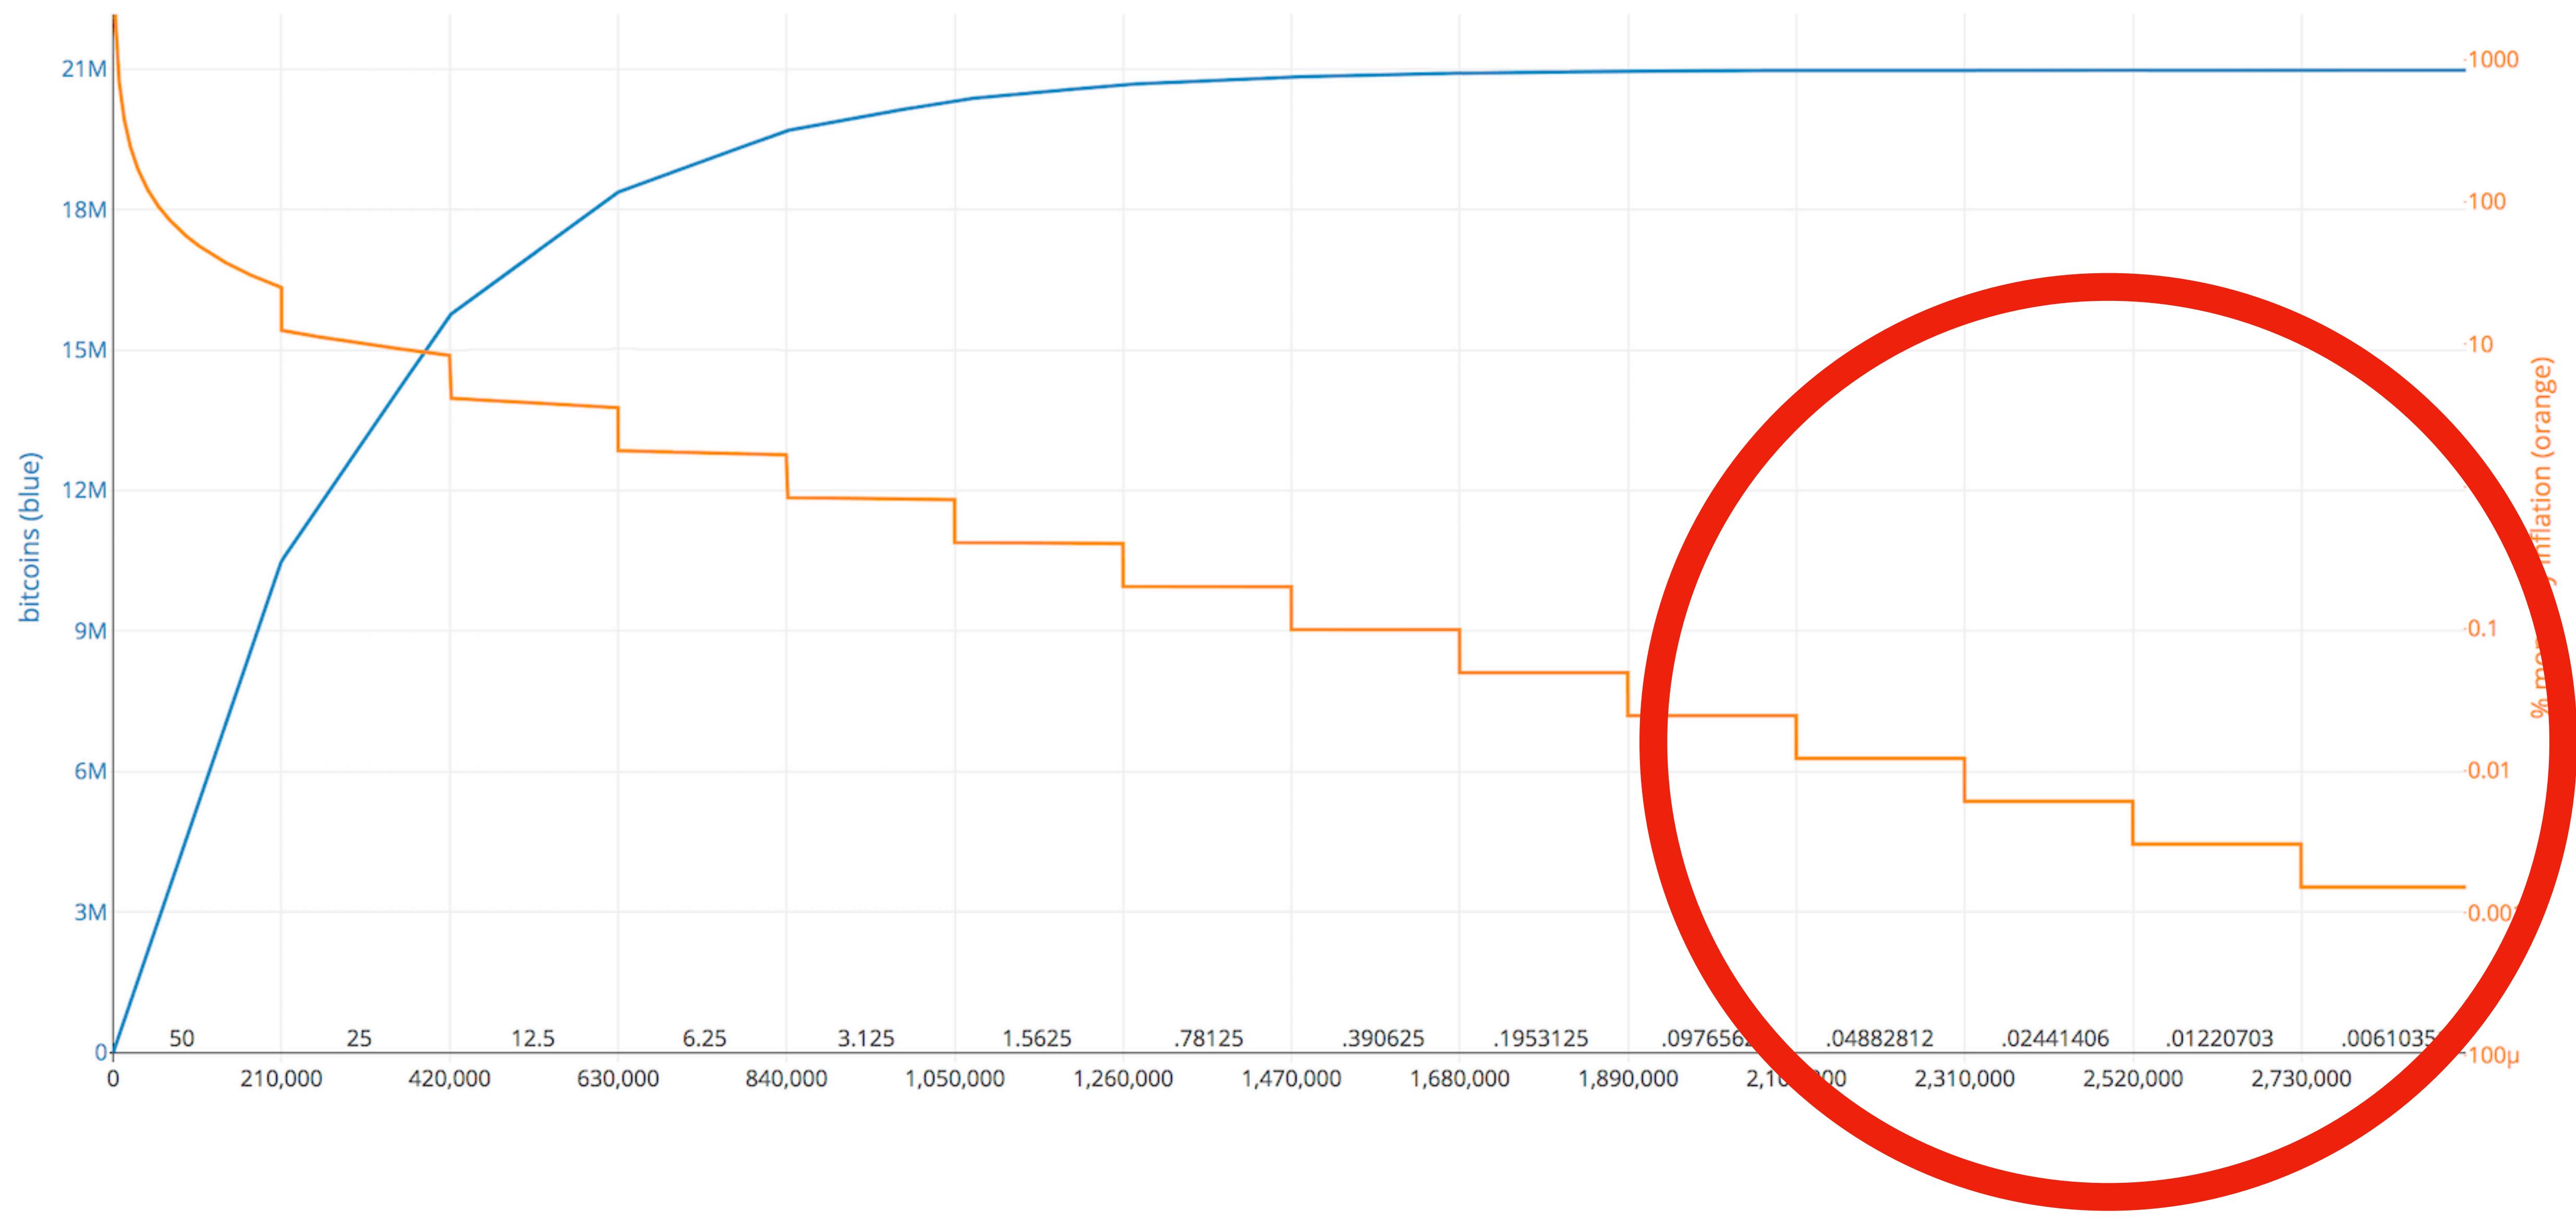
- Economical
- Technical

# Economical reason

## Supply and demand

- Are there gonna be enough fees when we increase the block size?





# Economical reason

## Supply and demand

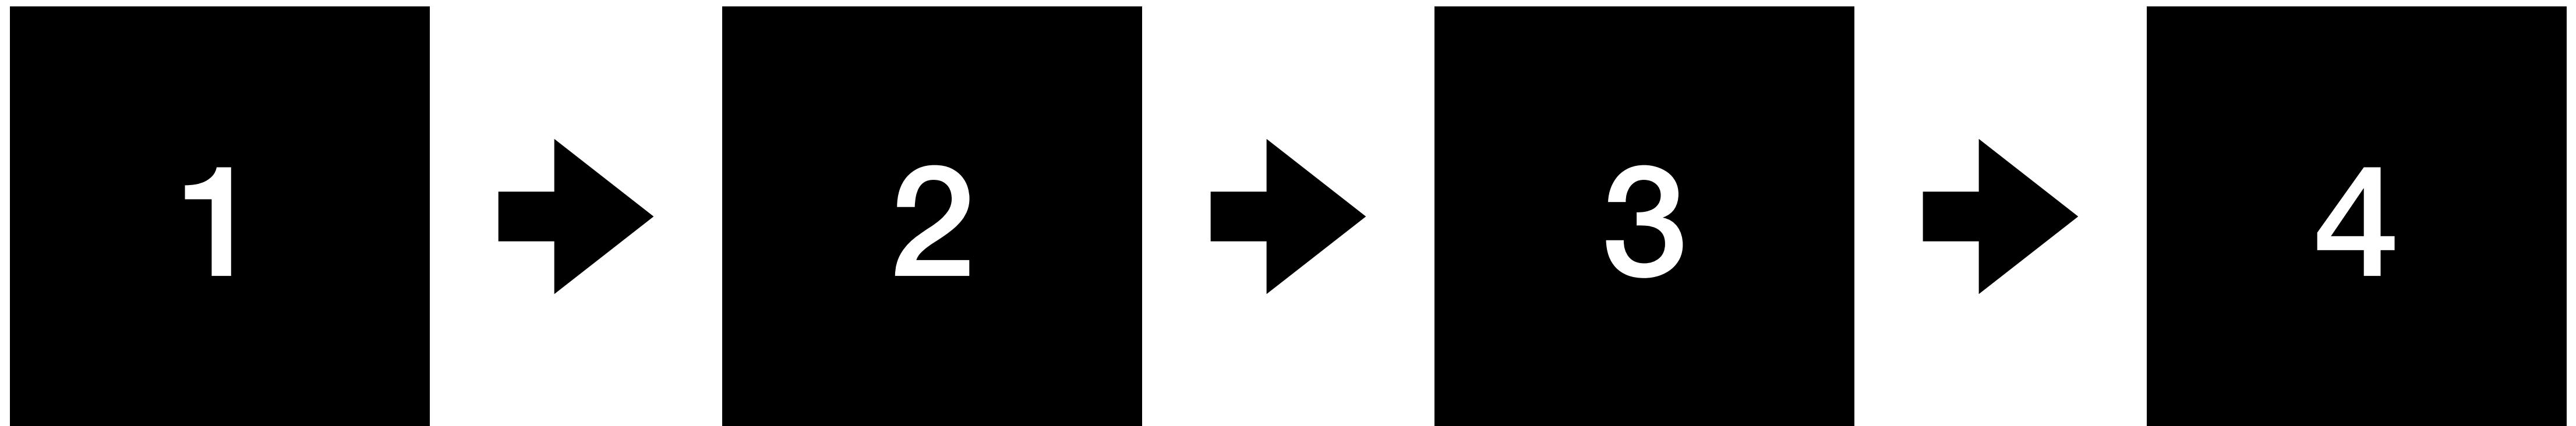
- Current *double spends preventions* are provided by the block subsidy

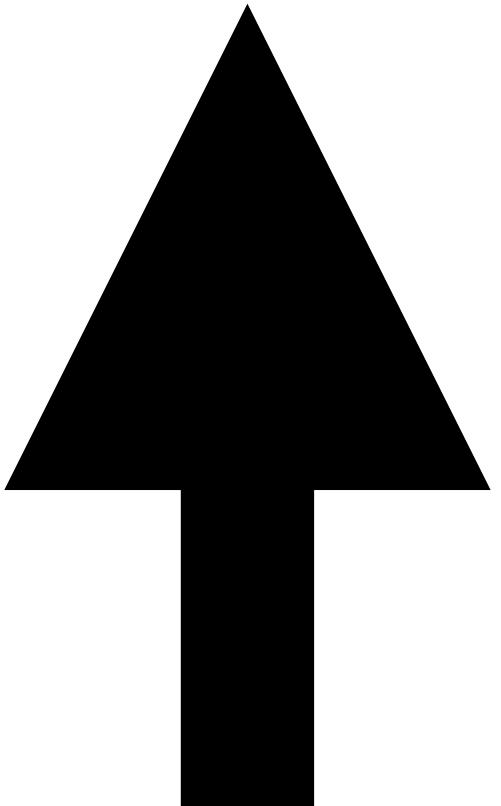
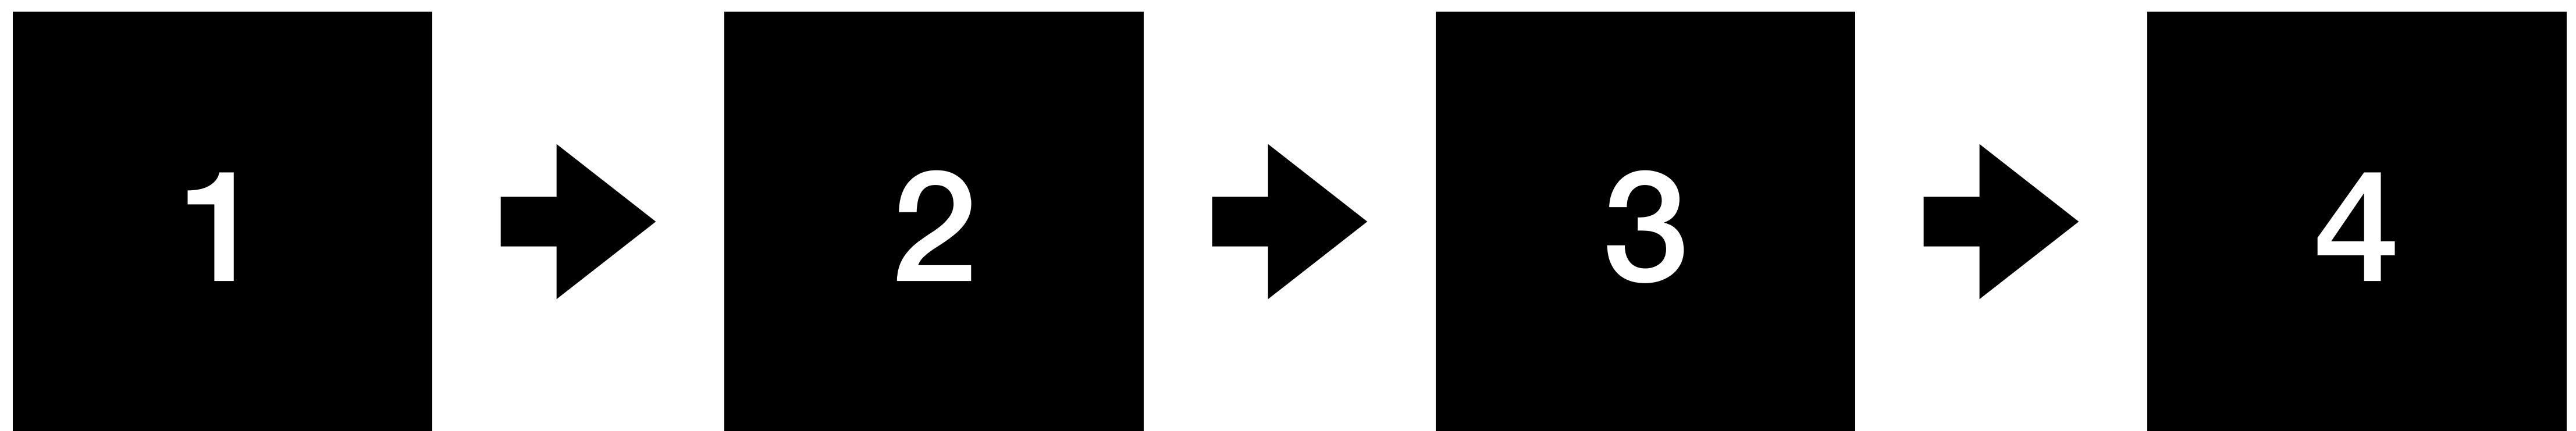
# Economical reason

## Supply and demand

- Current *double spends preventions* are provided by the block subsidy
- As the subsidy drops, the fees need to increase to keep the same amount of security

# Example of a double spend due to less fees





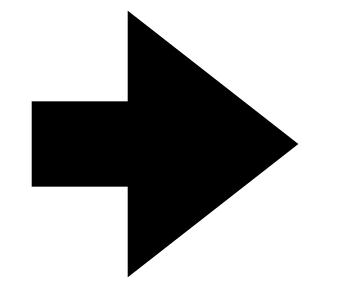
50 btc TX

- Alice buys gold from Bob with 50btc

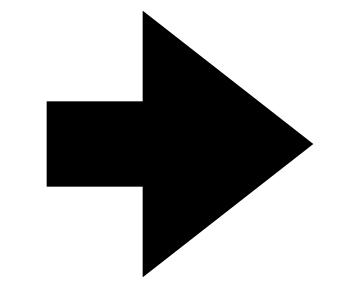
- Alice buys gold from Bob with 50btc
- Alice bribes miner Charlie 10btc to double spend Bob

- Alice buys gold from Bob with 50btc
- Bob bribes miner Charlie to double spend Alice
- Alice's profits - 40 btc + gold
- Miner Charlies profits - 10 btc + block subsidy

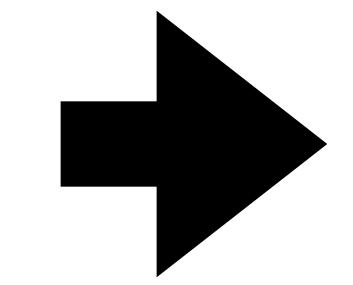
1



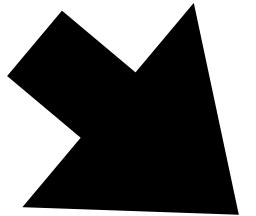
2



3

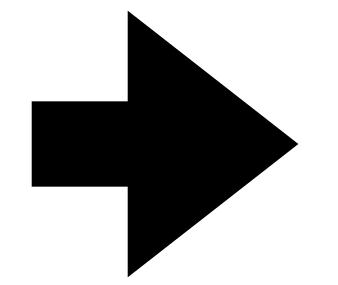


4

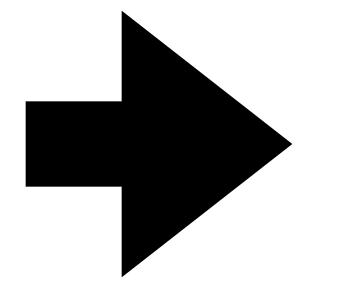


4b

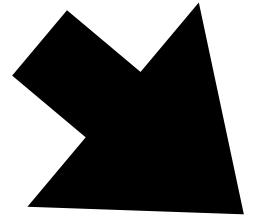
2



3

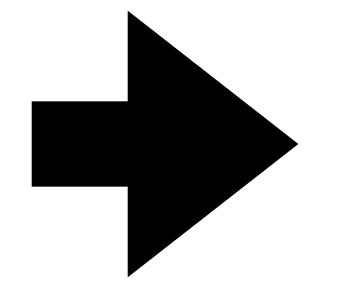


4

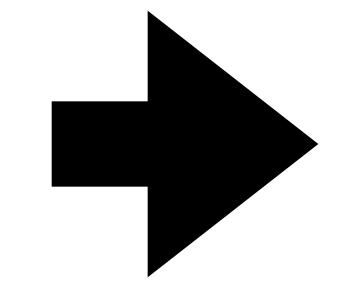


4b

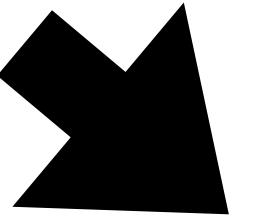
2



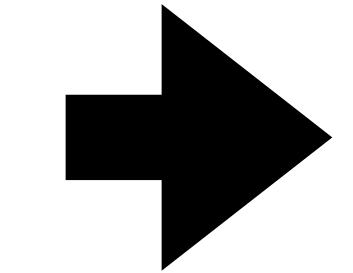
3



4



4b



5

# Economical reason

## Supply and demand

- Need enough fees to prevent these attacks

# Economical reason

## Supply and demand

- Need enough fees to prevent these attacks
- Block size ↑, Fees ↓

**Currently no research ongoing  
for the economical reason**

# **Technical reason**

## **Trustlessness**

- Can anyone that wants to use Bitcoin trustlessly do it?

# **Technical reason**

## **Trustlessness**

- Can anyone that wants to use Bitcoin trustlessly do it?
- Block size ↑, Cost of full nodes ↑

# Technical reason

## Trustlessness

- Download ~600GB worth of blocks
- Validate that the transaction is spending Bitcoins that exist
- Validate that the transaction is able to spend that Bitcoin

# **Technical reason**

## **The bottlenecks**

- Internet
- Disk i/o
- CPU

# Internet

## Downloading blocks/transactions

- Download ~600GB on first startup

# **Internet**

## **Downloading blocks/transactions**

- Download ~600GB on first startup
- **100mbps = ~800 min**

# **Internet**

## **Downloading blocks/transactions**

- Download ~600GB on first startup
- 100mbps = ~800 min
- 1gbps = ~80 min

# **Internet**

## **Downloading blocks/transactions**

- Download ~600GB on first startup
- 100mbps = ~800 min
- 1gbps = ~80 min
- 10gbps = ~8 min

**Internet bandwidth is ok**

# Disk performance

Read and write to the UTXO set

- UTXO set - Collection of all the Bitcoins available to be spent

# What's a UTXO set?

- TXO - Transaction Output

# Typical transaction

[8ad63210d059908229f00ad177469d9f16a5d37a7ea64cede9c92684ff7cf06e](#)

DETAILS



#0 [f758ad929fbefac99ef0d1f8b13990fa75efb517ef6c919178c8af7308f145](#) 0.04852489 BTC  
07:1

#0 [17dSwJytZBszAafyWrK8Pue7rsRh9s5xeJ](#) 0.04033963 BTC

#1 [bc1qwqdg6squ38e46795at95yu9atm8azzmyvckulcc7kytlcckxswvvzej](#) 0.00758526 BTC

1 CONFIRMATION 0.04792489 BTC

# The outputs

[8ad63210d059908229f00ad177469d9f16a5d37a7ea64cede9c92684ff7cf06e](#)

DETAILS

+

#0 [f758ad929fbefac99ef0d1f8b13990fa75efb517ef6c919178c8af7308f145](#) 0.04852489 BTC  
07:1

#0 [17dSwJytZBszAafyWrK8Pue7rsRh9s5xeJ](#) 0.04033963 BTC

#1 [bc1qwqdg6squ38e46795at95yu9atm8azzmyvckulcc7kytlcckxswvvzej](#) 0.00758526 BTC

1 CONFIRMATION 0.04792489 BTC

# What's a UTXO set?

- TXO - Transaction Output
- UTXO - Unspent TXO

# What's a UTXO set?

- TXO - Transaction Output
- UTXO - Unspent TXO
- UTXO set - All the UTXOs

# Disk performance

Read and write to the UTXO set

- UTXO set uses a key-value database (leveldb)
- leveldb does random read and writes

# Disk performance

## Read and write to the UTXO set

- At least 59 bytes per UTXO
- Average tx is 1 input, 2 output
- For 2000txs: ~2000 read, ~4000 write. (2000-3000 txs per block)

# 기술의 경이로움이 시작된다

## Platinum P41 SSD



SK하이닉스 Platinum P41 SSD는 업계 최고 수준의 성능으로  
차세대 컴퓨팅 시스템을 완성하는 제품입니다. PCIe 4세대 NVMe 인터페이스와  
176단 NAND의 한층 더 강화된 기술력을 토대로 최강의 데이터 전송 속도와  
성능을 제공합니다. 이제 차원이 다른 성능을 경험해 보세요.

## 성능

---

2TB

순차성능  
(최대)

읽기 : 7,000 MB/s  
쓰기 : 6,500 MB/s

랜덤성능  
(최대)

읽기 : 1,400K IOPS  
쓰기 : 1,300K IOPS

TBW

최대 1,200TBW

1TB

읽기 : 7,000 MB/s  
쓰기 : 6,500 MB/s

읽기 : 1,400K IOPS  
쓰기 : 1,300K IOPS

최대 750TBW

500GB

읽기 : 7,000 MB/s  
쓰기 : 4,700 MB/s

읽기 : 960K IOPS  
쓰기 : 1,000K IOPS

최대 500TBW

- 1,400 IOPS = ~5.6MB/s read
- 1,300 IOPS = ~5.2MB/s write
- Per block 0.12MB read, 0.24MB write

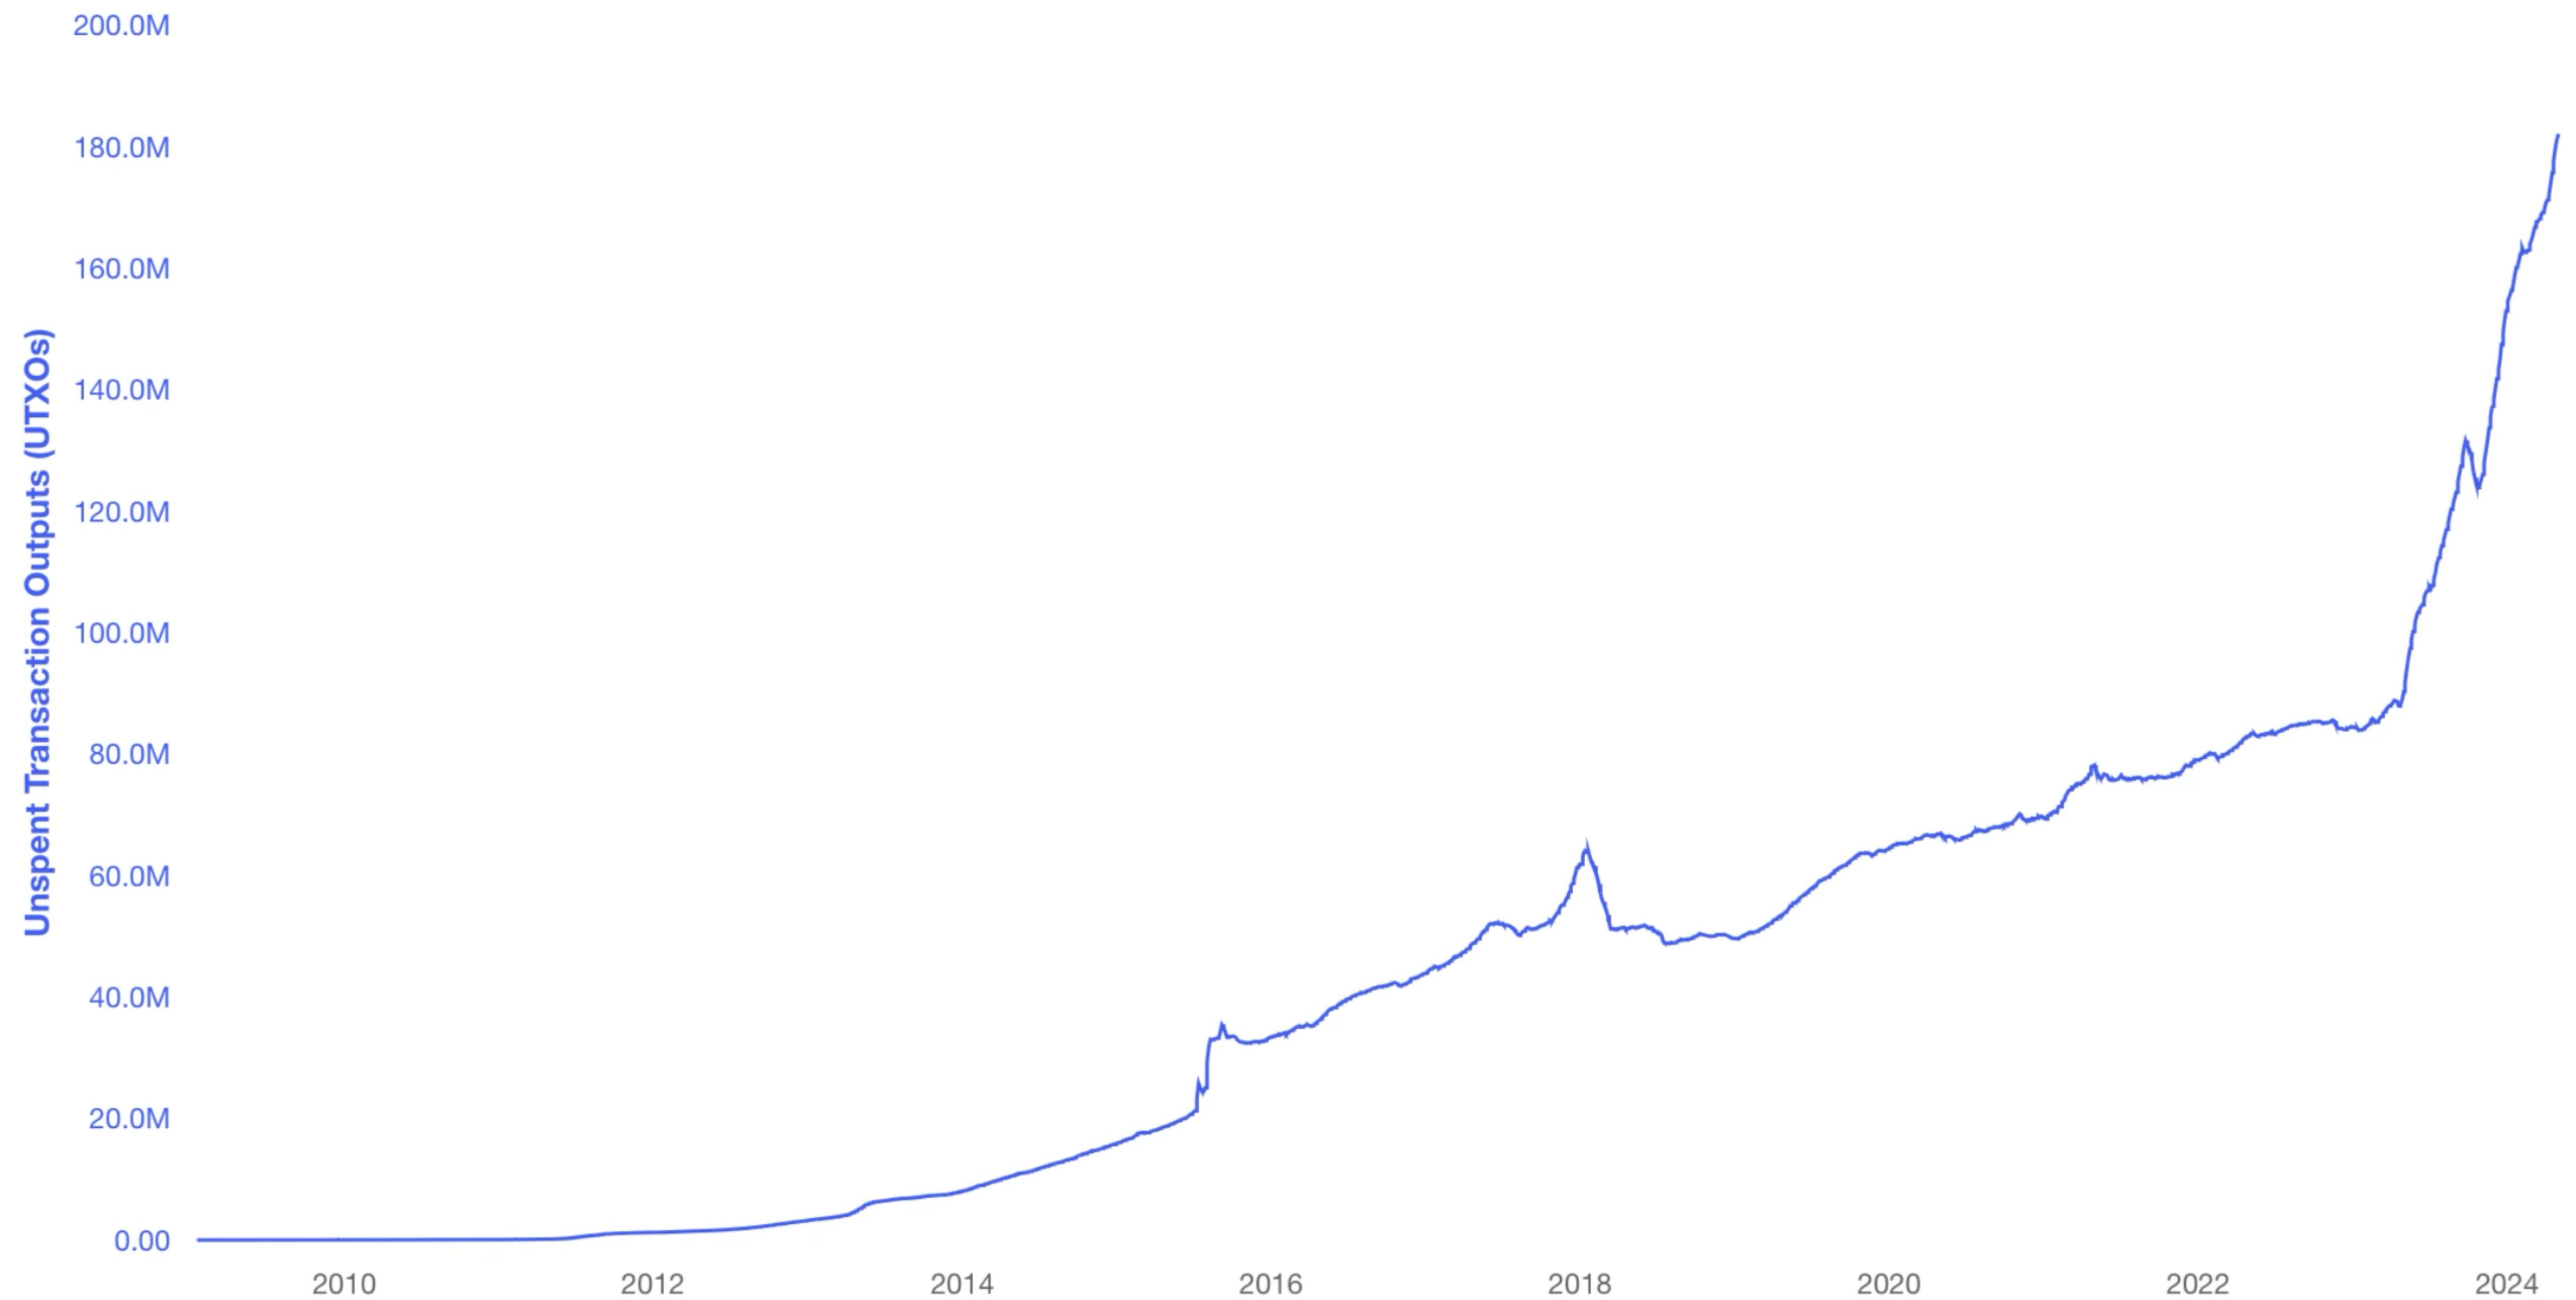
(1,400 \* 4) / 1000. OS page size of 4(mb). Divide by 1000 to get MB/s instead of byte/s

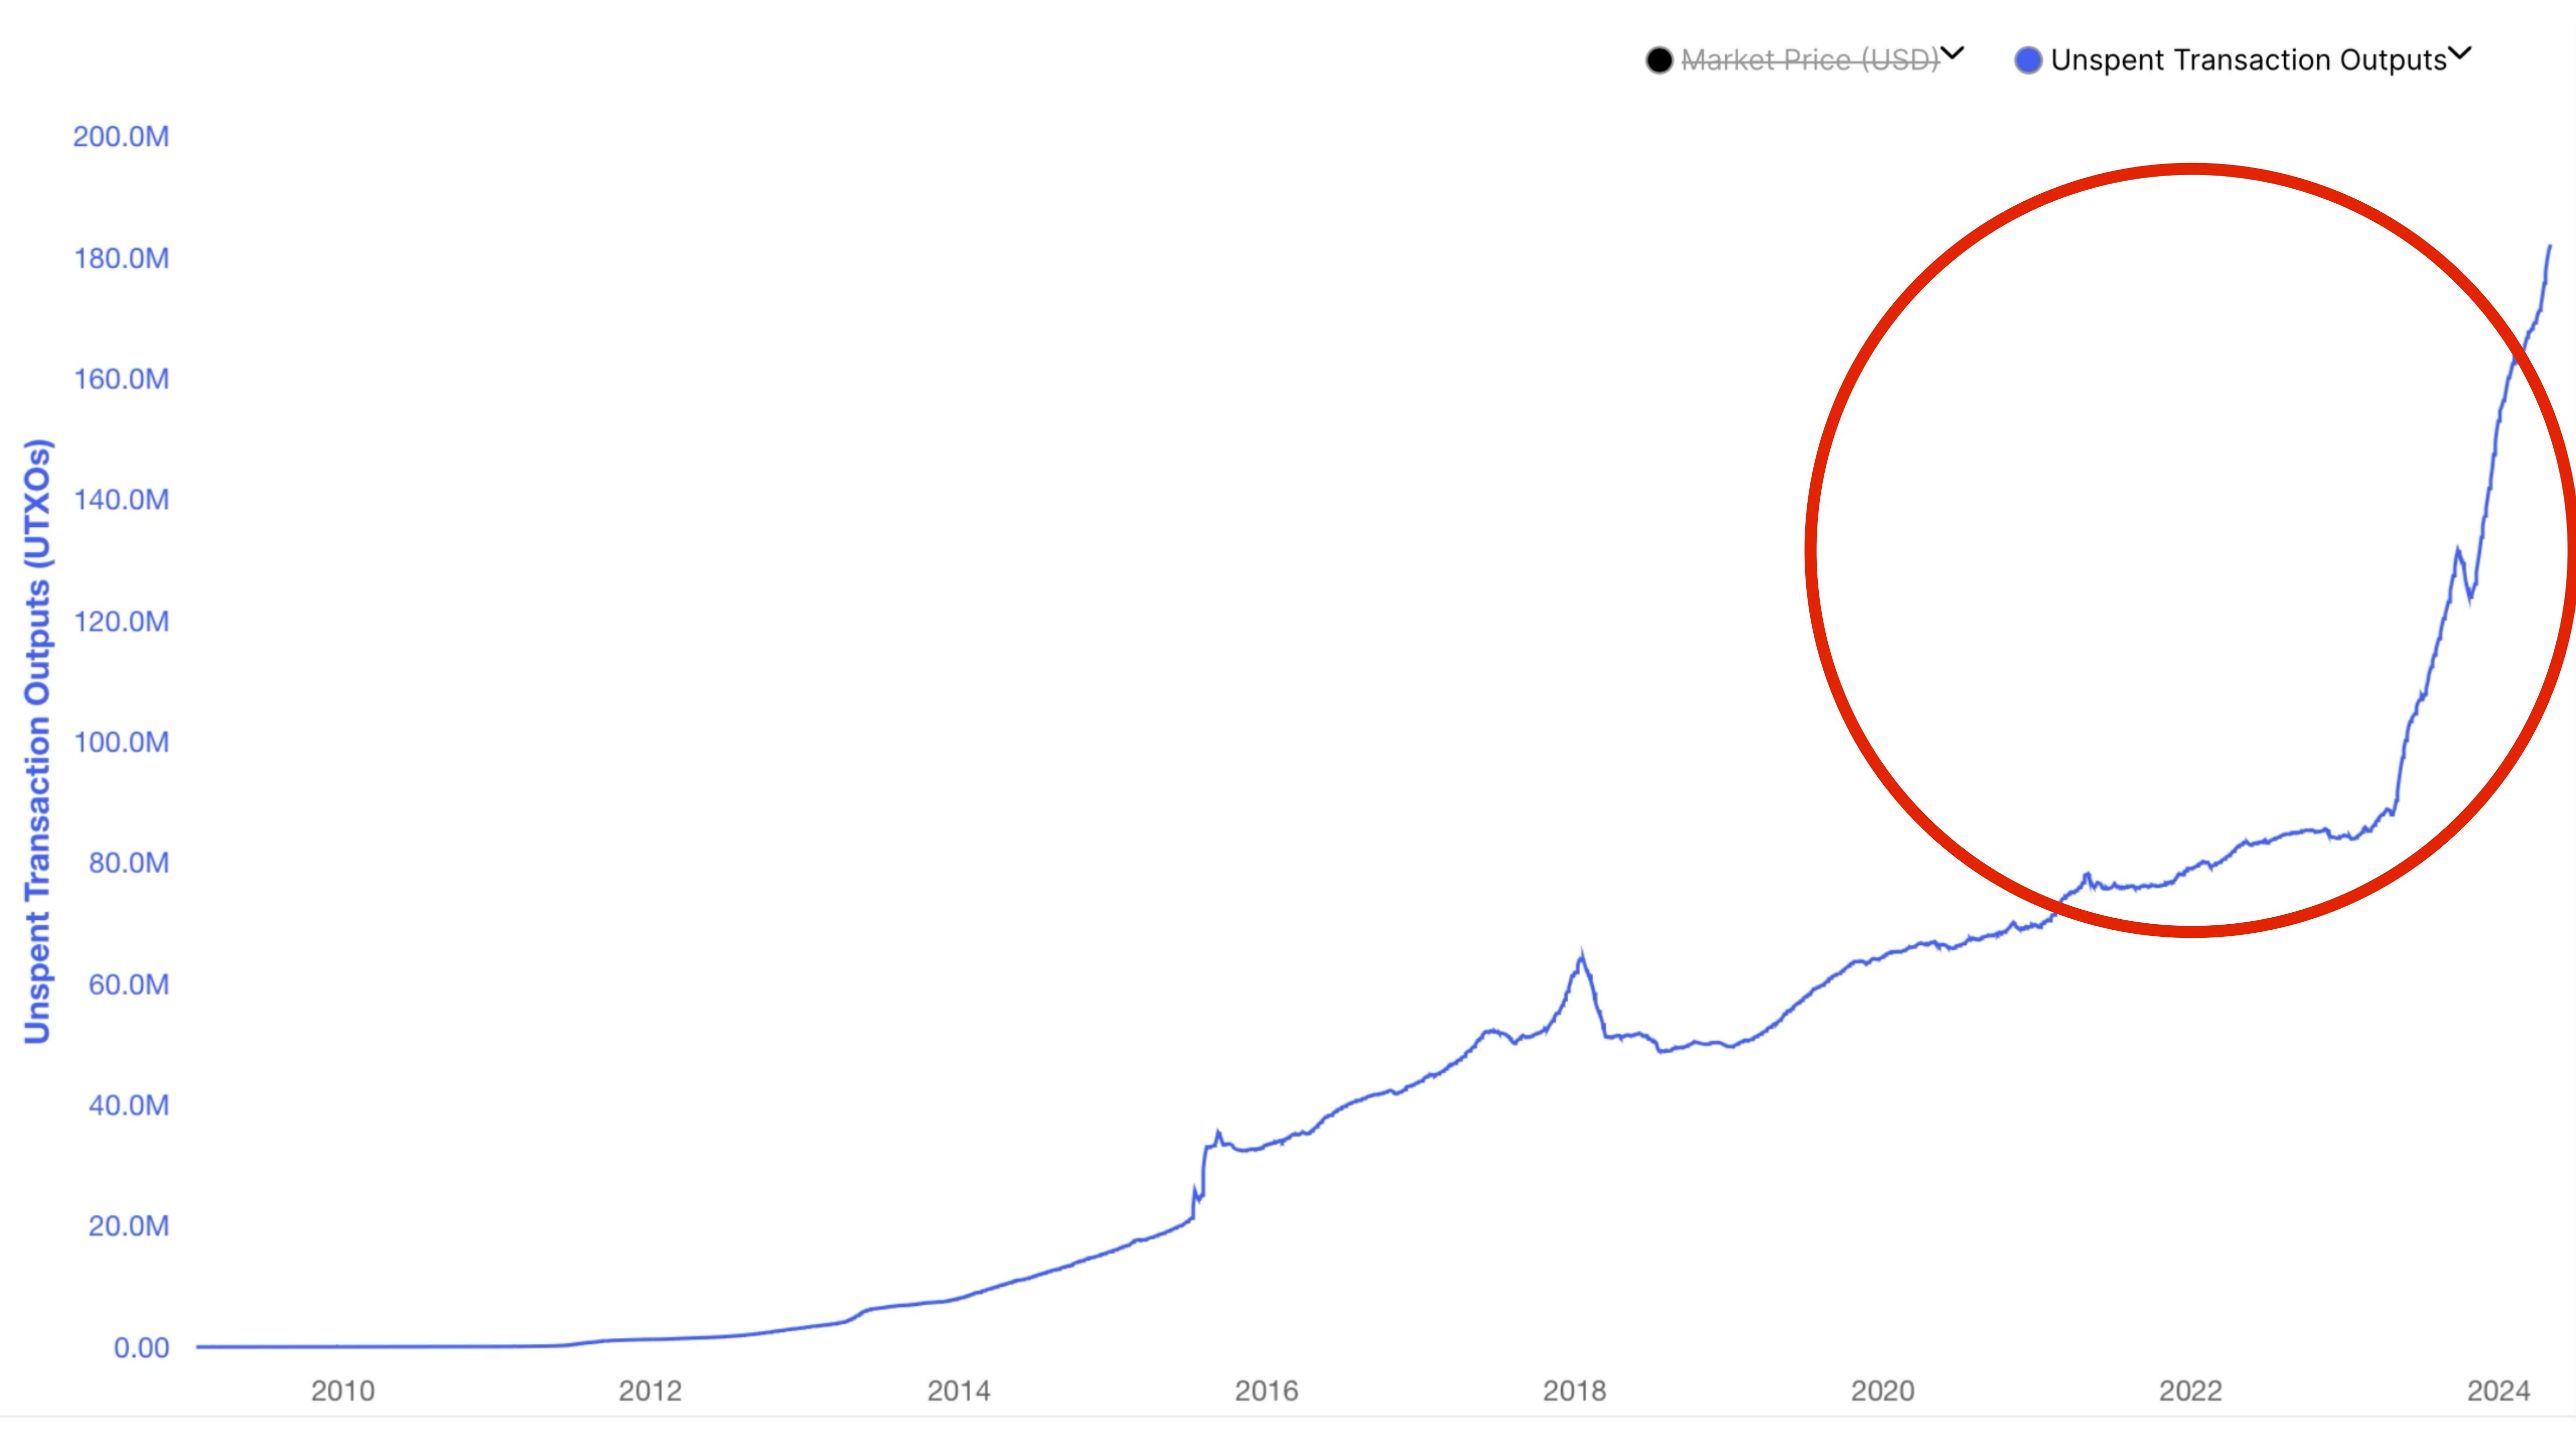


- 73 IOPS = ~0.29MB/s read&write
- Per block: 0.12MB read, 0.24MB write

Market Price (USD) ▾

Unspent Transaction Outputs ▾





**As the UTXO set grows,  
leveldb performance drops**

**UTXO set size is unlimited**

**8billion UTXO, 59 bytes each**

**472GB**

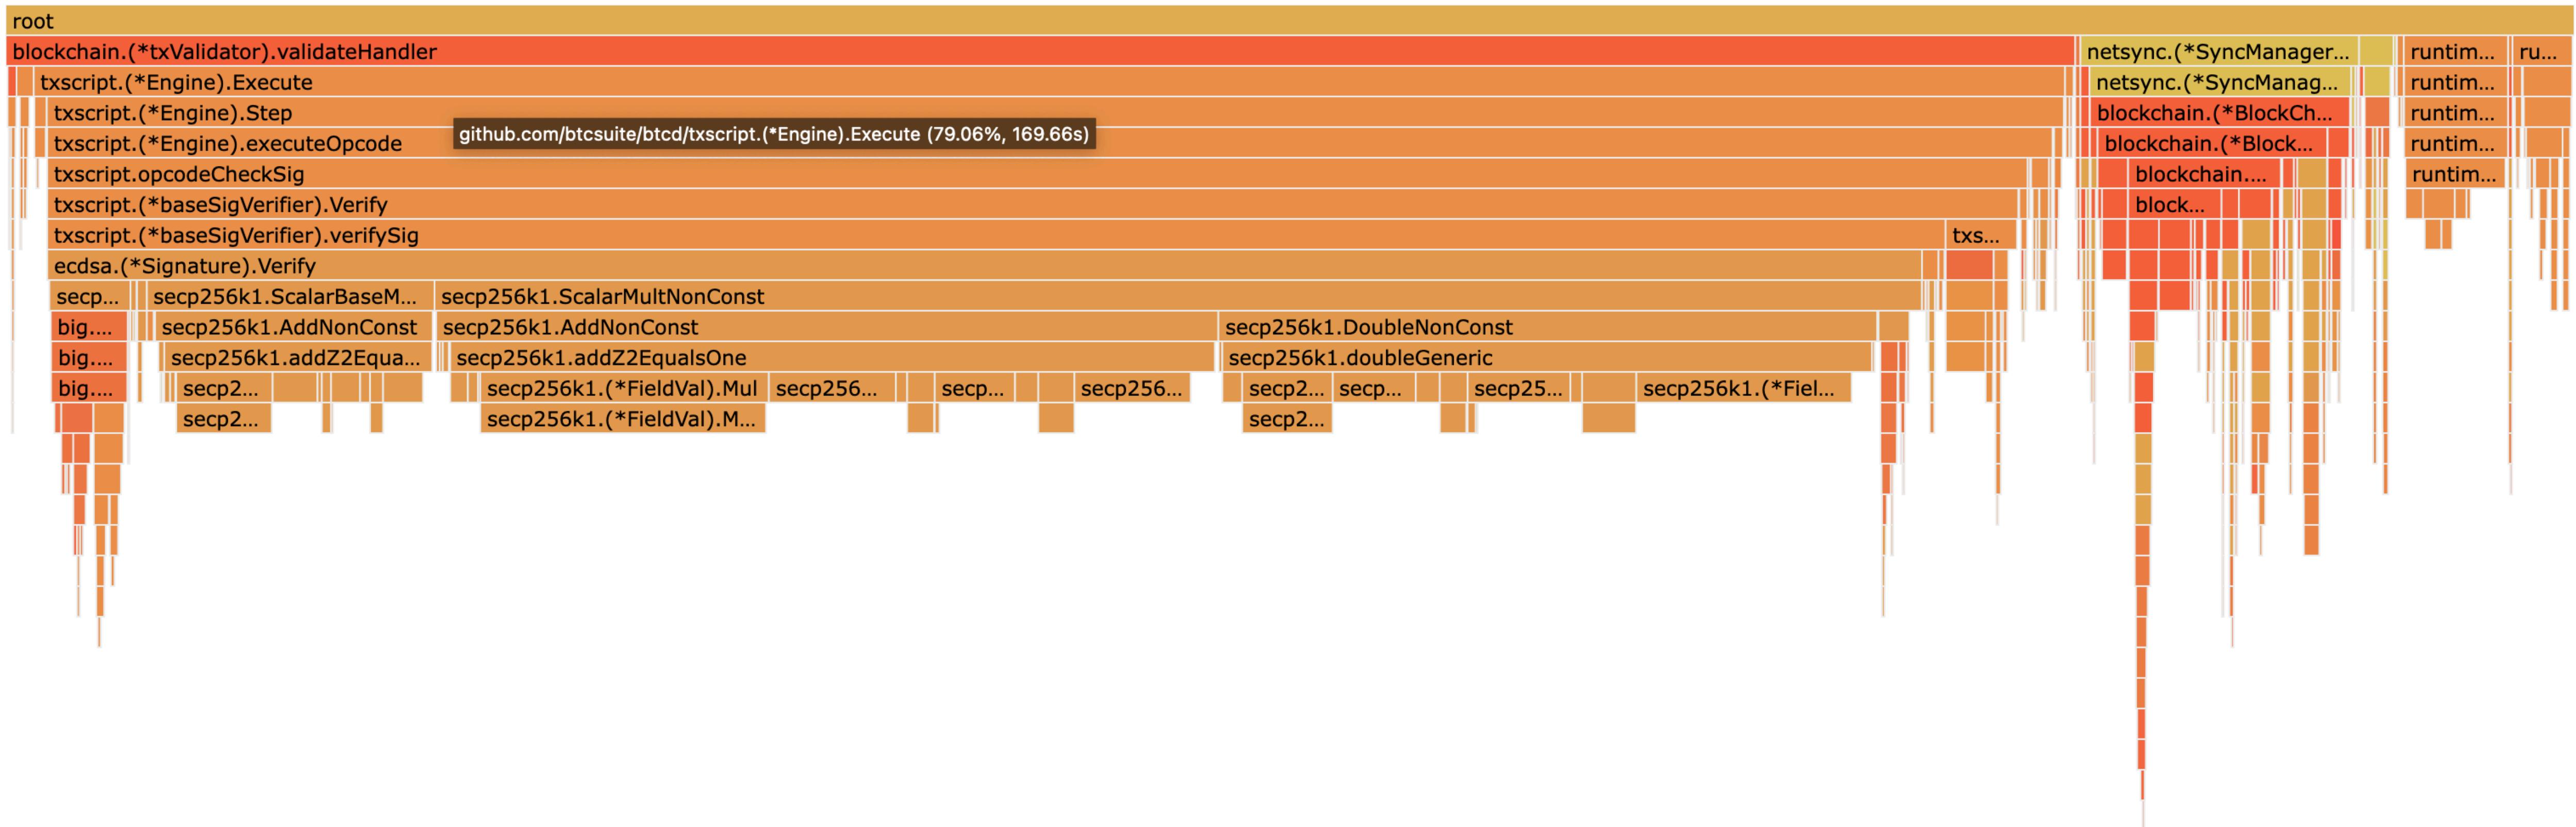
**HDDs are already  
bottlenecked**

# CPU Performance

## Transaction signature validation

- Currently most of the block validation is spent on signature validation

github.com/btcsuite/btcd/txscript.(\*Engine).Execute (79.06%, 169.66s)



**80% of the time  
spent on tx validation**

**Based on btcd. Bitcoin Core does much better.**

# CPU Performance

## Transaction signature validation

- CPU cores are still increasing
- Taproot txs can be validated together

# libsecp256k1

[build](#) [passing](#) [dependencies](#) [none](#) [irc.libera.chat](#) [#secp256k1](#)

Optimized C library for ECDSA signatures and secret/public key operations on curve secp256k1.

This library is intended to be the highest quality publicly available library for cryptography on the secp256k1 curve. However, the primary focus of its development has been for usage in the Bitcoin system and usage unlike Bitcoin's may be less well tested, verified, or suffer from a less well thought out interface. Correct usage requires some care and consideration that the library is fit for your application's purpose.

## Features:

- secp256k1 ECDSA signing/verification and key generation.
- Additive and multiplicative tweaking of secret/public keys.
- Serialization/parsing of secret keys, public keys, signatures.
- Constant time, constant memory access signing and public key generation.
- Derandomized ECDSA (via RFC6979 or with a caller provided function.)
- Very efficient implementation.
- Suitable for embedded systems.
- No runtime dependencies.
- Optional module for public key recovery.
- Optional module for ECDH key exchange.
- Optional module for Schnorr signatures according to [BIP-340](#).

**CPU already bottlenecked  
Improvements possible**

What do we need to do to increase the block size?

**Signature validation speed ↑**  
**UTXO read/write speed ↑**

**How do we improve the UTXO  
read/write speed?**

# **Utreexo**

**What I work on**

**Let's turn the UTXO set into  
a merkle tree**

**Same performance on  
HDD and SSD**

# Disk i/o X

**Utreexo eliminates random reads/writes**

# Benchmarking this Testing methodology

- Compare the performance drop on a pruned node vs a non-pruned node

**36% slower  
on testnet**

**Bitcoin Core**

**3% slower on  
testnet  
Utreexo node**

# What we write with Utreexo

## No reads

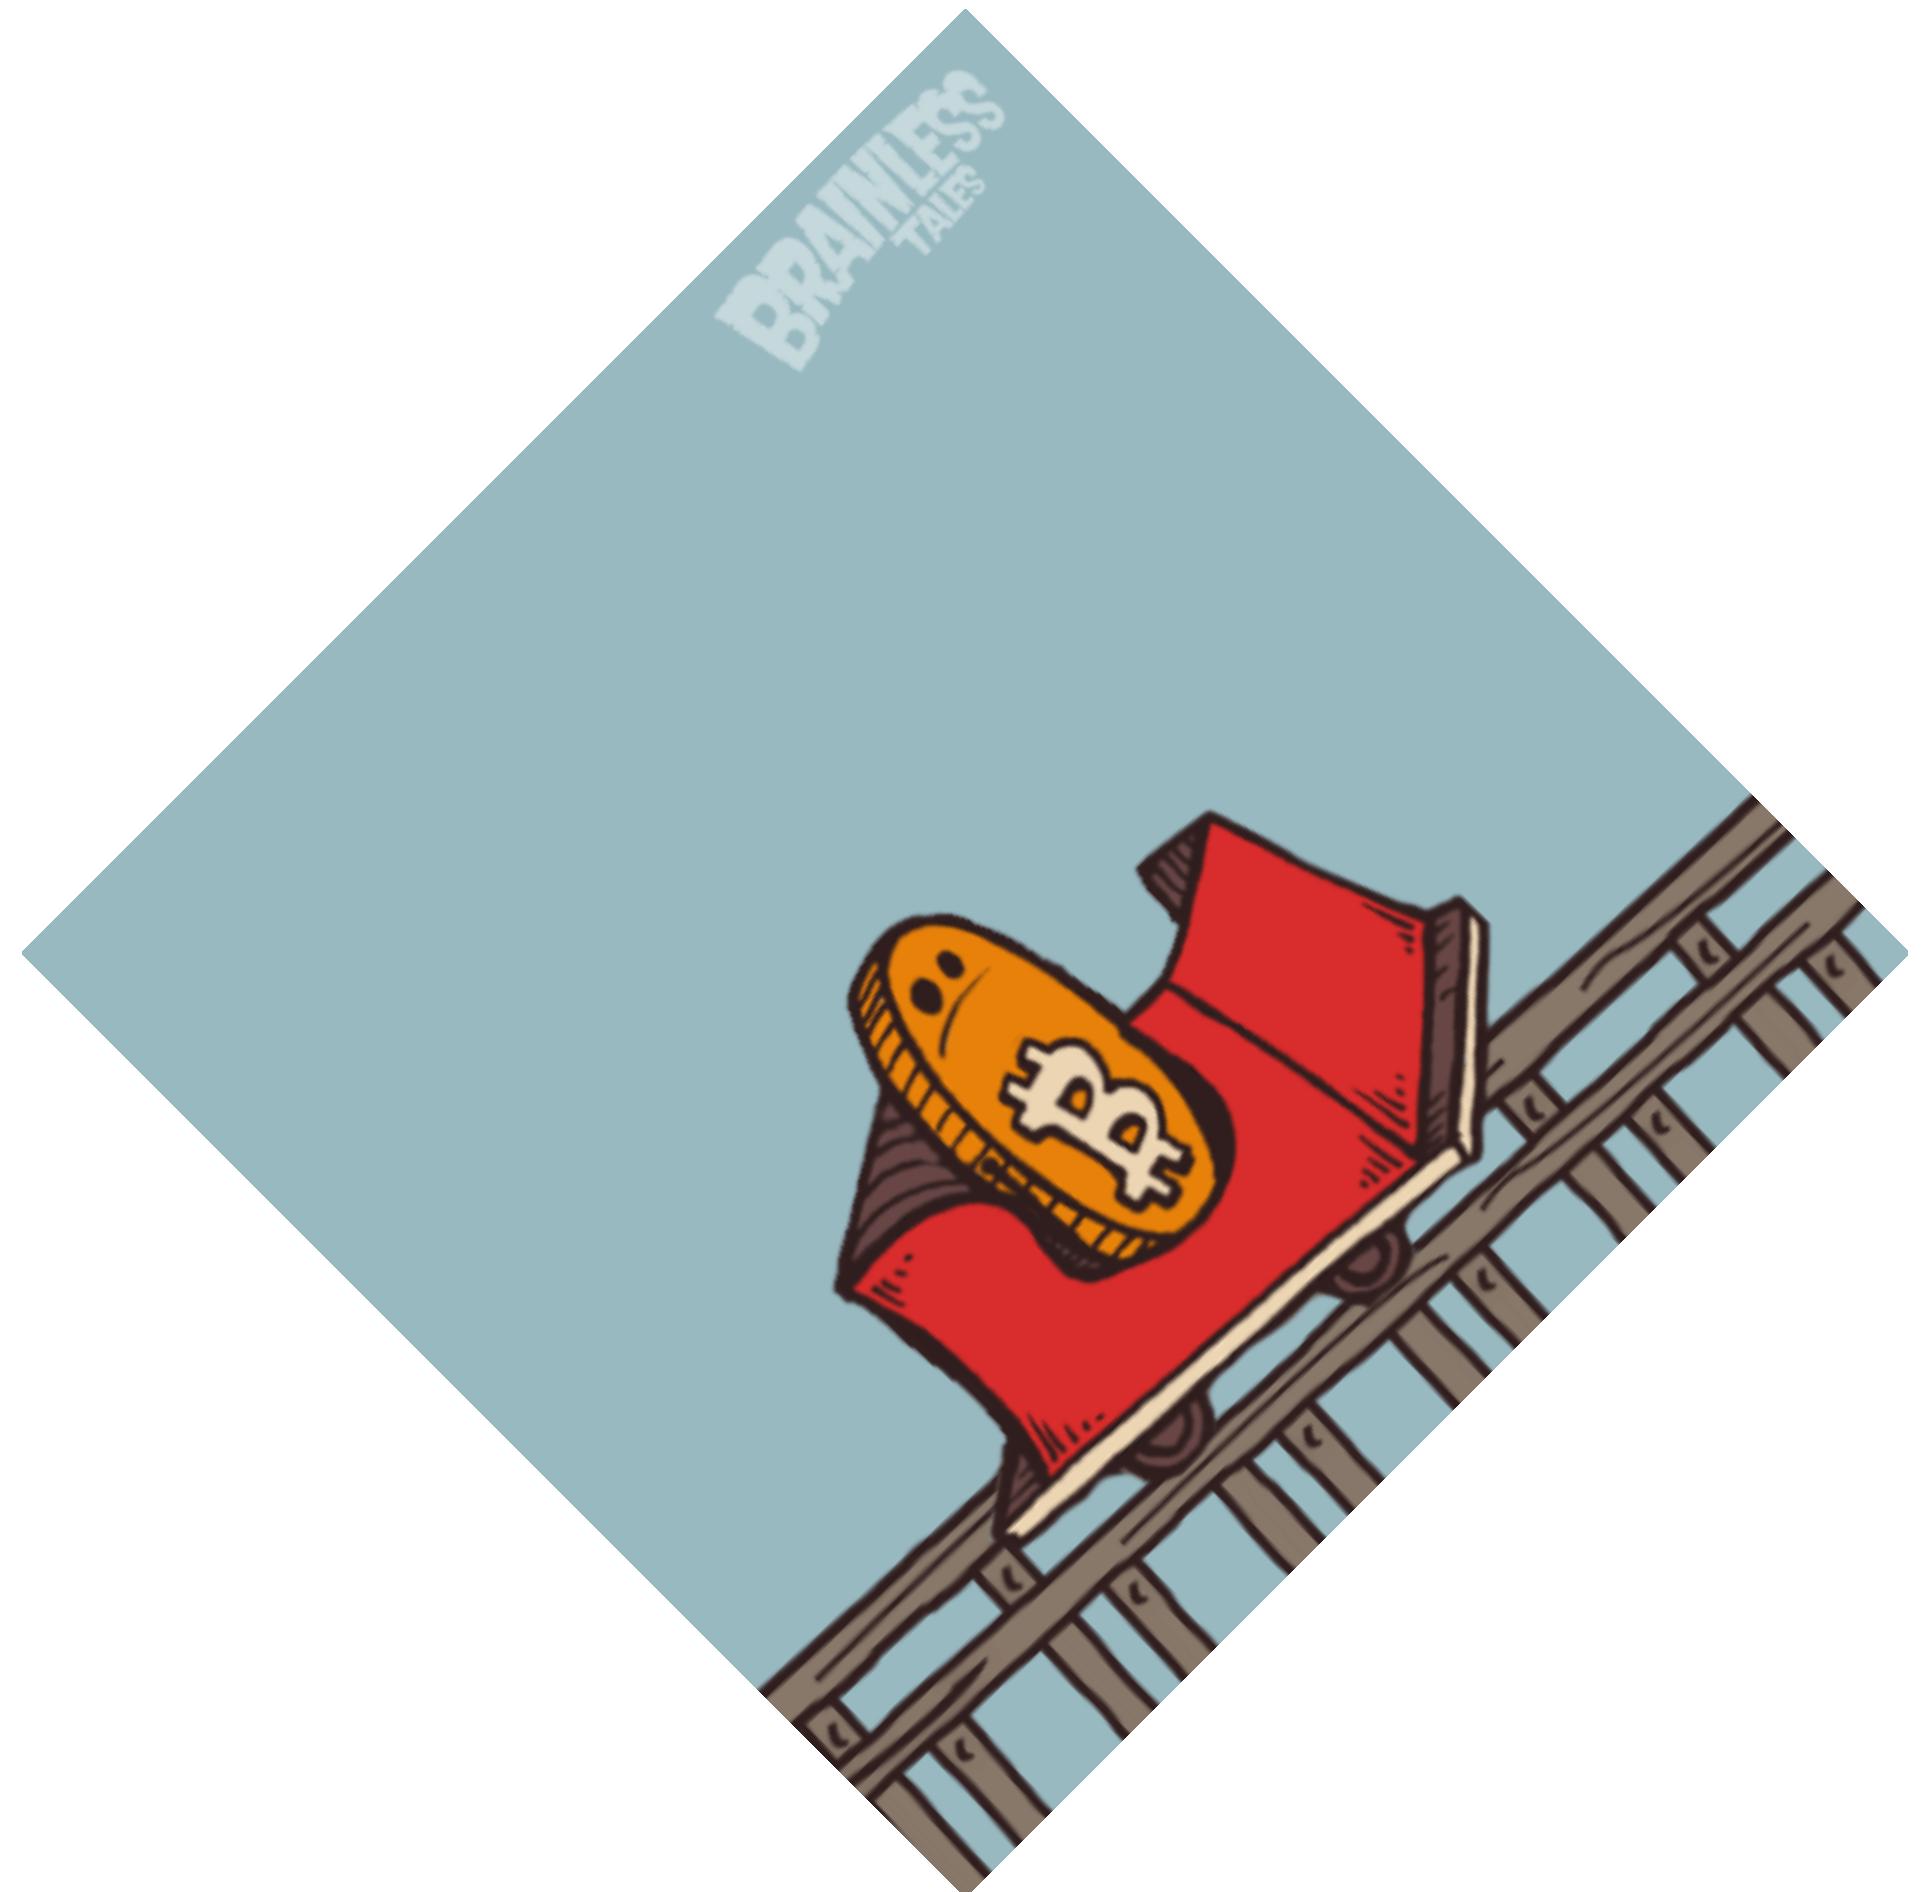
```
{443000, newHashFromStr("000000000000000047efc3c126c3398cfdef1609682d1492cd2909e3c0a17f"),  
    43243051, []*chainhash.Hash{  
        newLeafHashFromStr("df31c496a54a8021c38c77809da15a8bd3968970315f5a56c07f0fd9763262ff"),  
        newLeafHashFromStr("755f4cfbf5727c3e14aa6dcc6e5ab150c61101cead06fade5a886ee8d1dd377"),  
        newLeafHashFromStr("c572a92c7cd9e14ce5d7c94473ef90bee21856654eed5535314b3c117006bc83"),  
        newLeafHashFromStr("9e85567347ef0ea38c1797d8ed8e399ed0d99a07025388f6e06791ff6b88643c"),  
        newLeafHashFromStr("ab4efc738bd7b8281689e899979d44ed0ca8fe1489efa0f5a09f5b1c276fe792"),  
        newLeafHashFromStr("1c763d8f490baef0b756597dfc2da5a31d0aee324da6314b52f21707351cf51a"),  
        newLeafHashFromStr("29fcd9989939a17e5d113fb1f247f3bd05606488bcb48c9cfad55a57006f2248"),  
        newLeafHashFromStr("e79f840fe1bcd6b638a53b637885dda7ea3da2b2b5f1f5447ebf4d54fc0a2e"),  
        newLeafHashFromStr("02f33c3ce0d684e644410fa35fee4c7572454facc710389ea0f69a00e27c1780"),  
        newLeafHashFromStr("4e29ad940088ab42991bd99d7943a68ea38d498f8435c4bea21b8252788b5a7d"),  
        newLeafHashFromStr("feabfc6565e918827df8ade12d0885e7da28bb81410c71a0fc92da28e5538e"),  
        newLeafHashFromStr("42275481edcd150c0891f8f7486941e7ecc9dbec0d1554d1ae3027d8523571b4"),  
        newLeafHashFromStr("84dd84bbb564ab9e878ea774d66a2cb818c2b74680974e01a07dde68bd064bfd"),  
        newLeafHashFromStr("19a64f2f77403d4a1e3482b3a1801aa360120f6429f5ed47261fc8676920d9d"),  
    },  
},
```

# Tradeoffs?

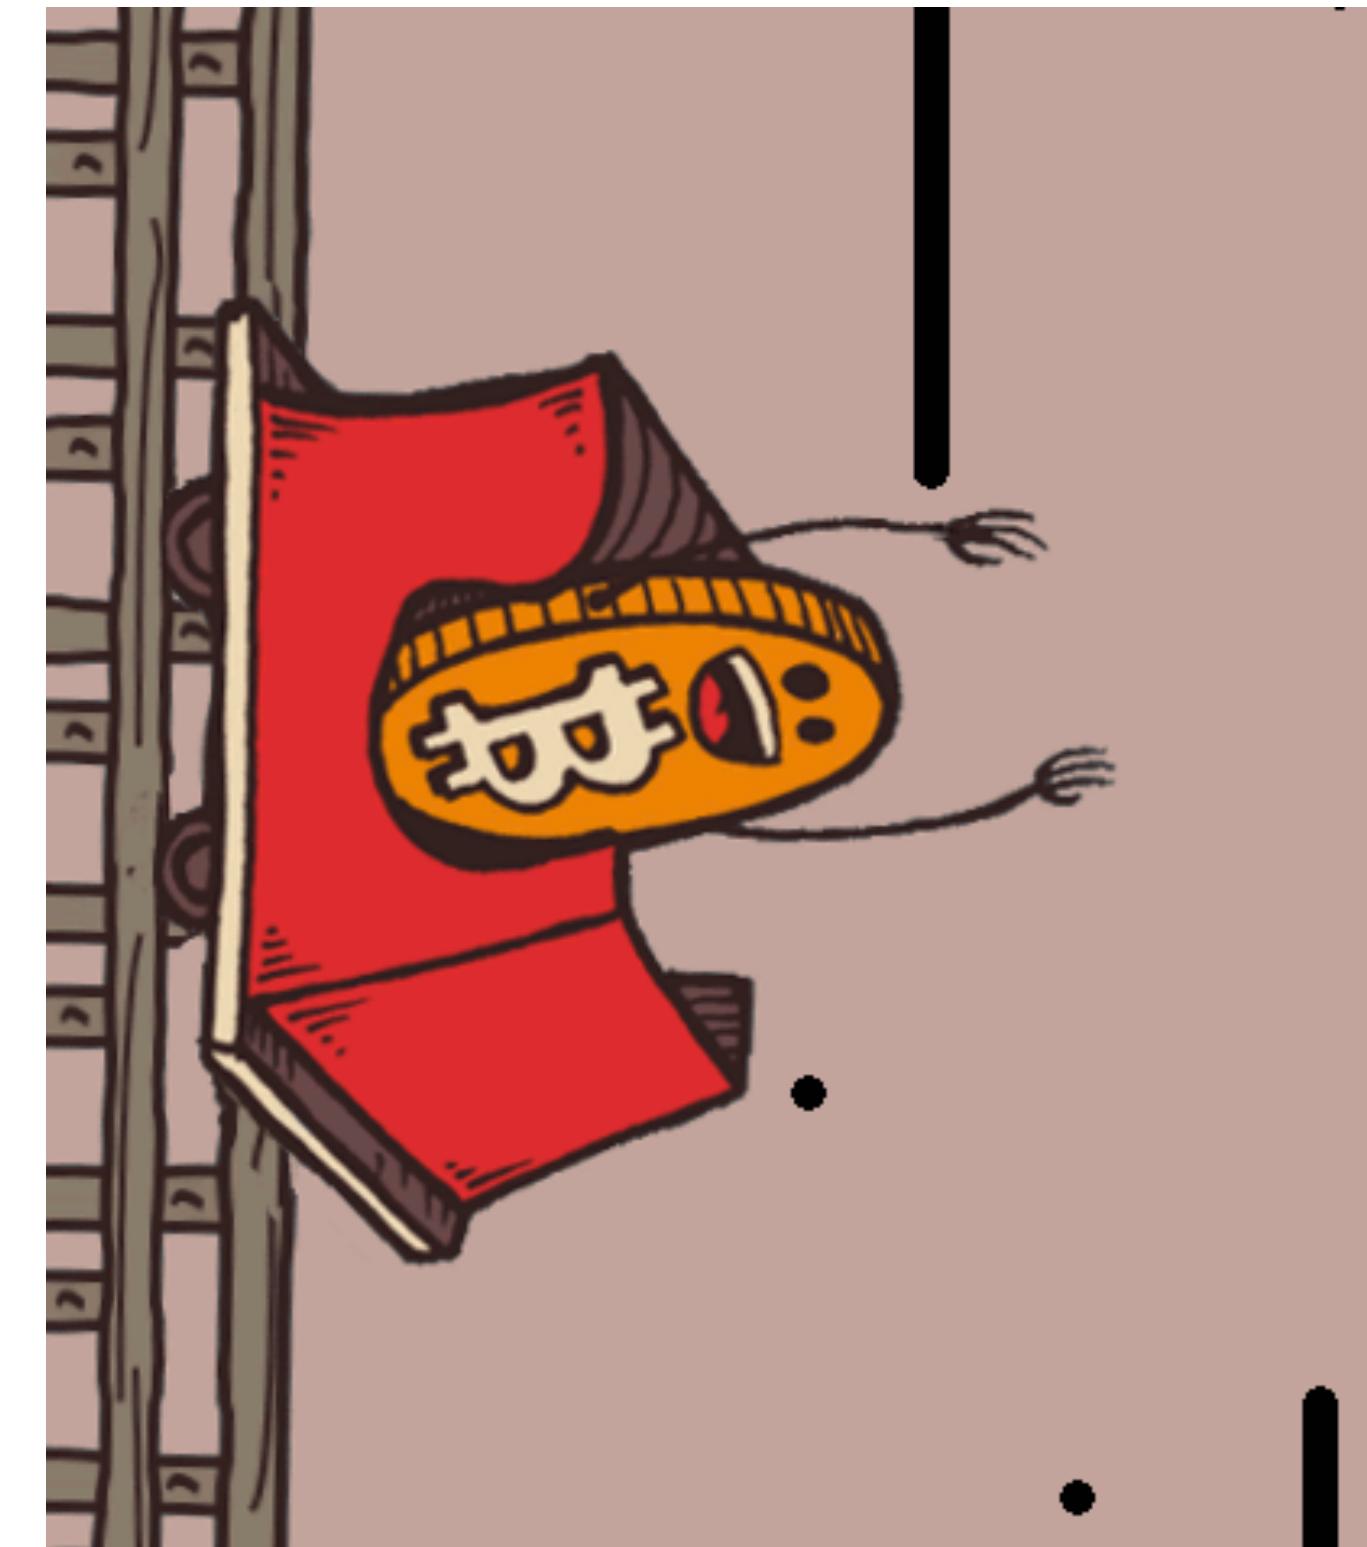
# The internet usage increases Utreeexo nodes need to download extra data

- 500GB of proof data needs to be downloaded (as of height 841,642). Block data is 613GB
- Currently working on caching to reduce this

# Bandwidth CPU



# disk i/o RAM usage

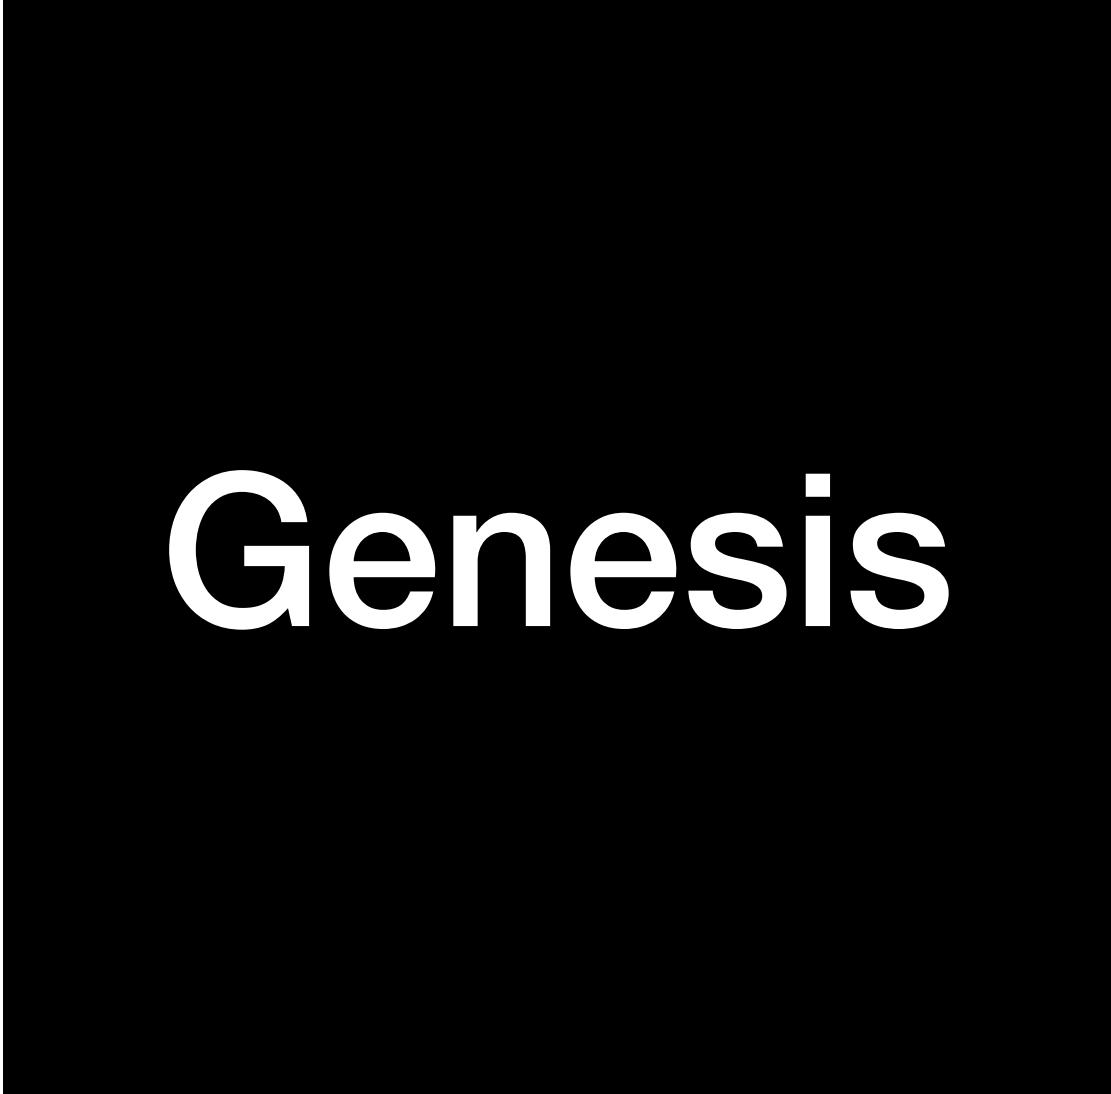


# Utilizing Utreetexo for faster block verification

# **Parallel block validation**

# **Current block validation**

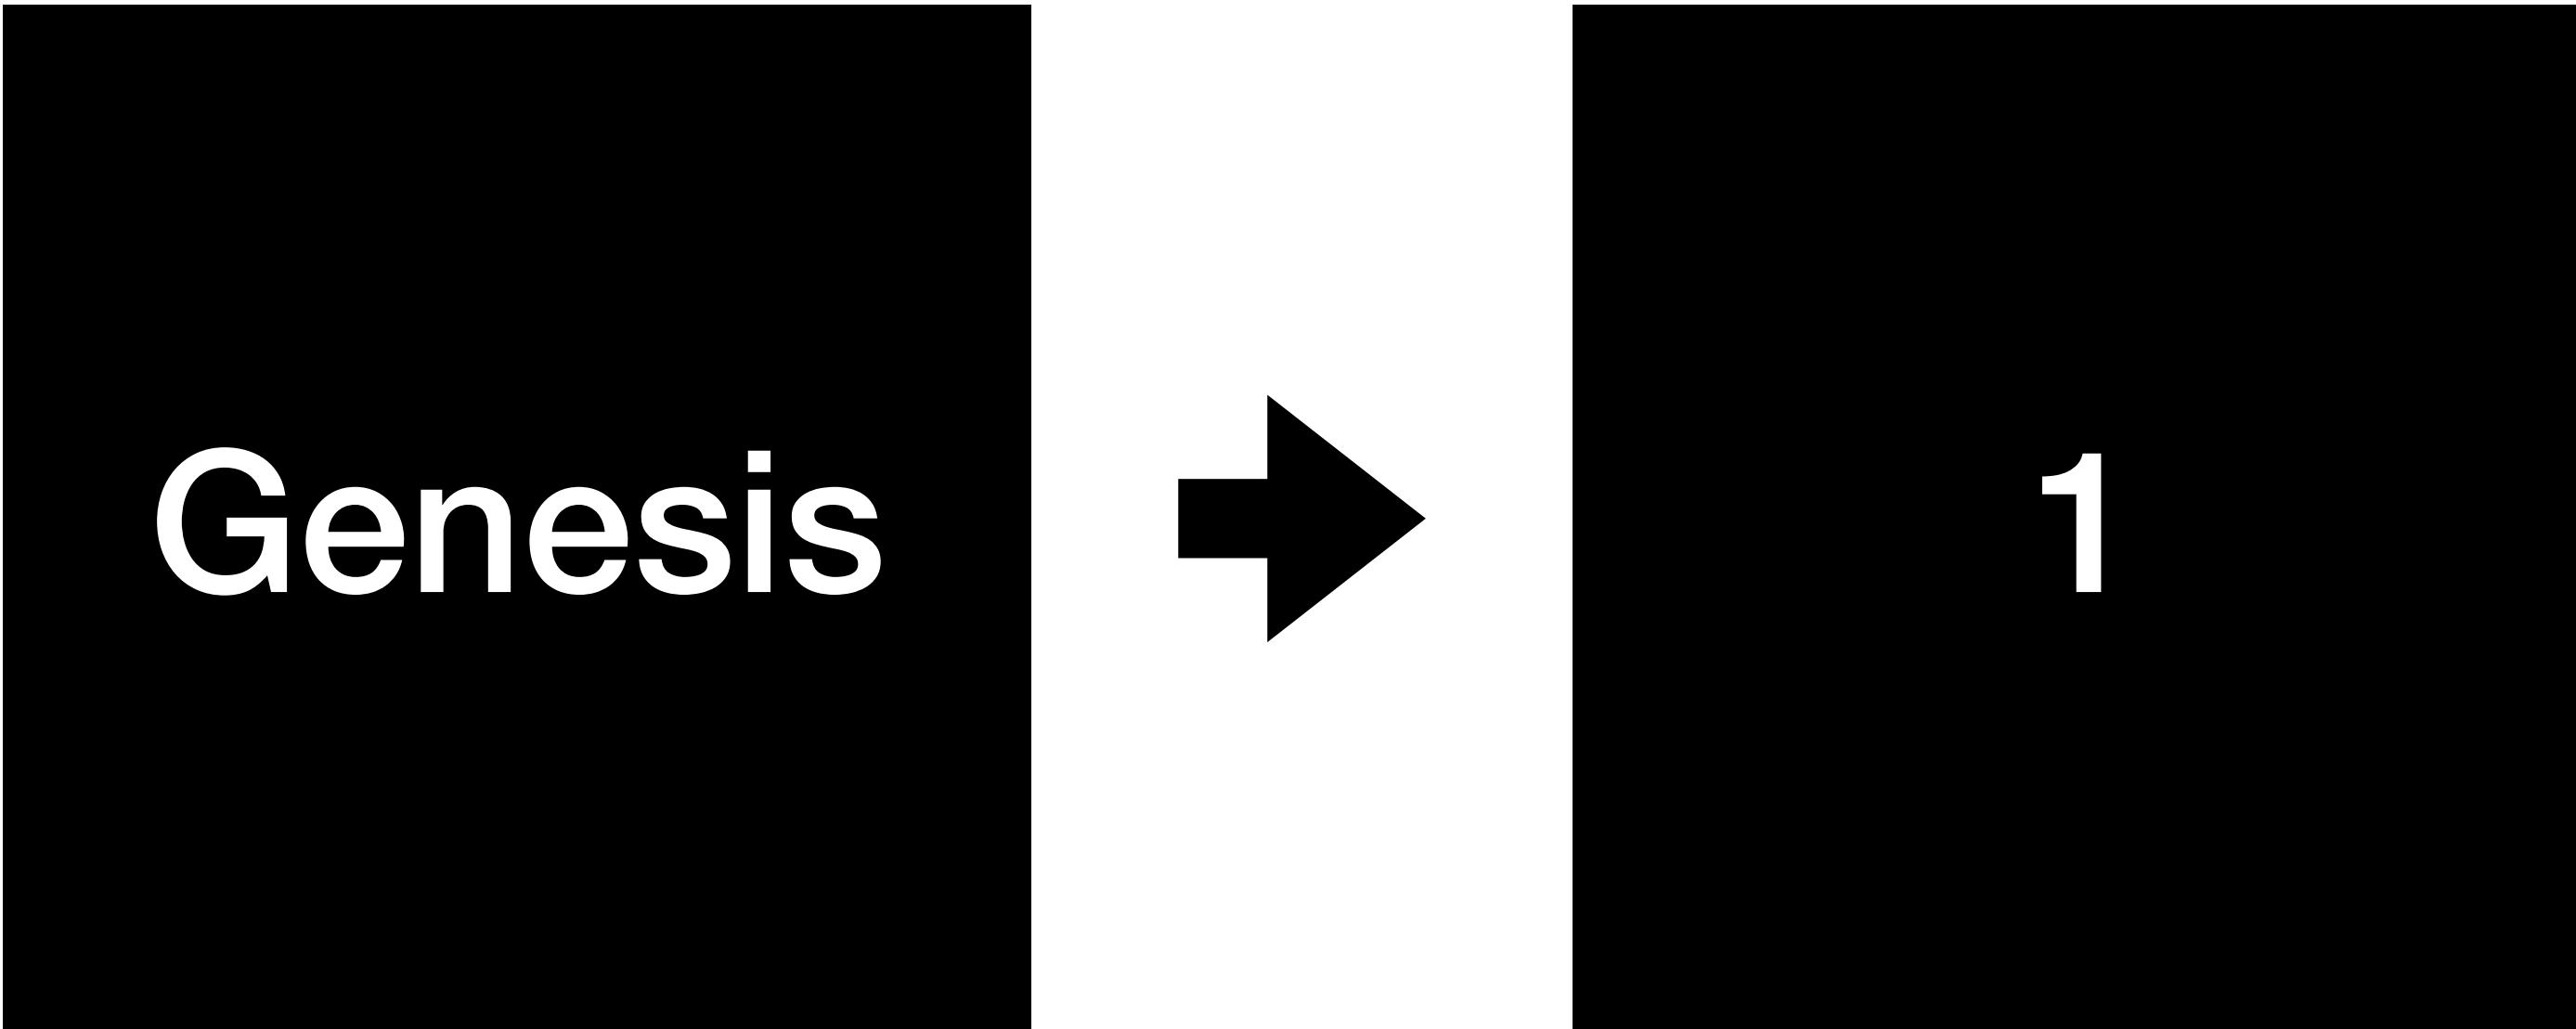
## **Sequential validation**



**Genesis**

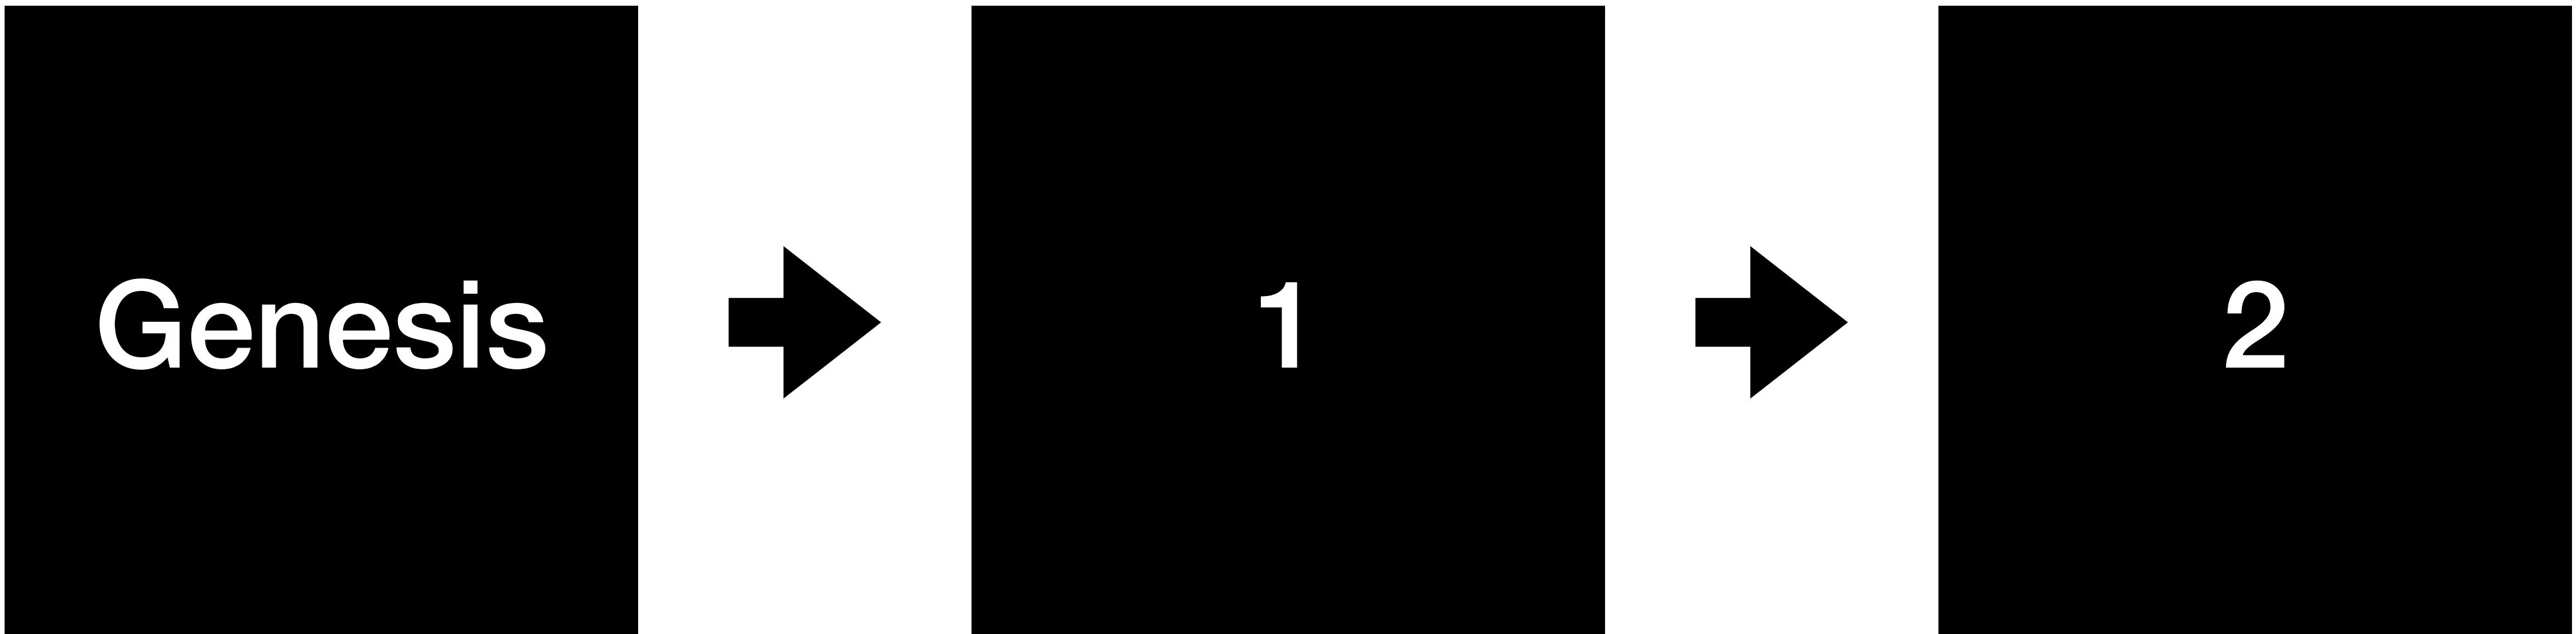
# **Current block validation**

**Sequential validation**



# Current block validation

## Sequential validation



# Parallel block validation

Replay blocks in any order

- Utreexo roots committed into the binary
- Process blocks starting from any of the roots that are committed
- <https://github.com/mit-dci/utcd/blob/master/chaincfg/mainnetroots.go>

# With Utreexo

## Efficient parallel validation

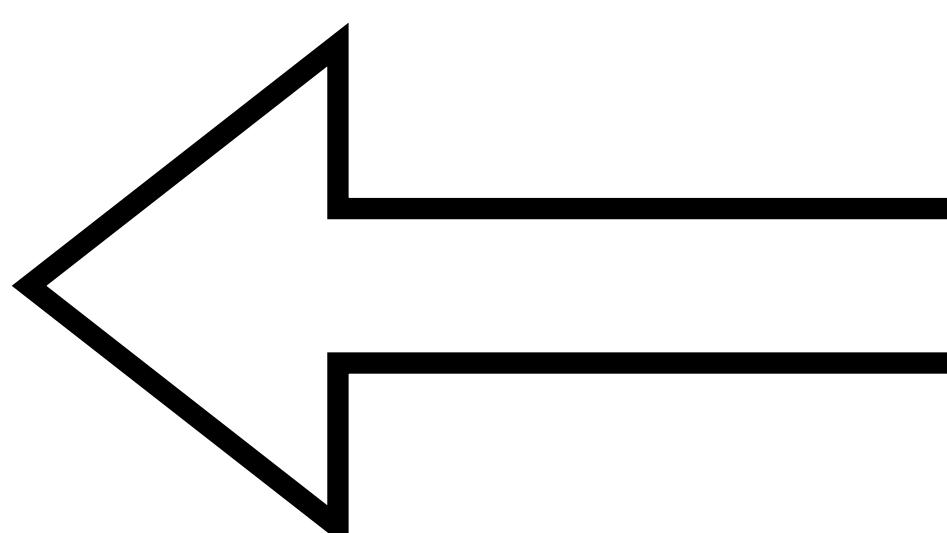
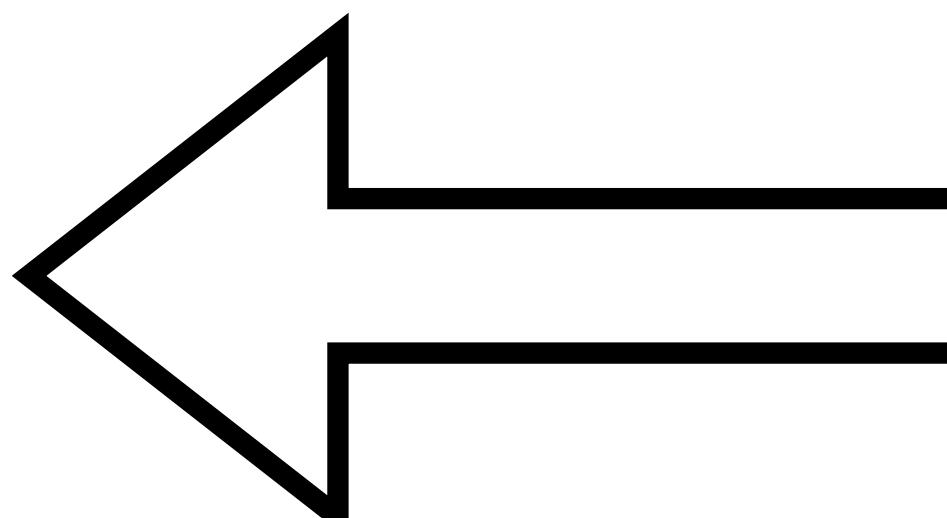
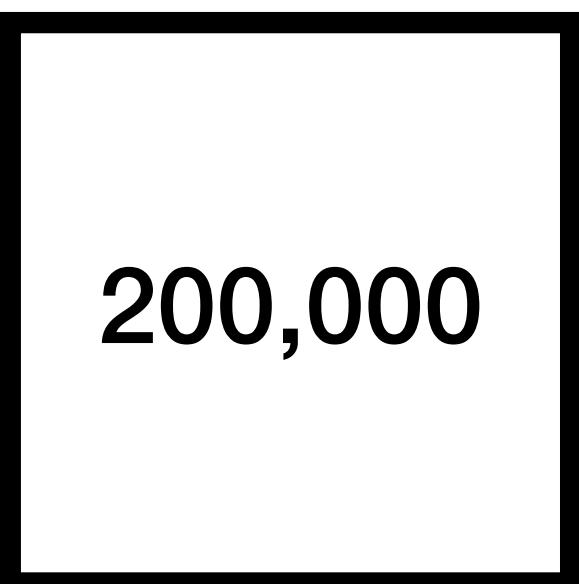
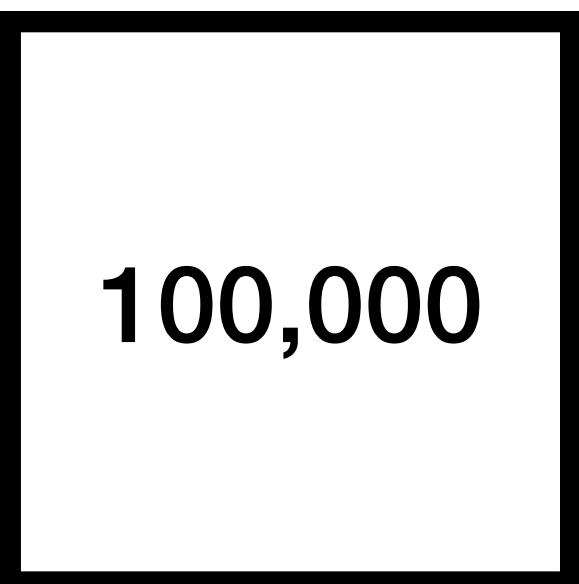
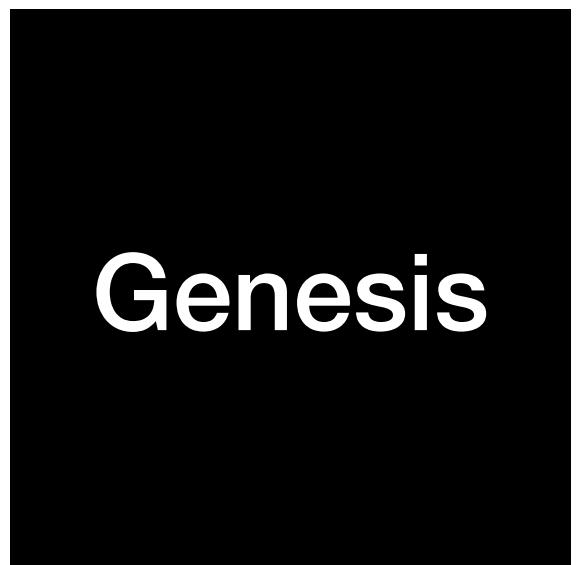
Genesis

100,000

200,000

# With Utreexo

Efficient parallel validation



Untrusted roots at  
the start

# With Utreexo

## Efficient parallel validation

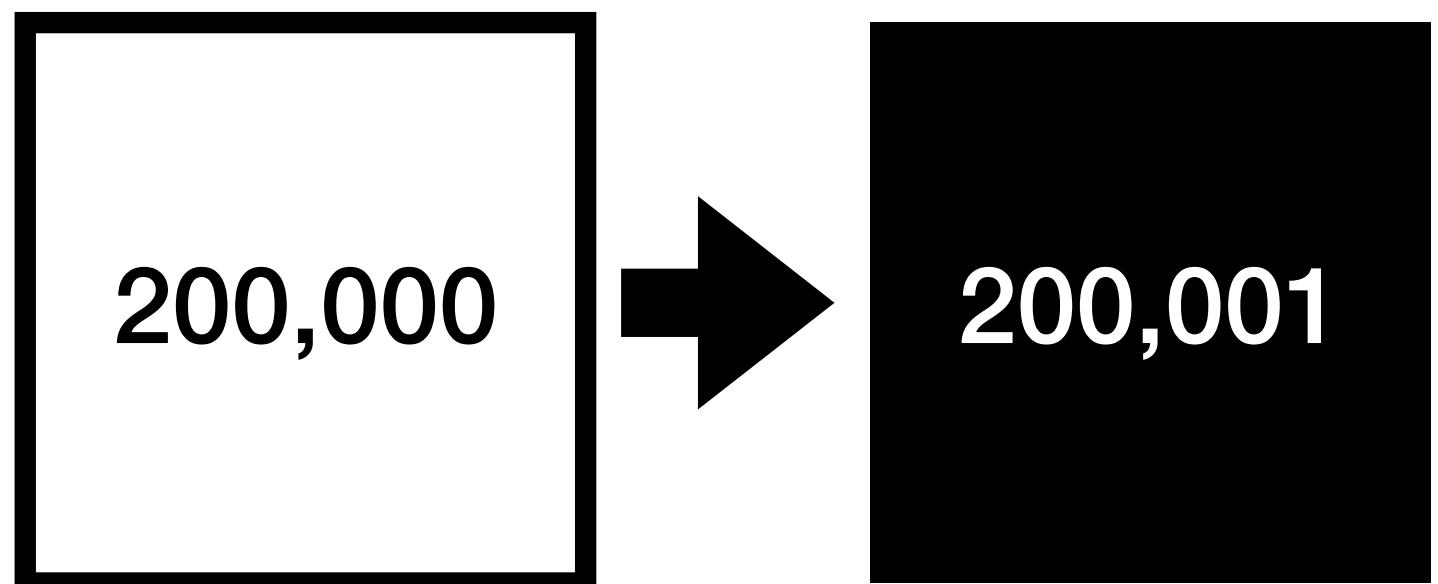
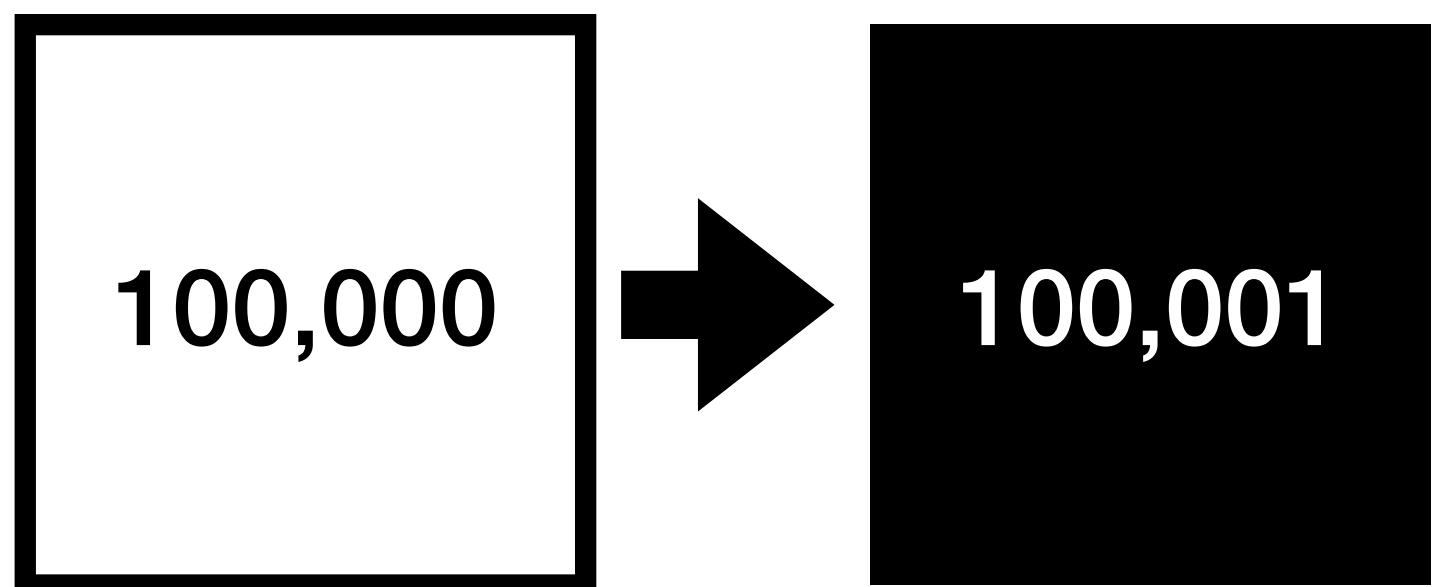
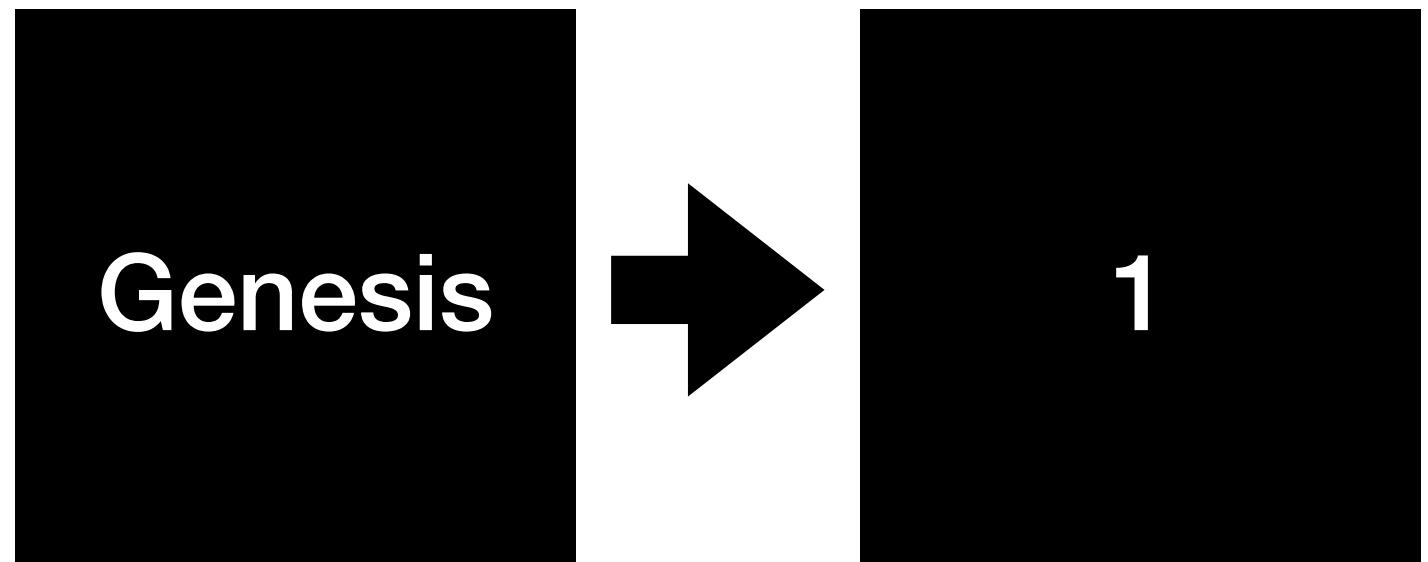
Genesis

100,000

200,000

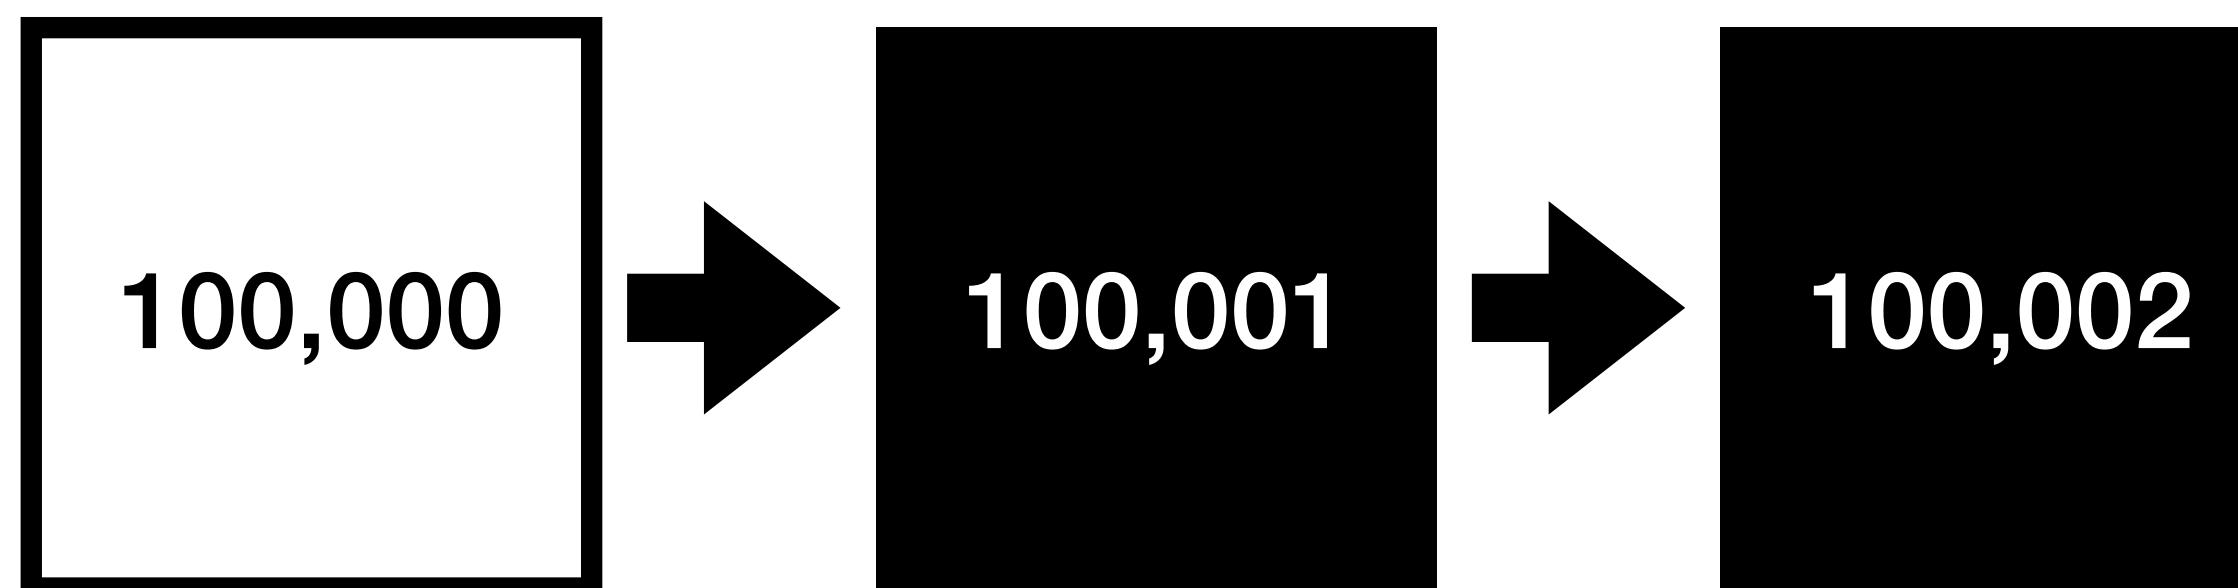
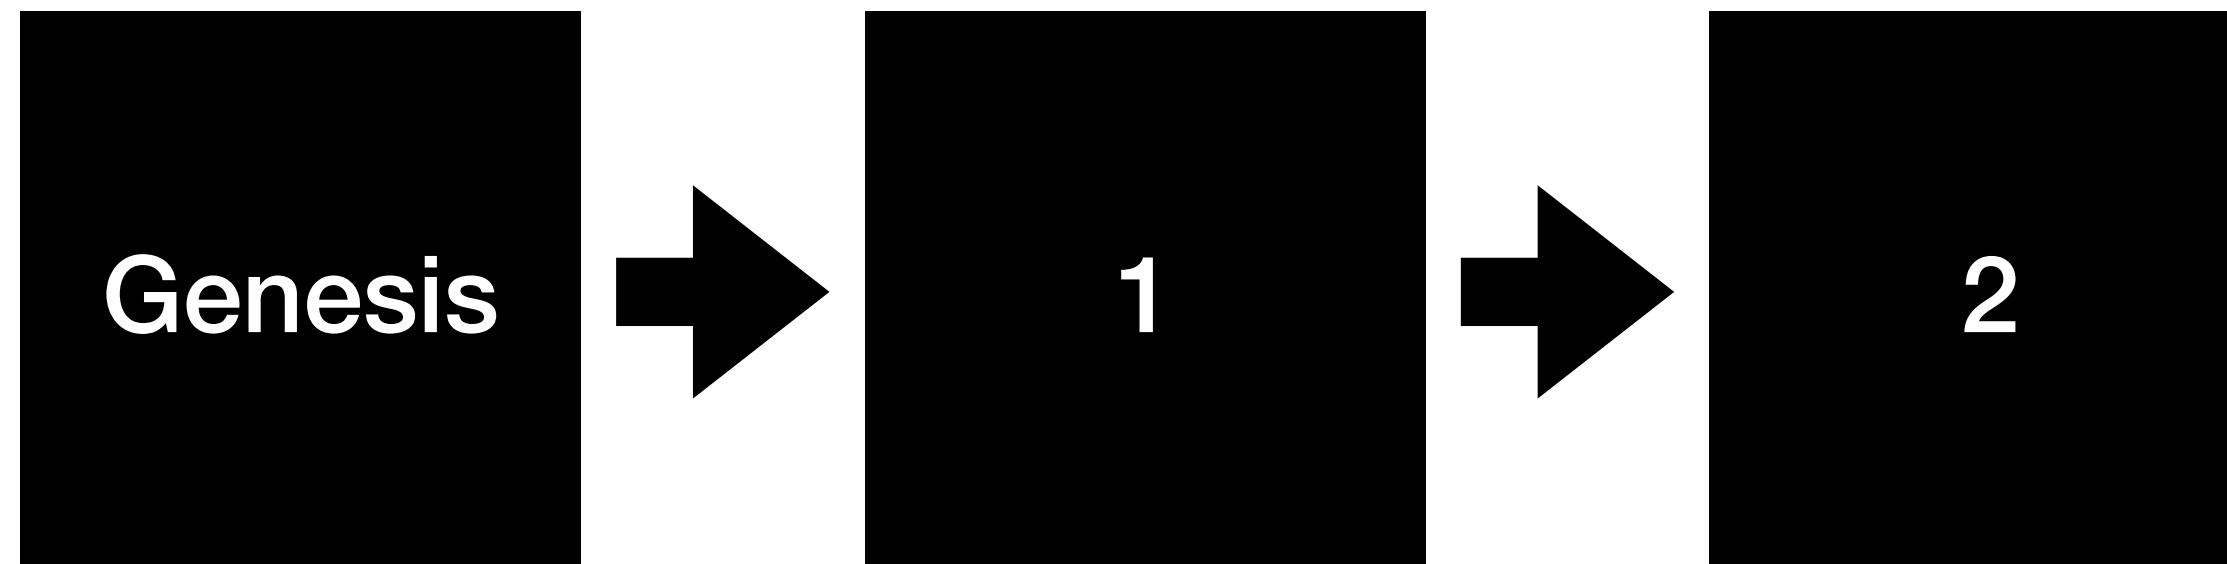
# With Utreexo

## Efficient parallel validation



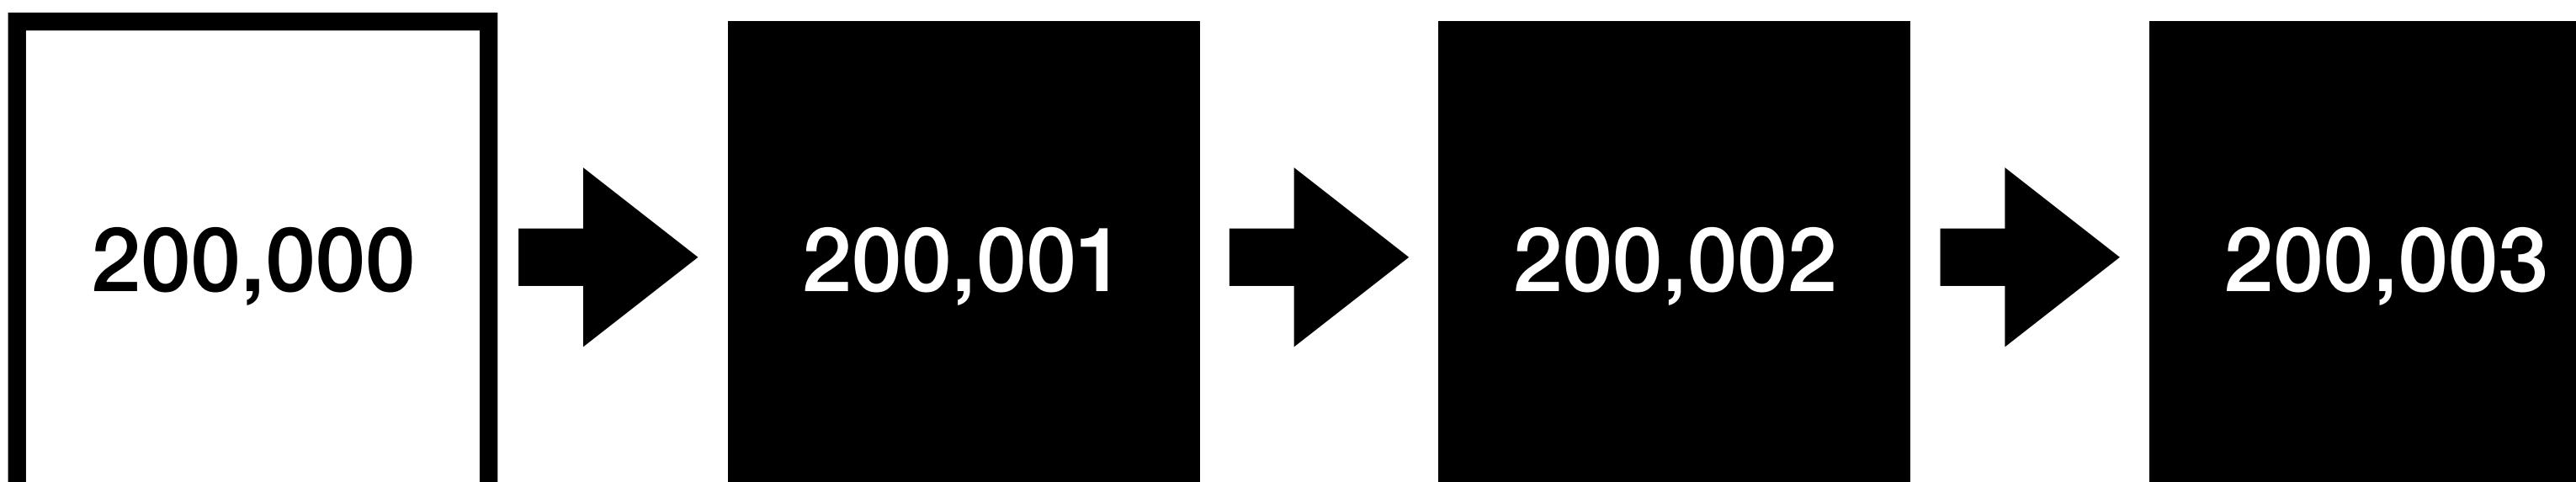
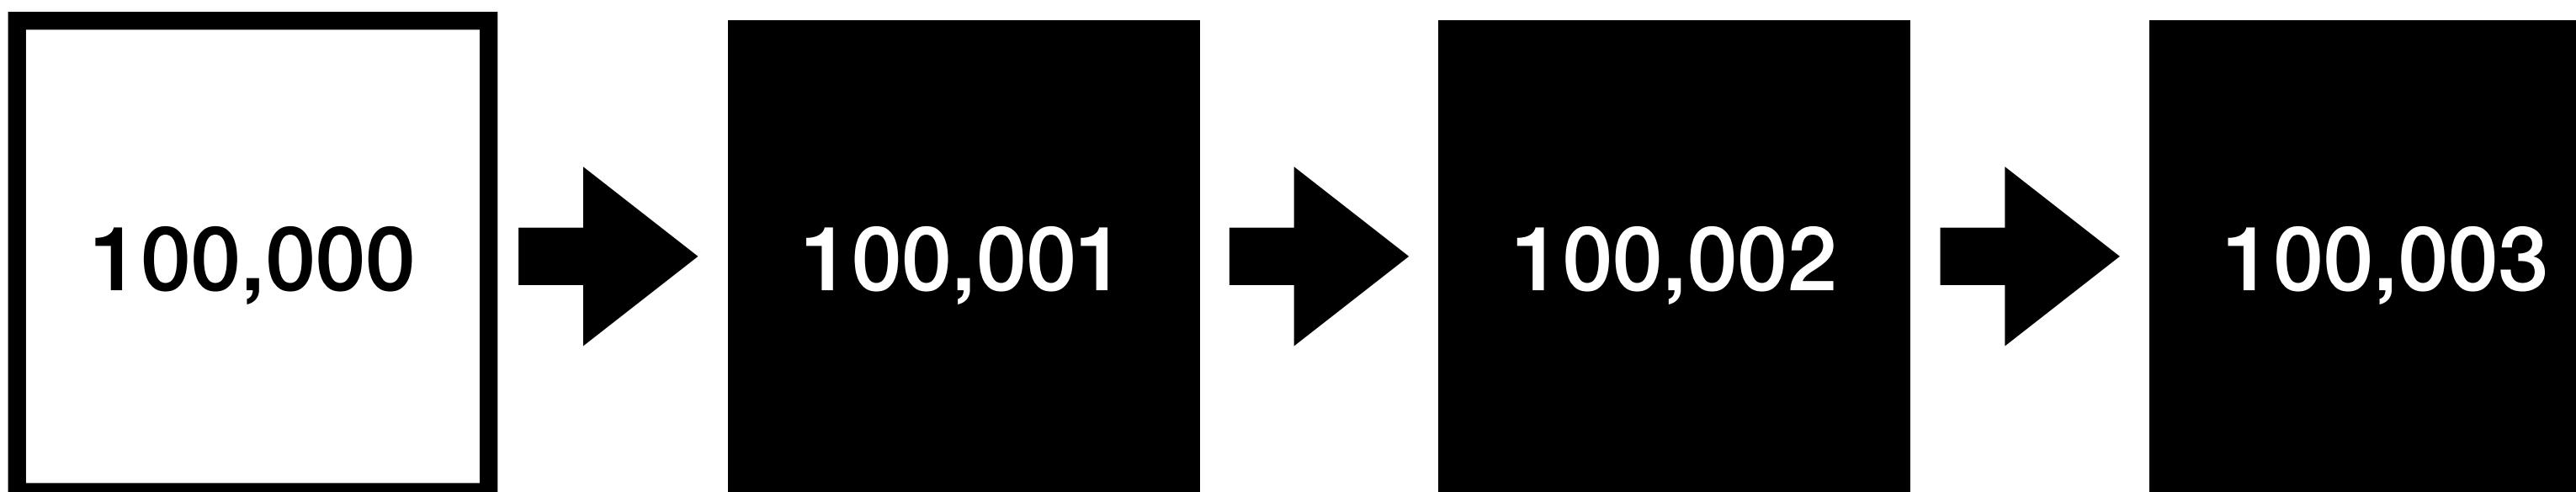
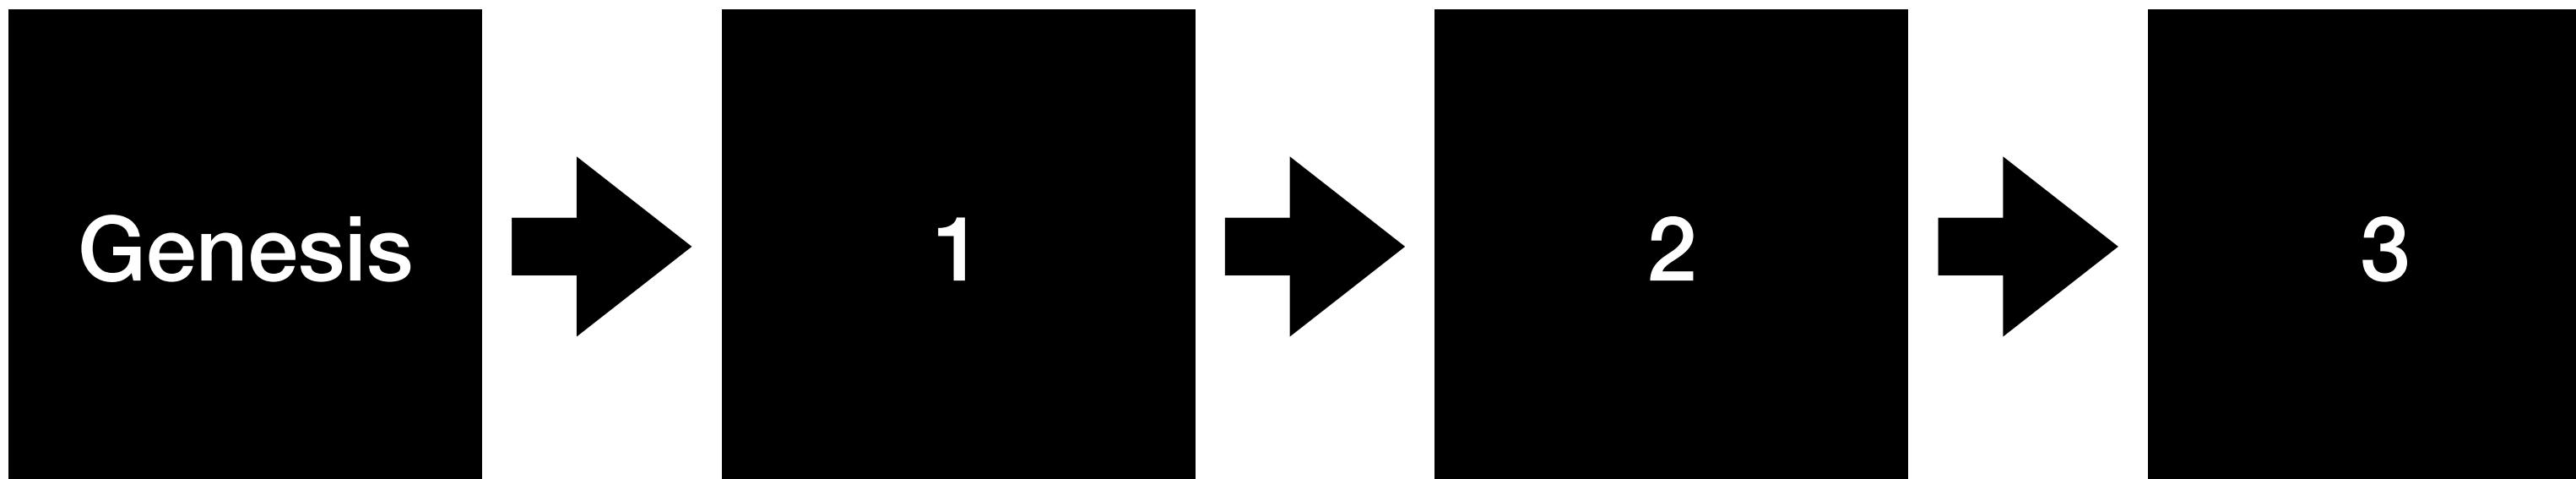
# With Utreexo

## Efficient parallel validation



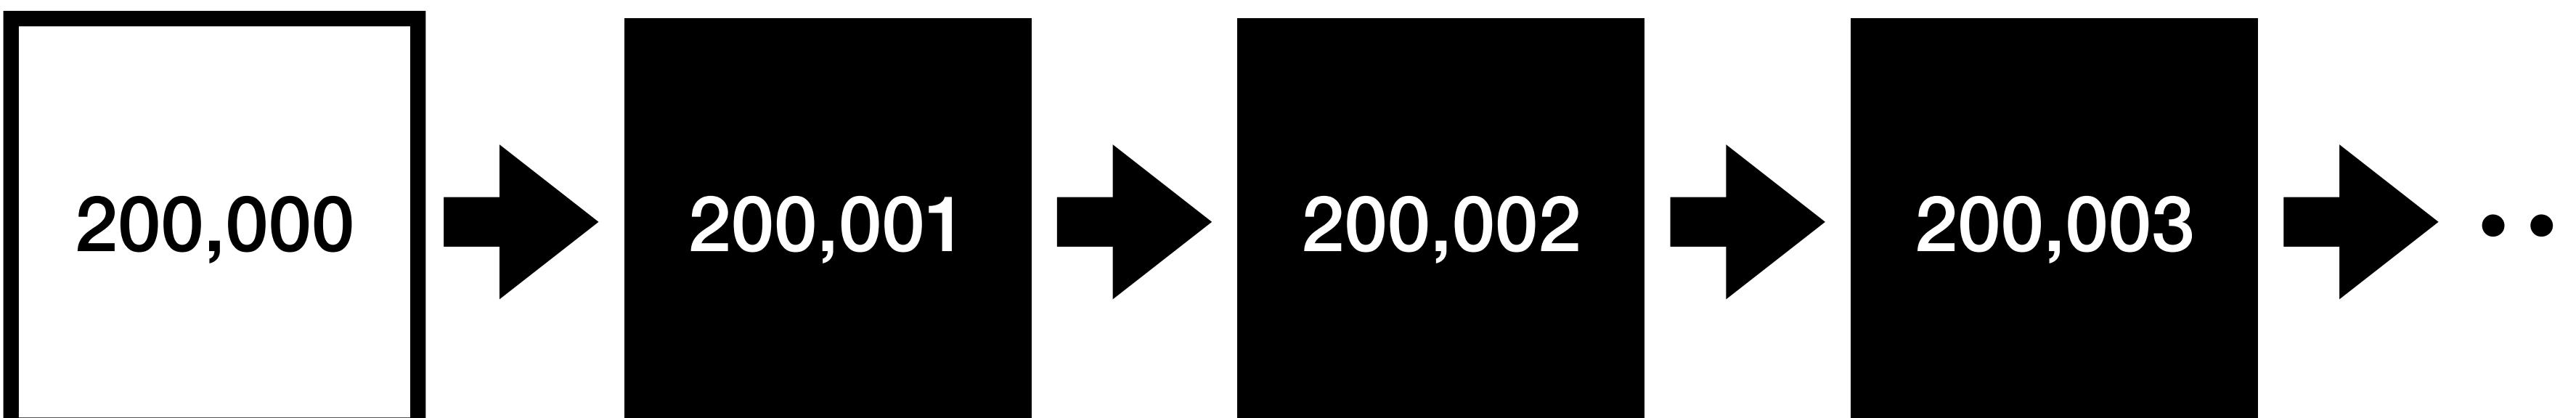
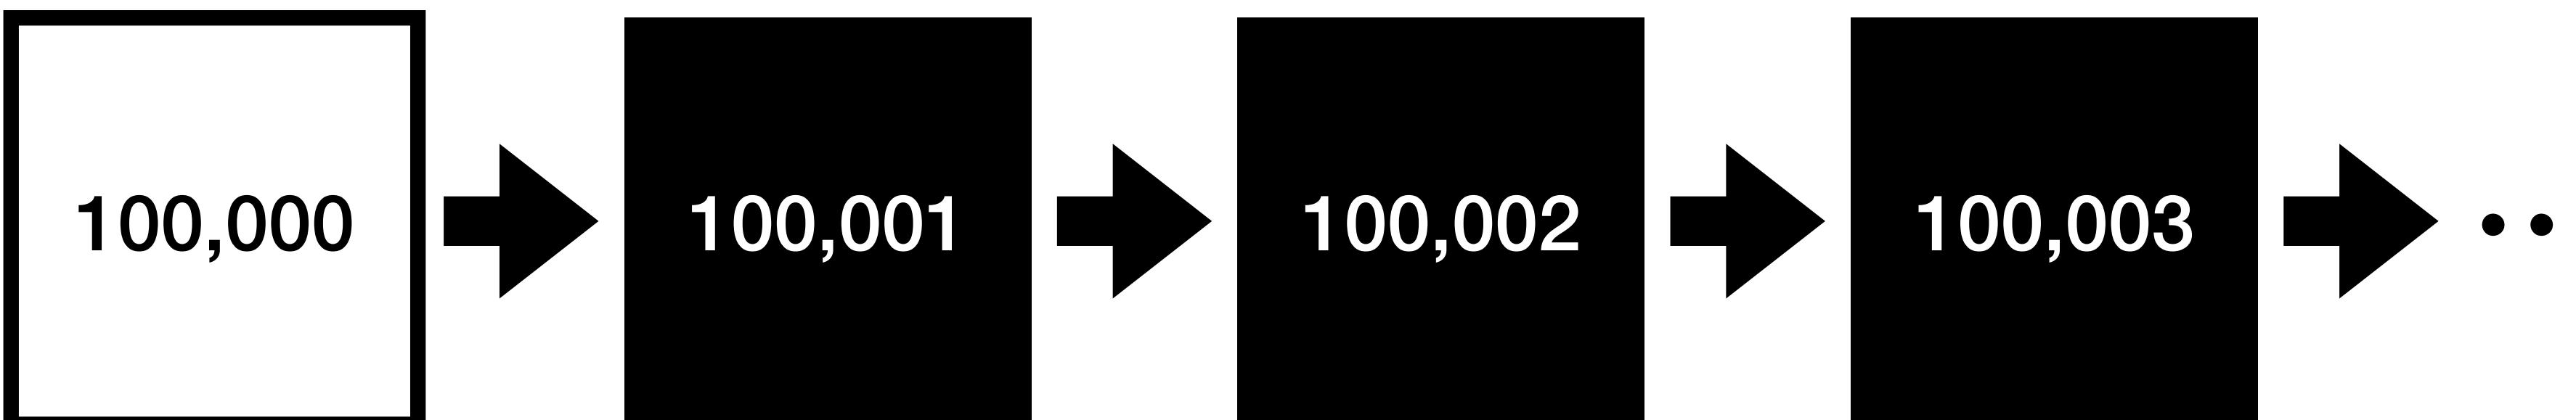
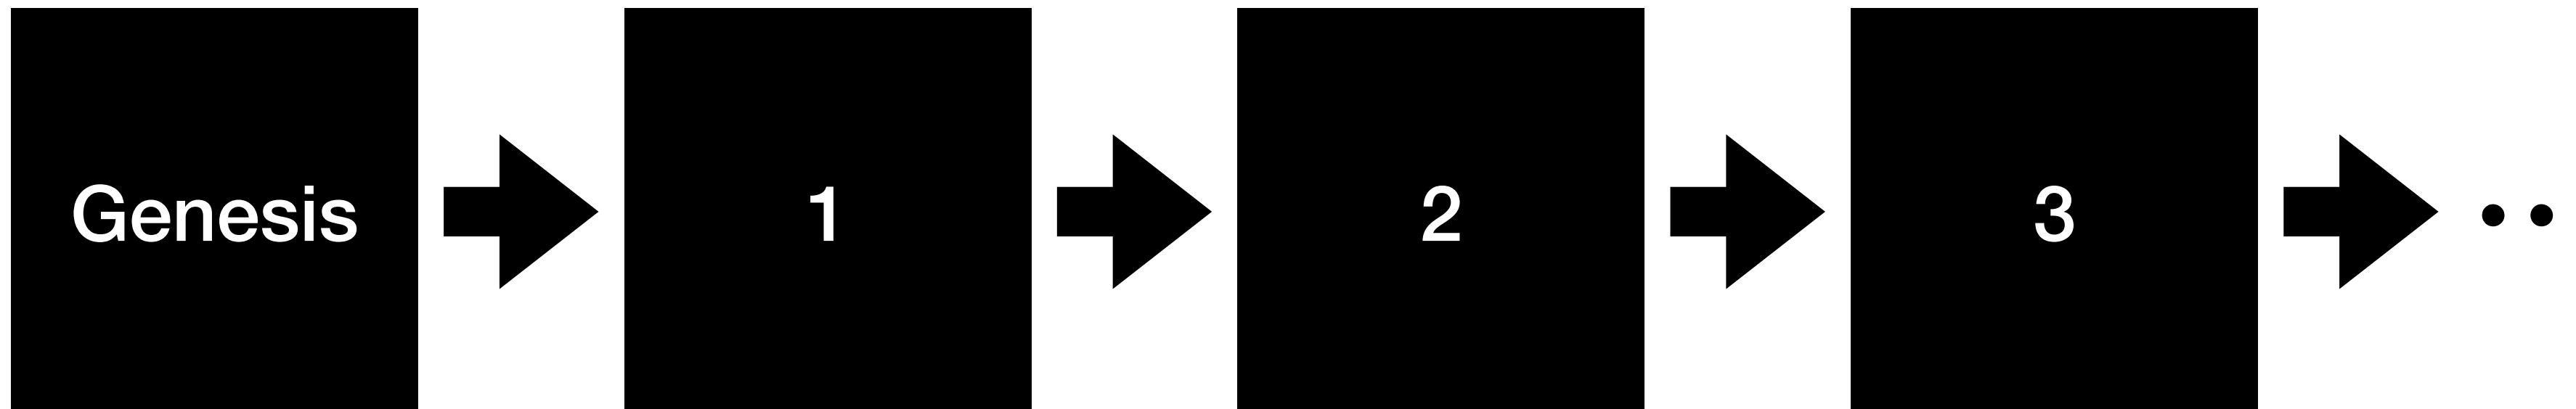
# With Utreexo

## Efficient parallel validation



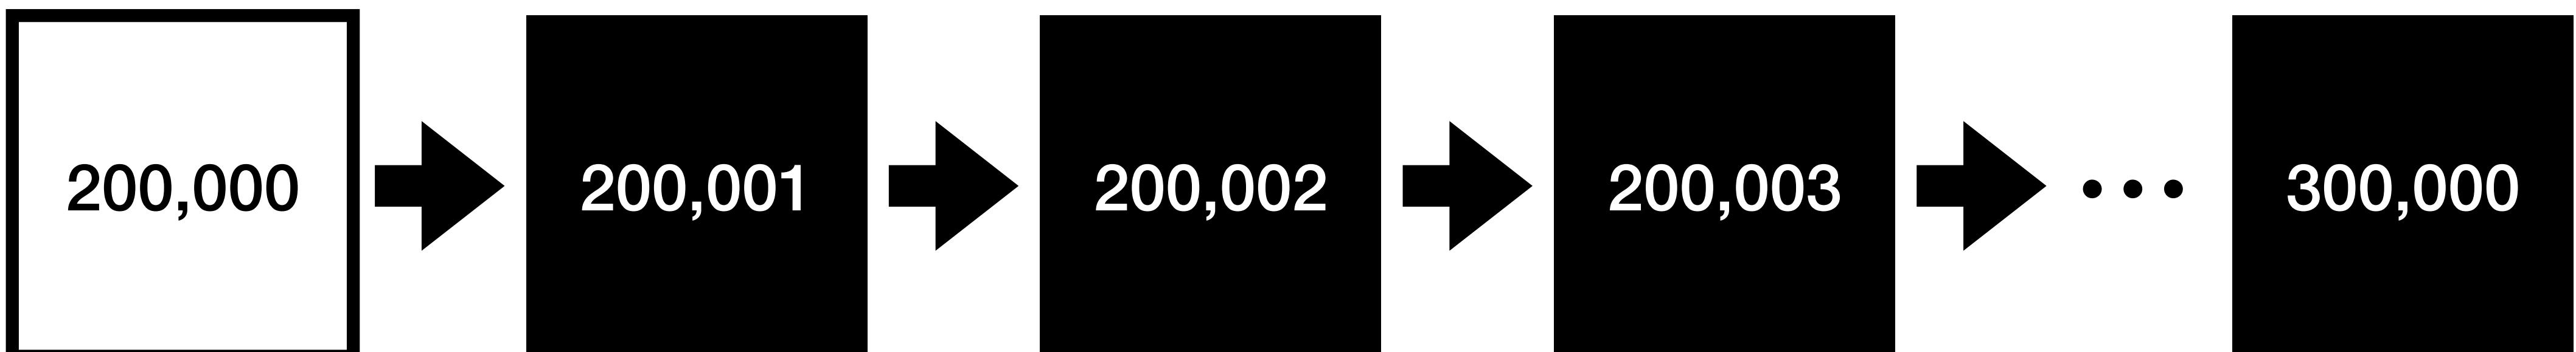
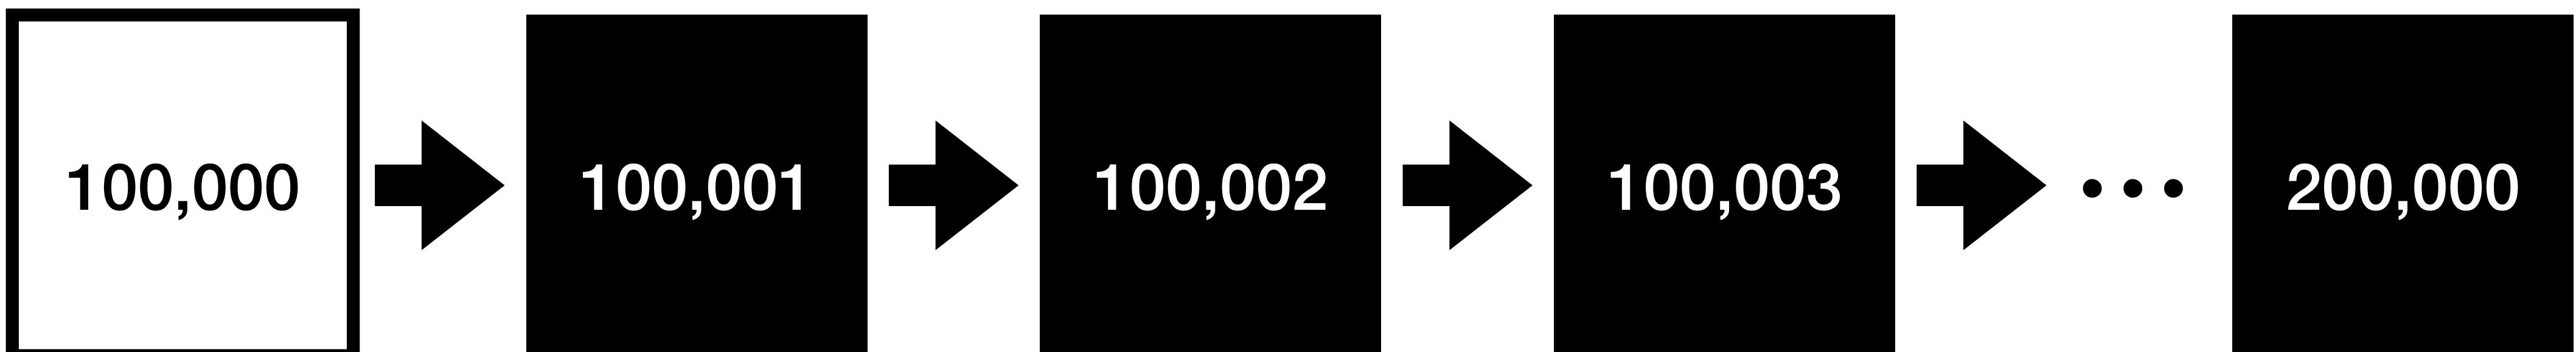
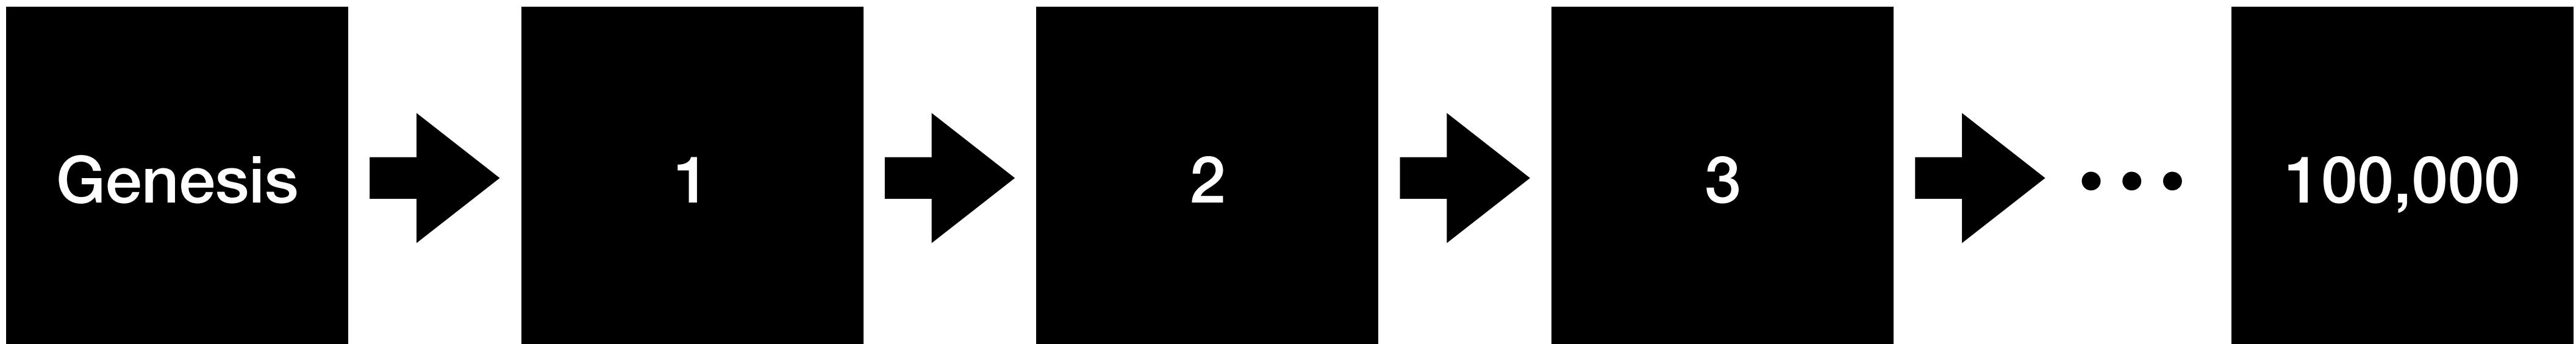
# With Utreexo

## Efficient parallel validation



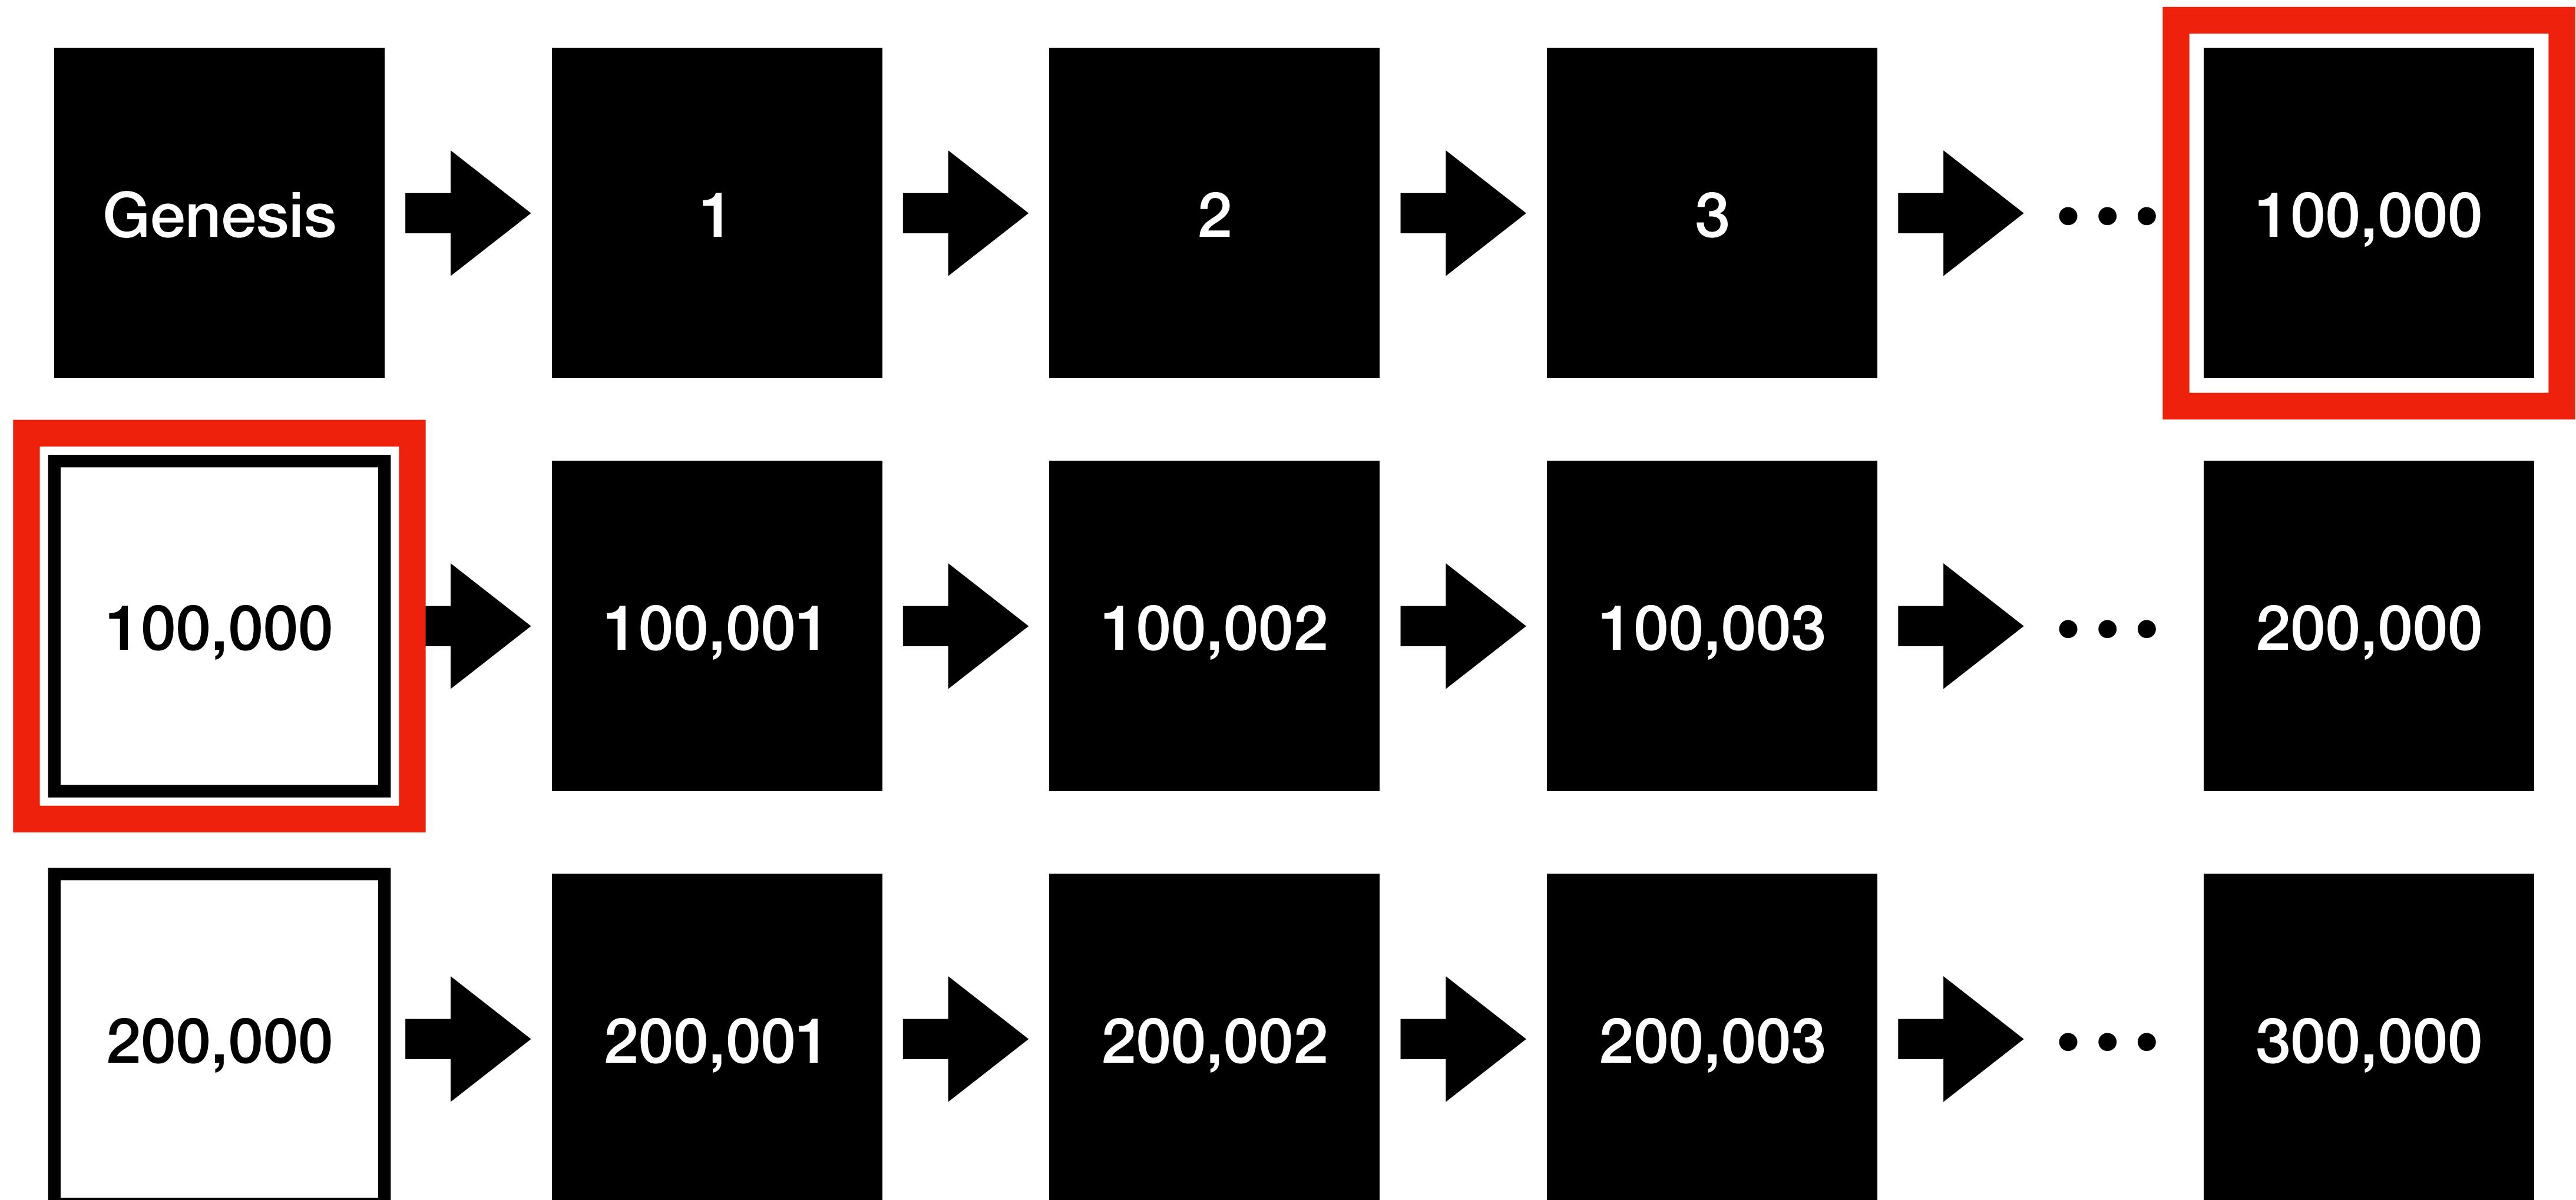
# With Utreexo

## Efficient parallel validation



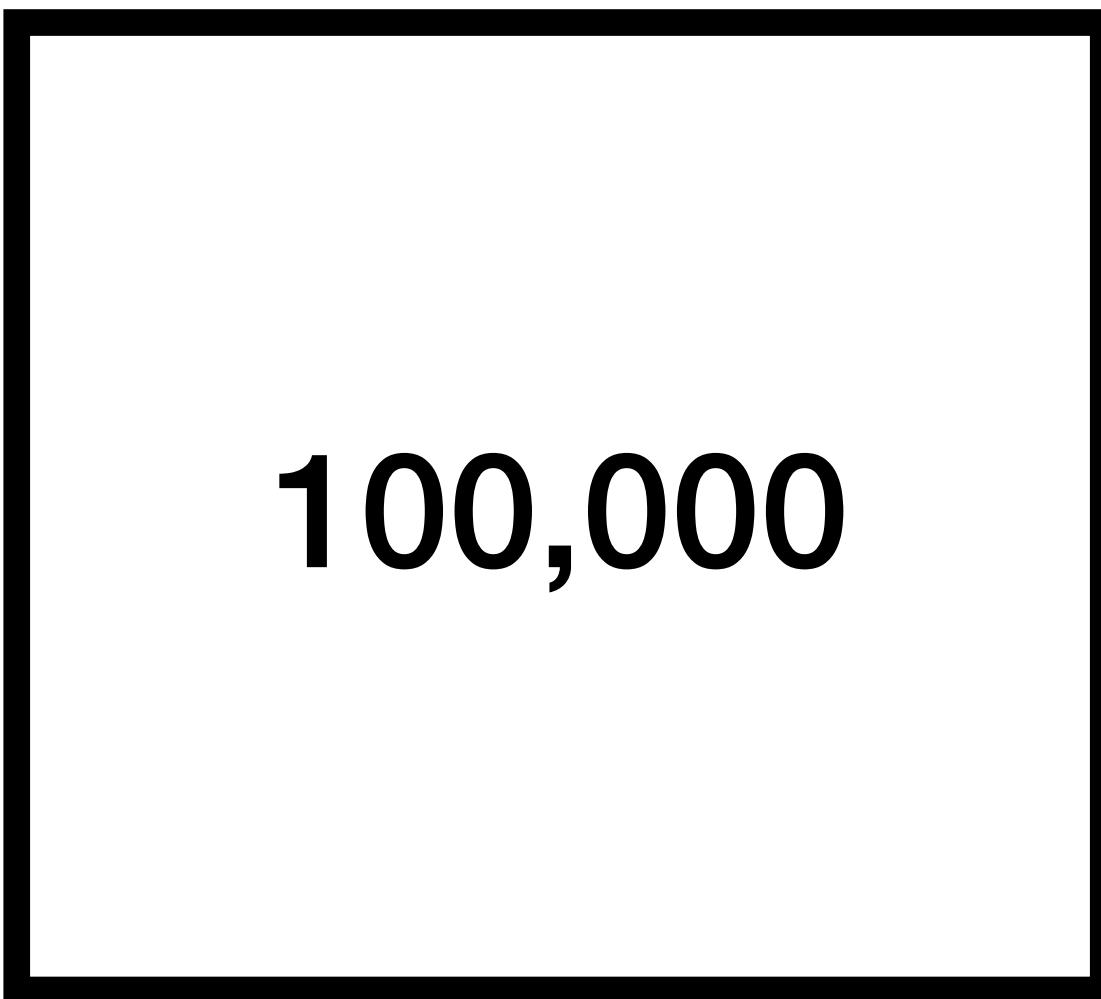
# With Utreexo

## Efficient parallel validation

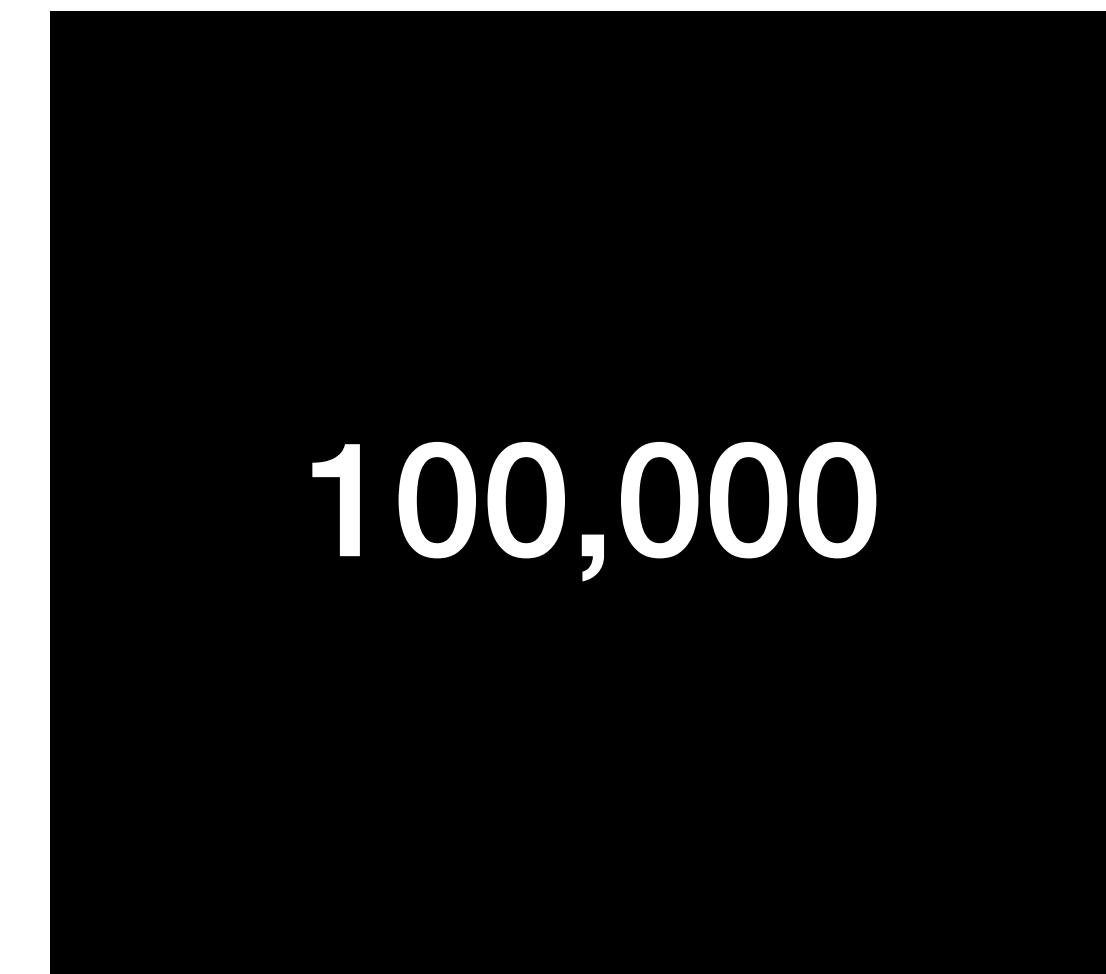


# Compare roots

Root at 100,000 becomes trusted if equal

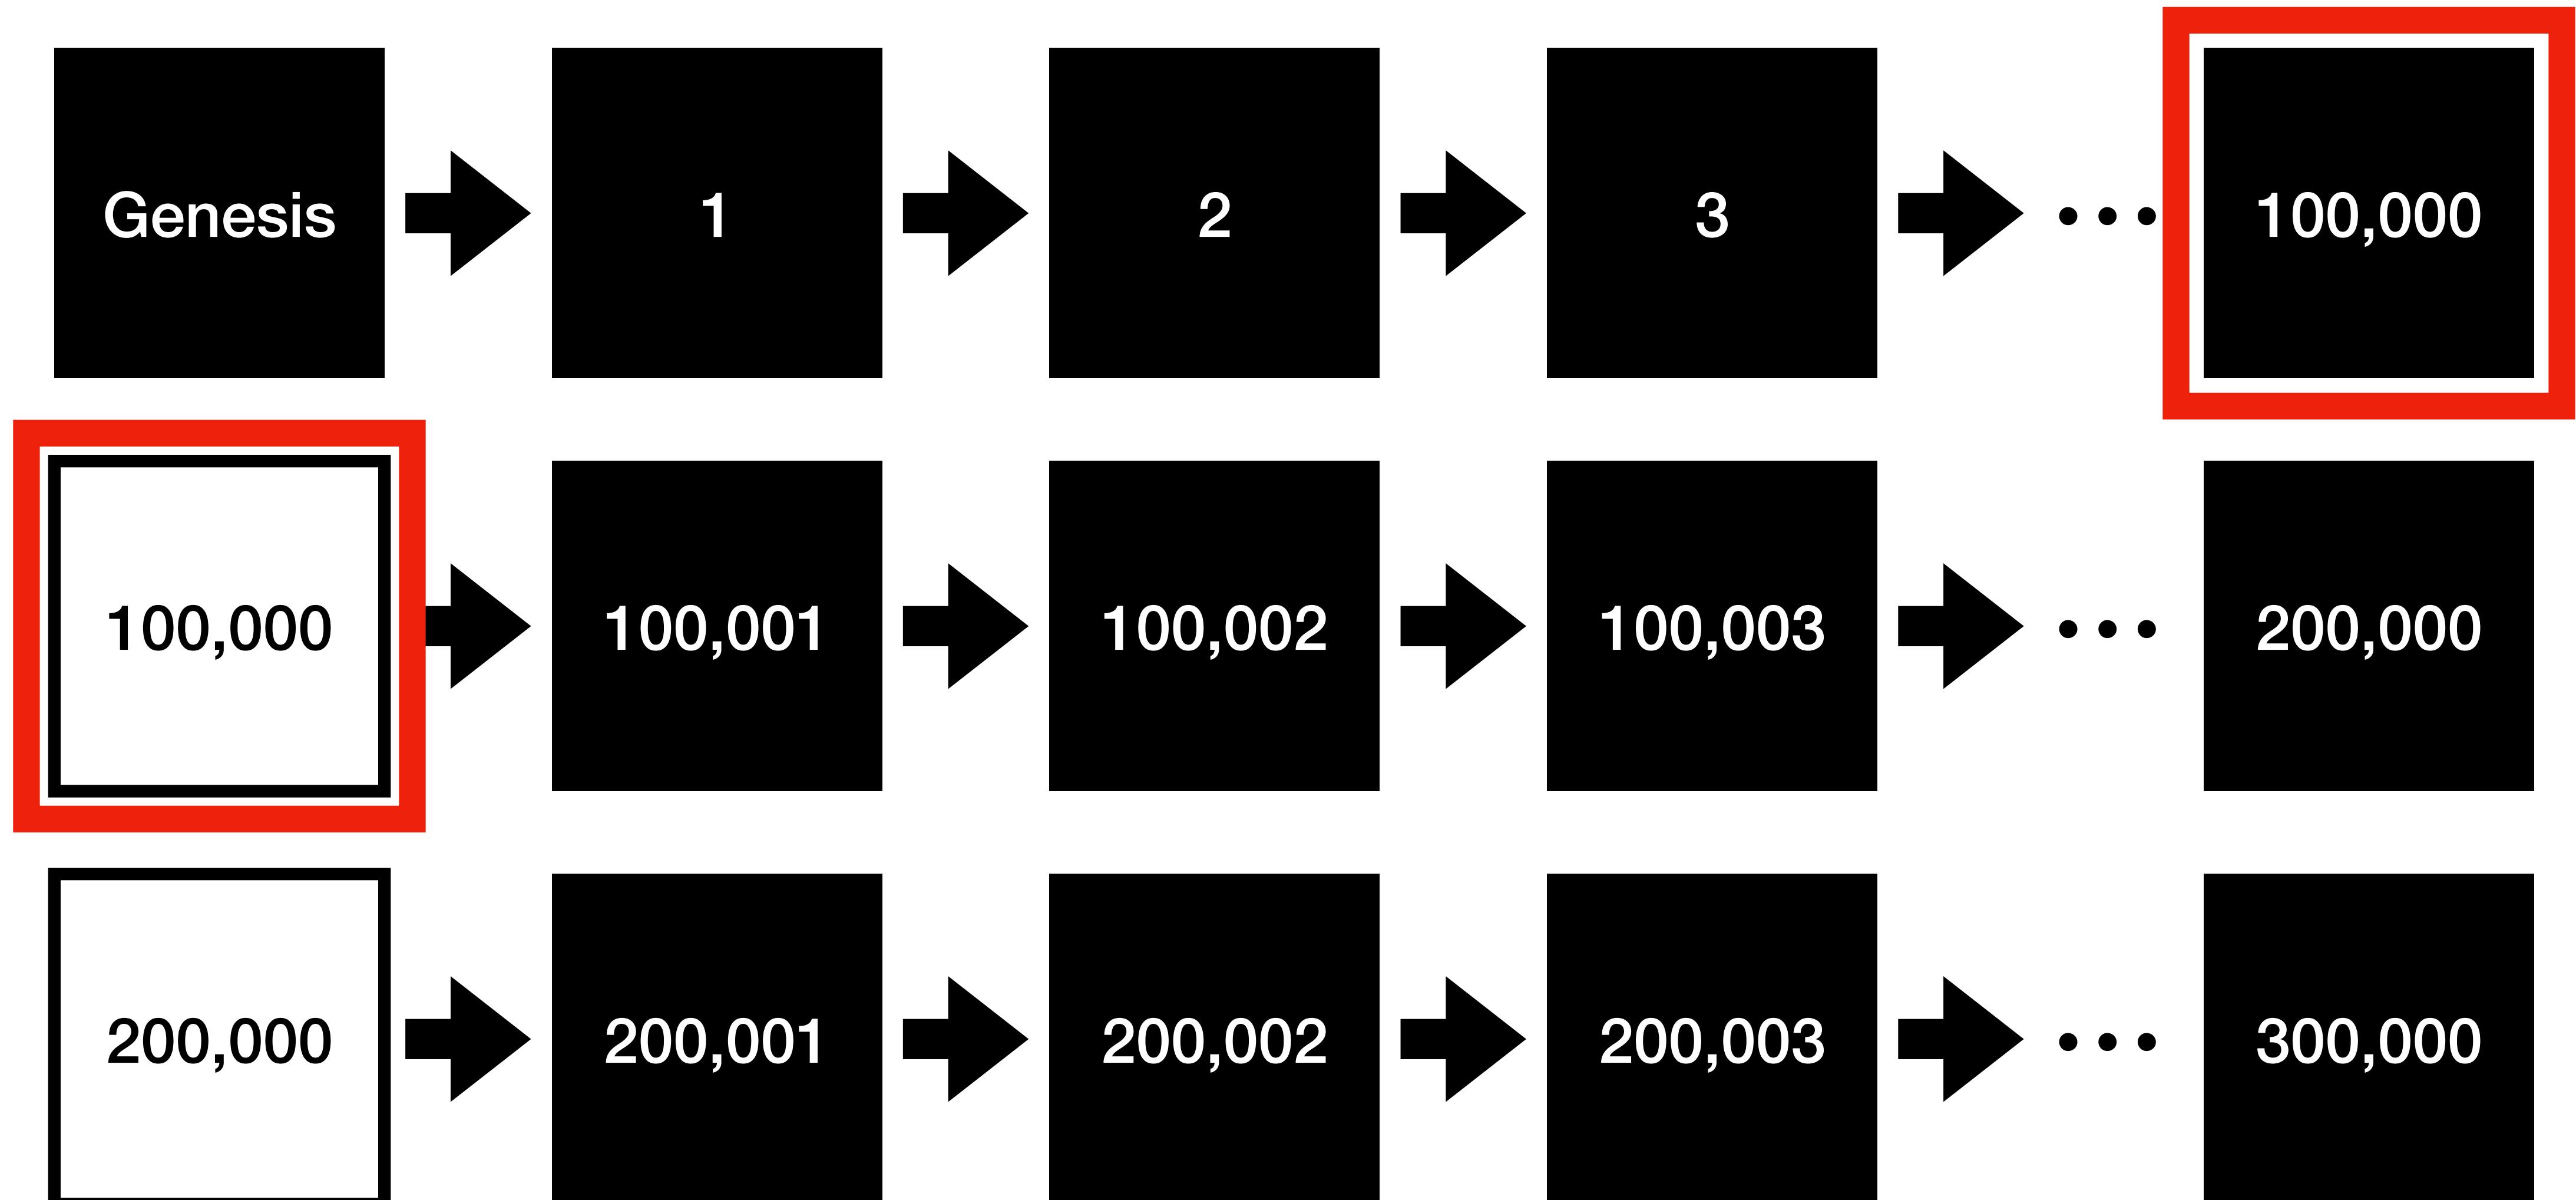


==



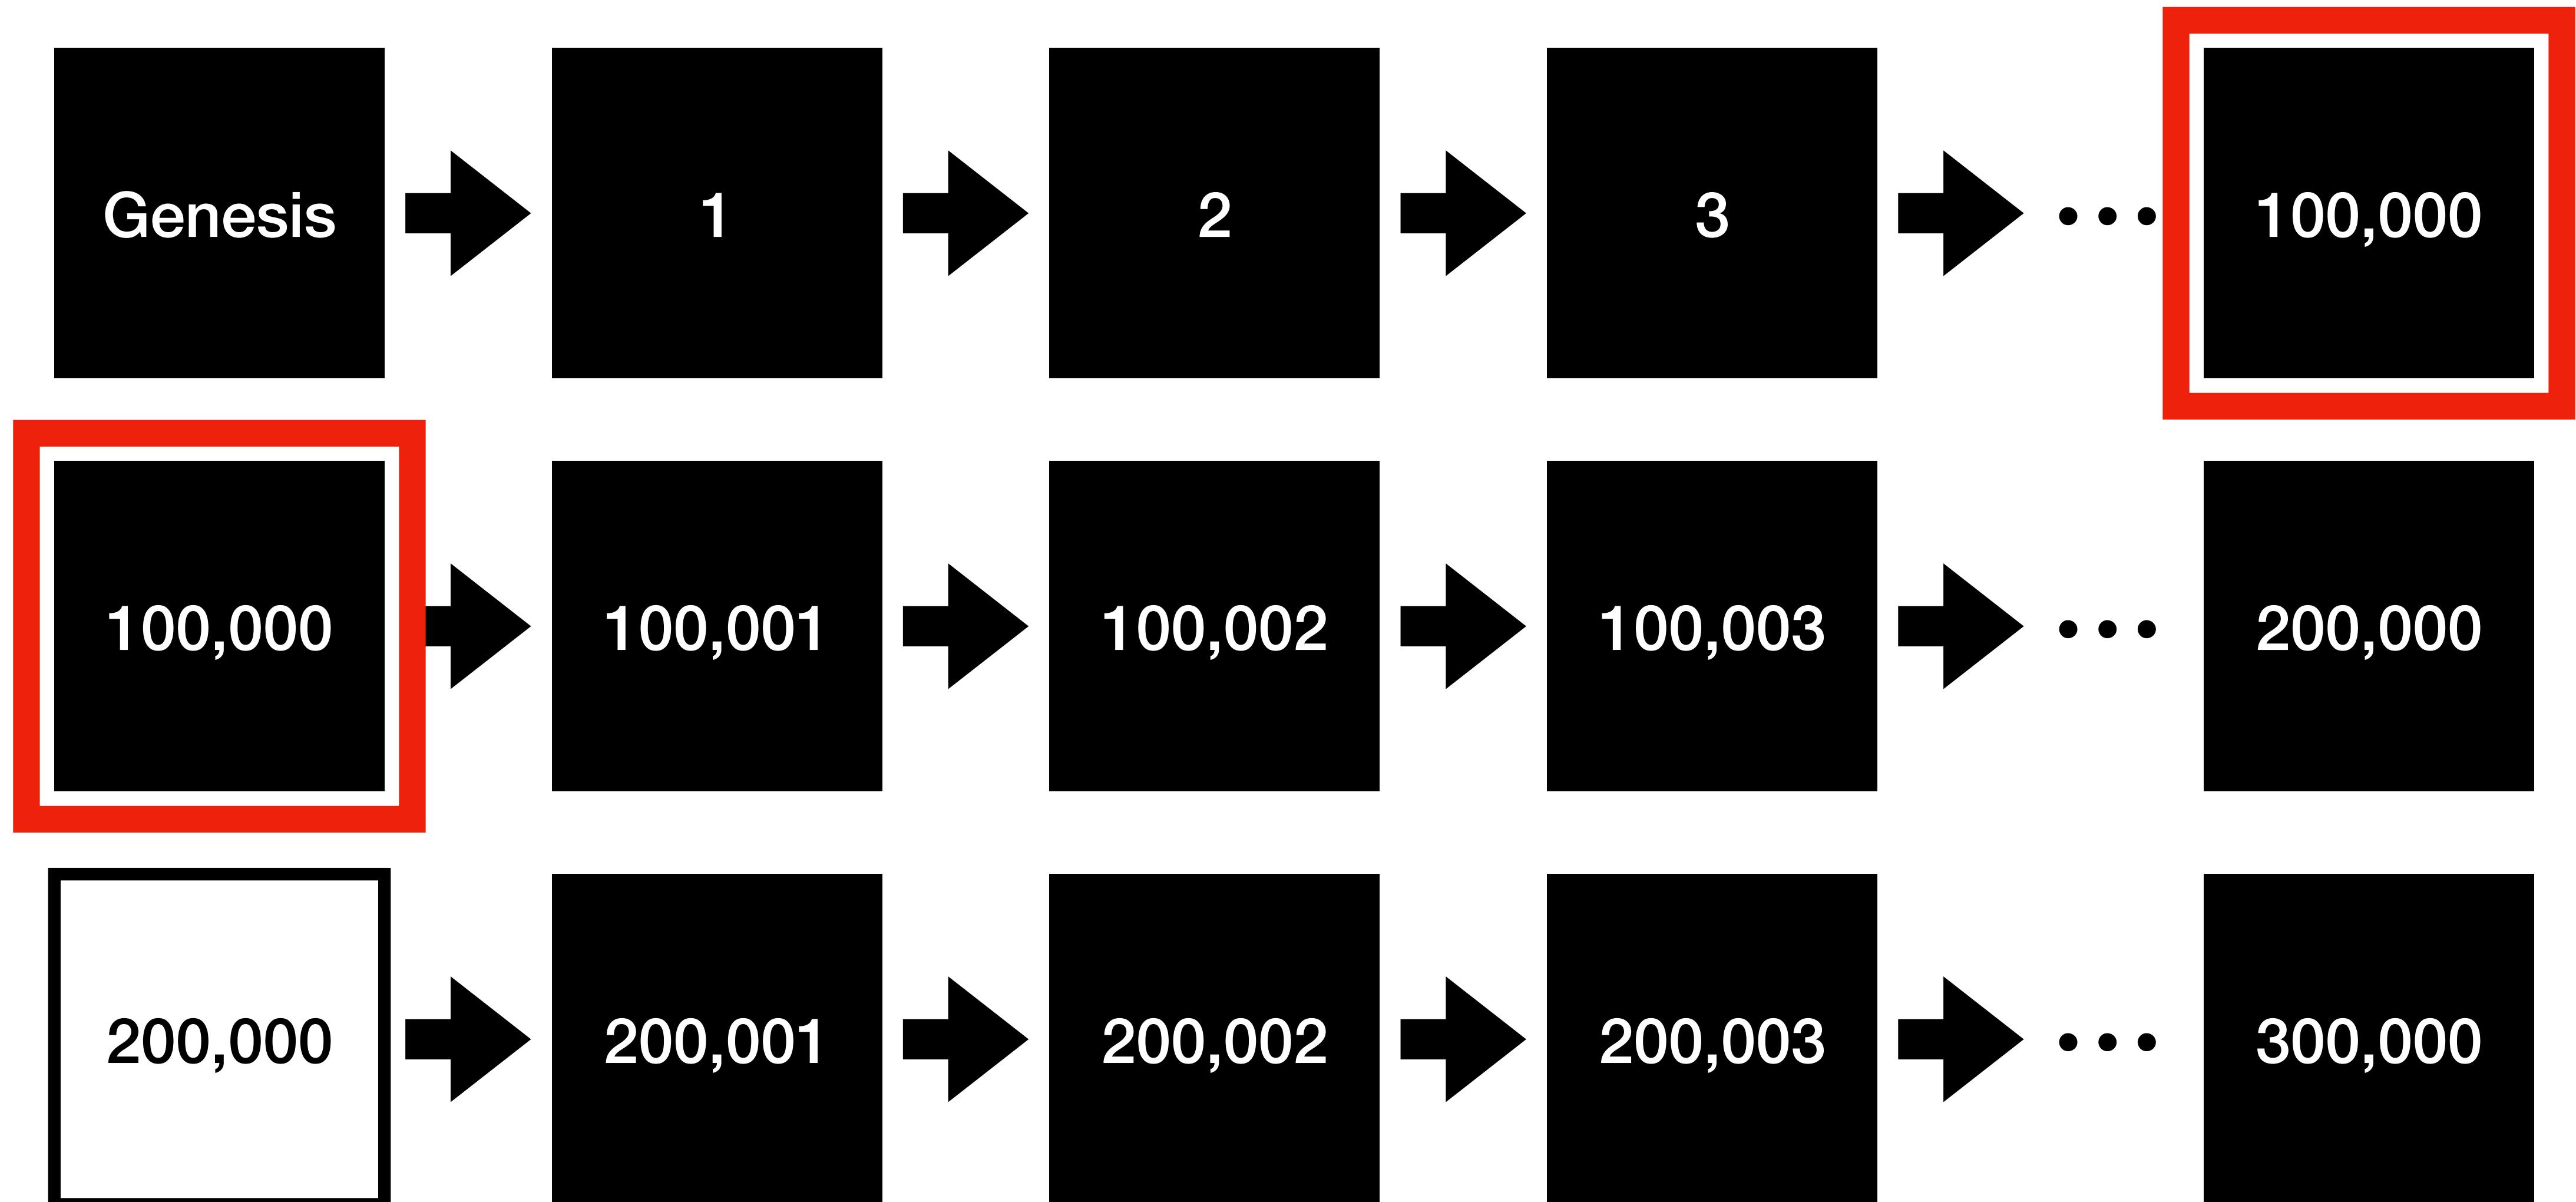
# With Utreexo

## Efficient parallel validation



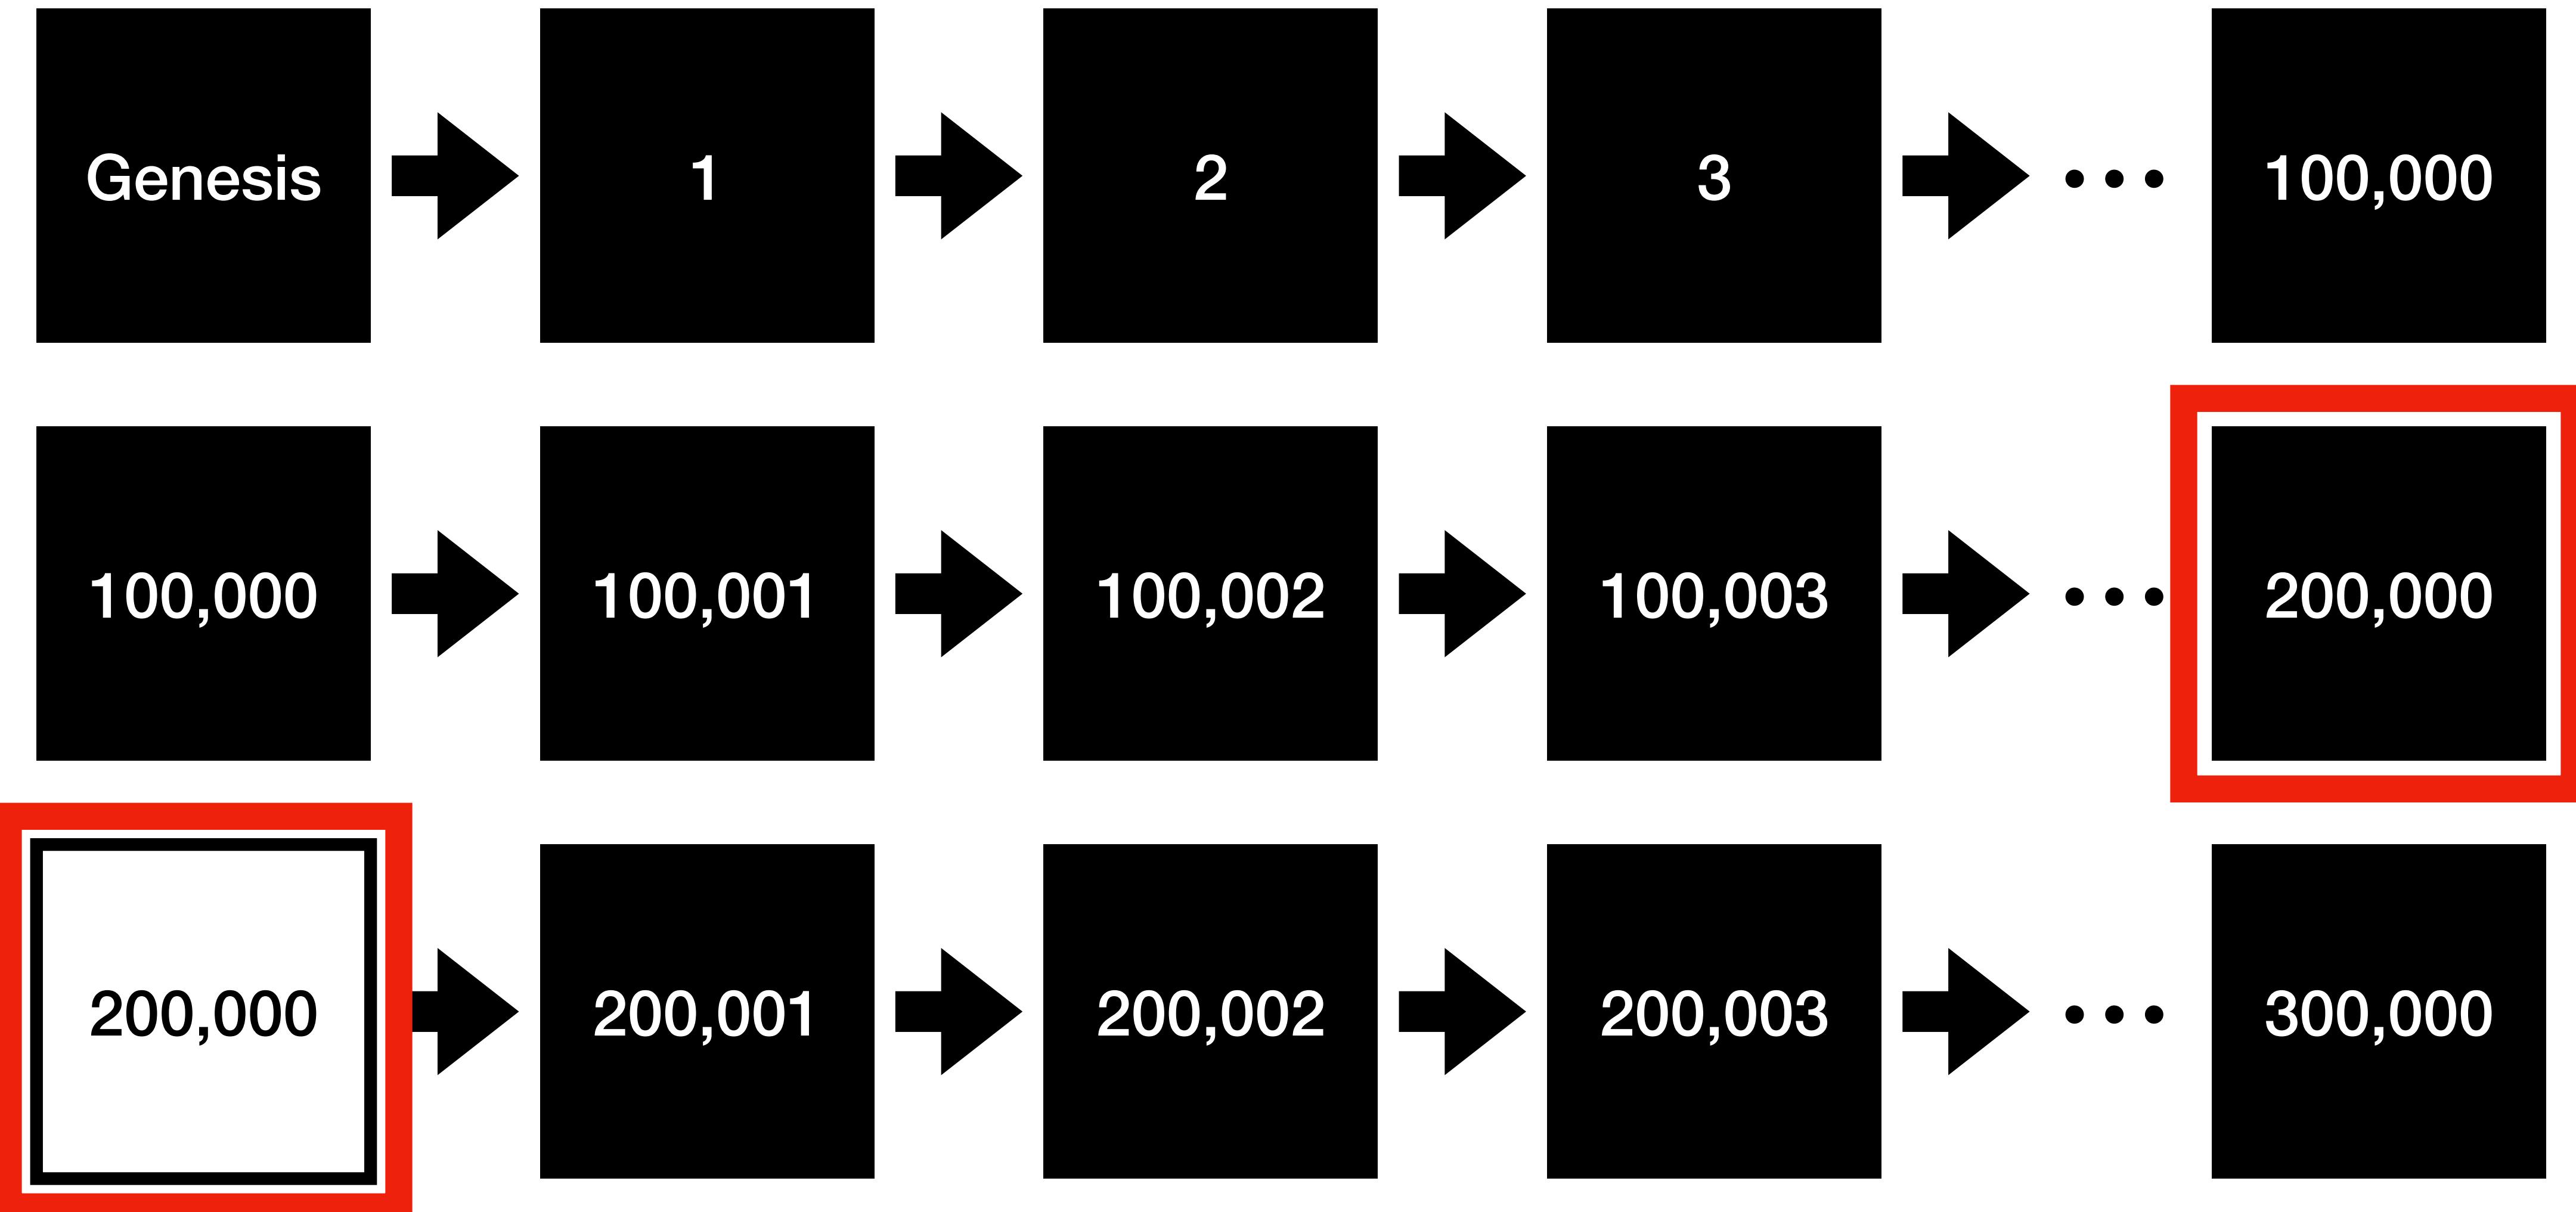
# With Utreexo

## Efficient parallel validation



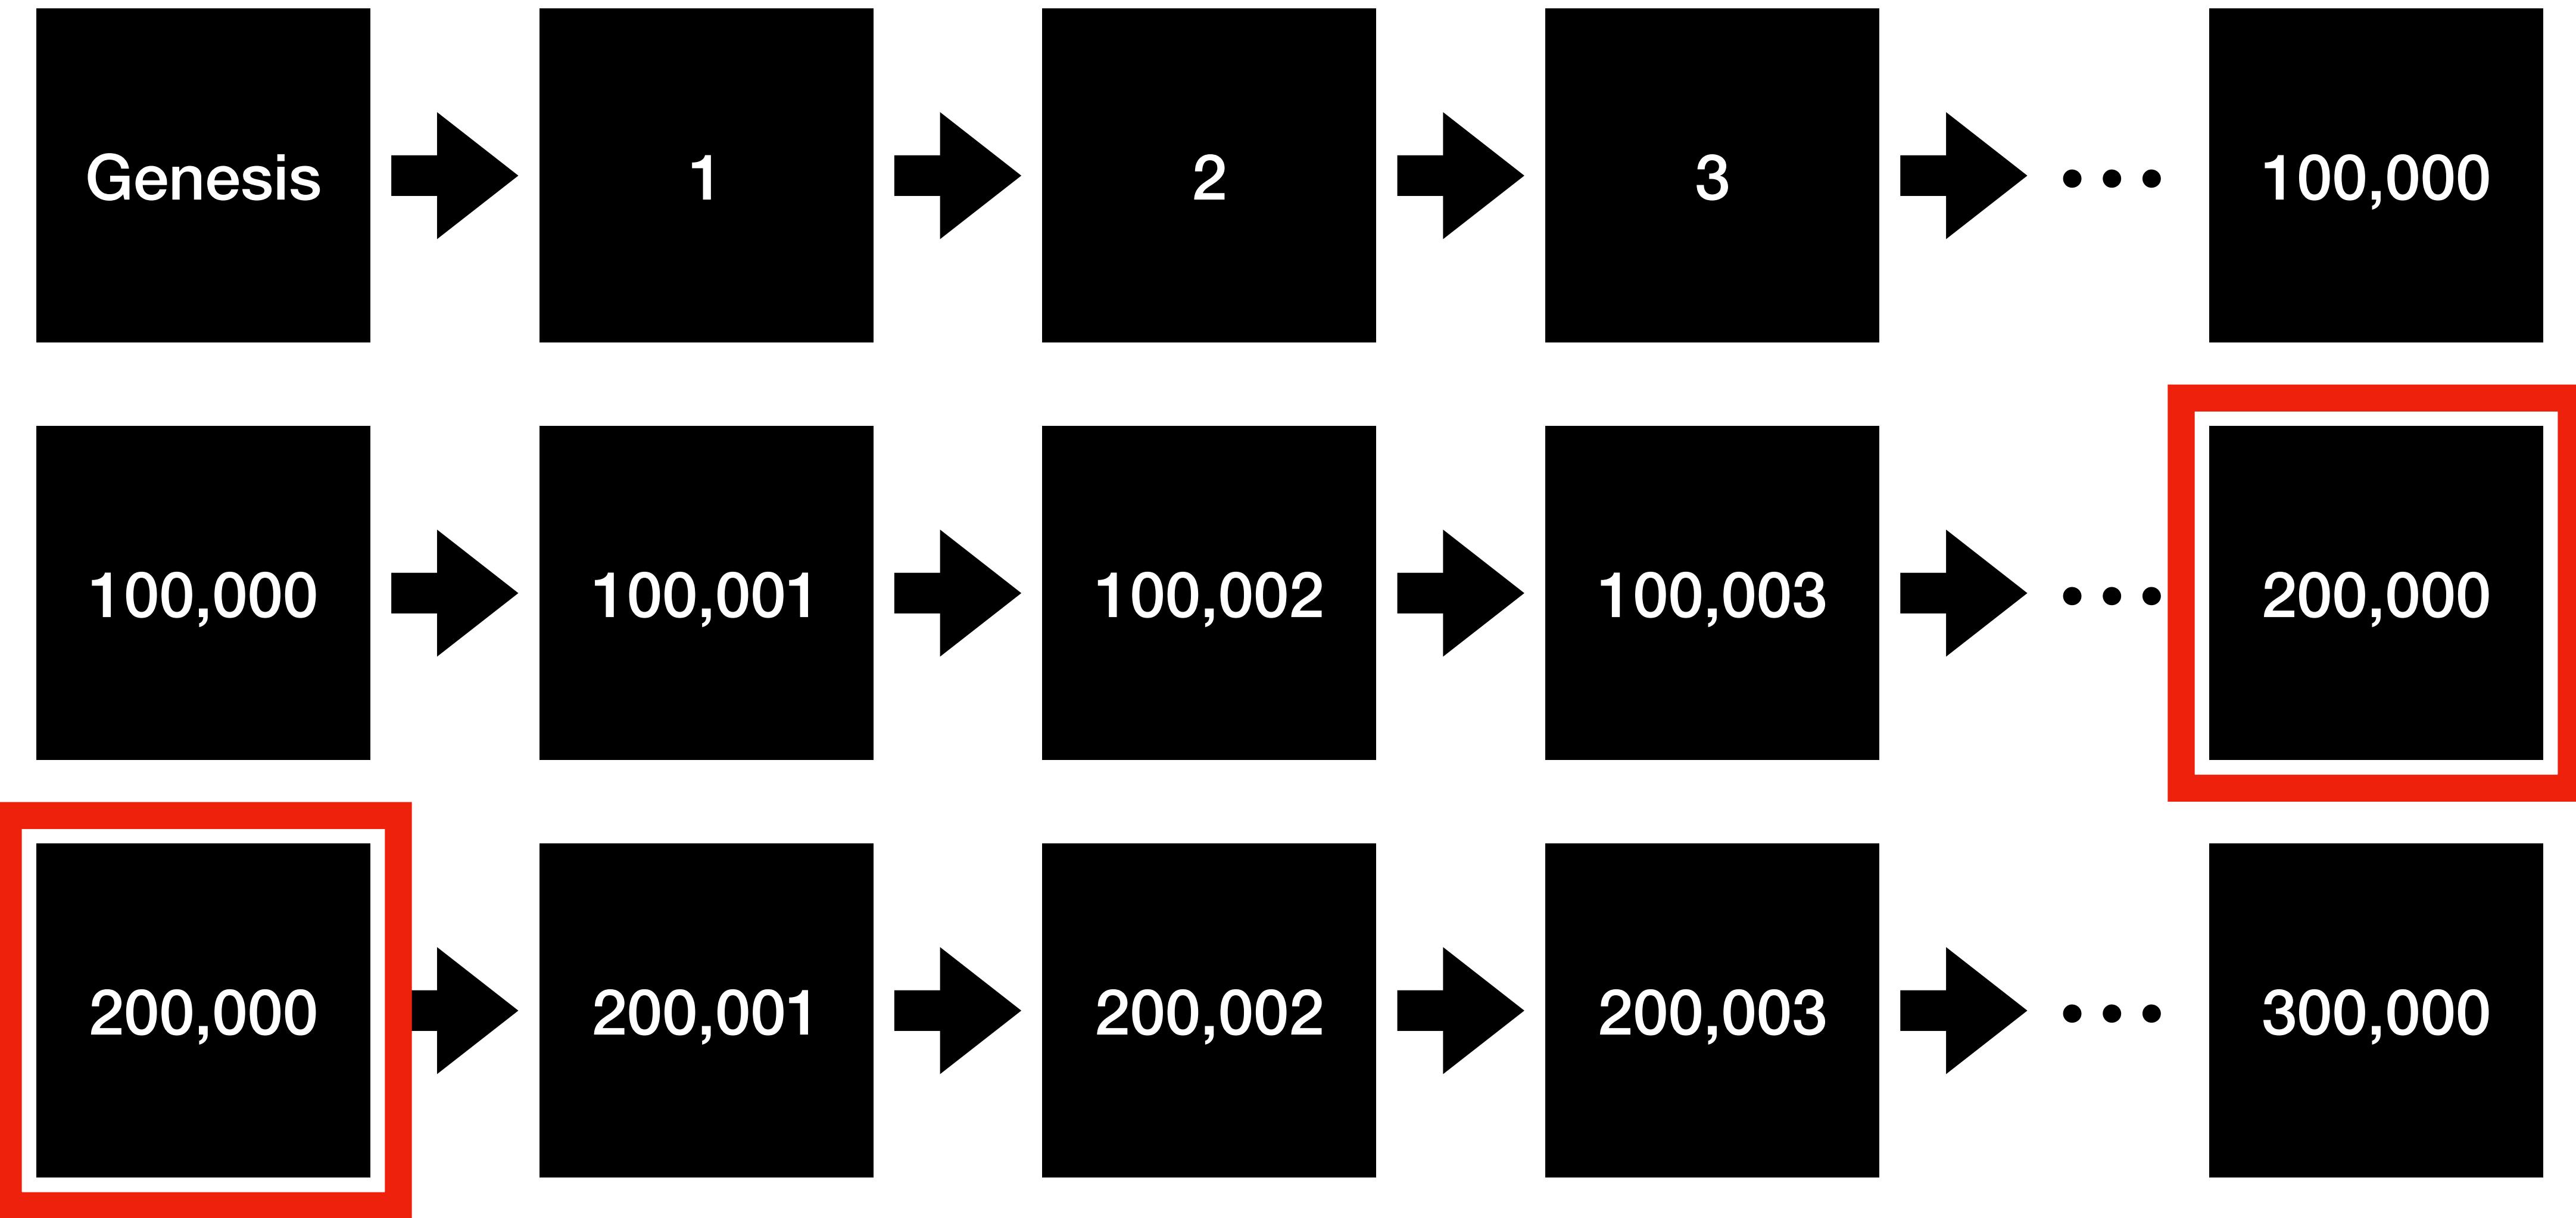
# With Utreexo

## Efficient parallel validation



# With Utreexo

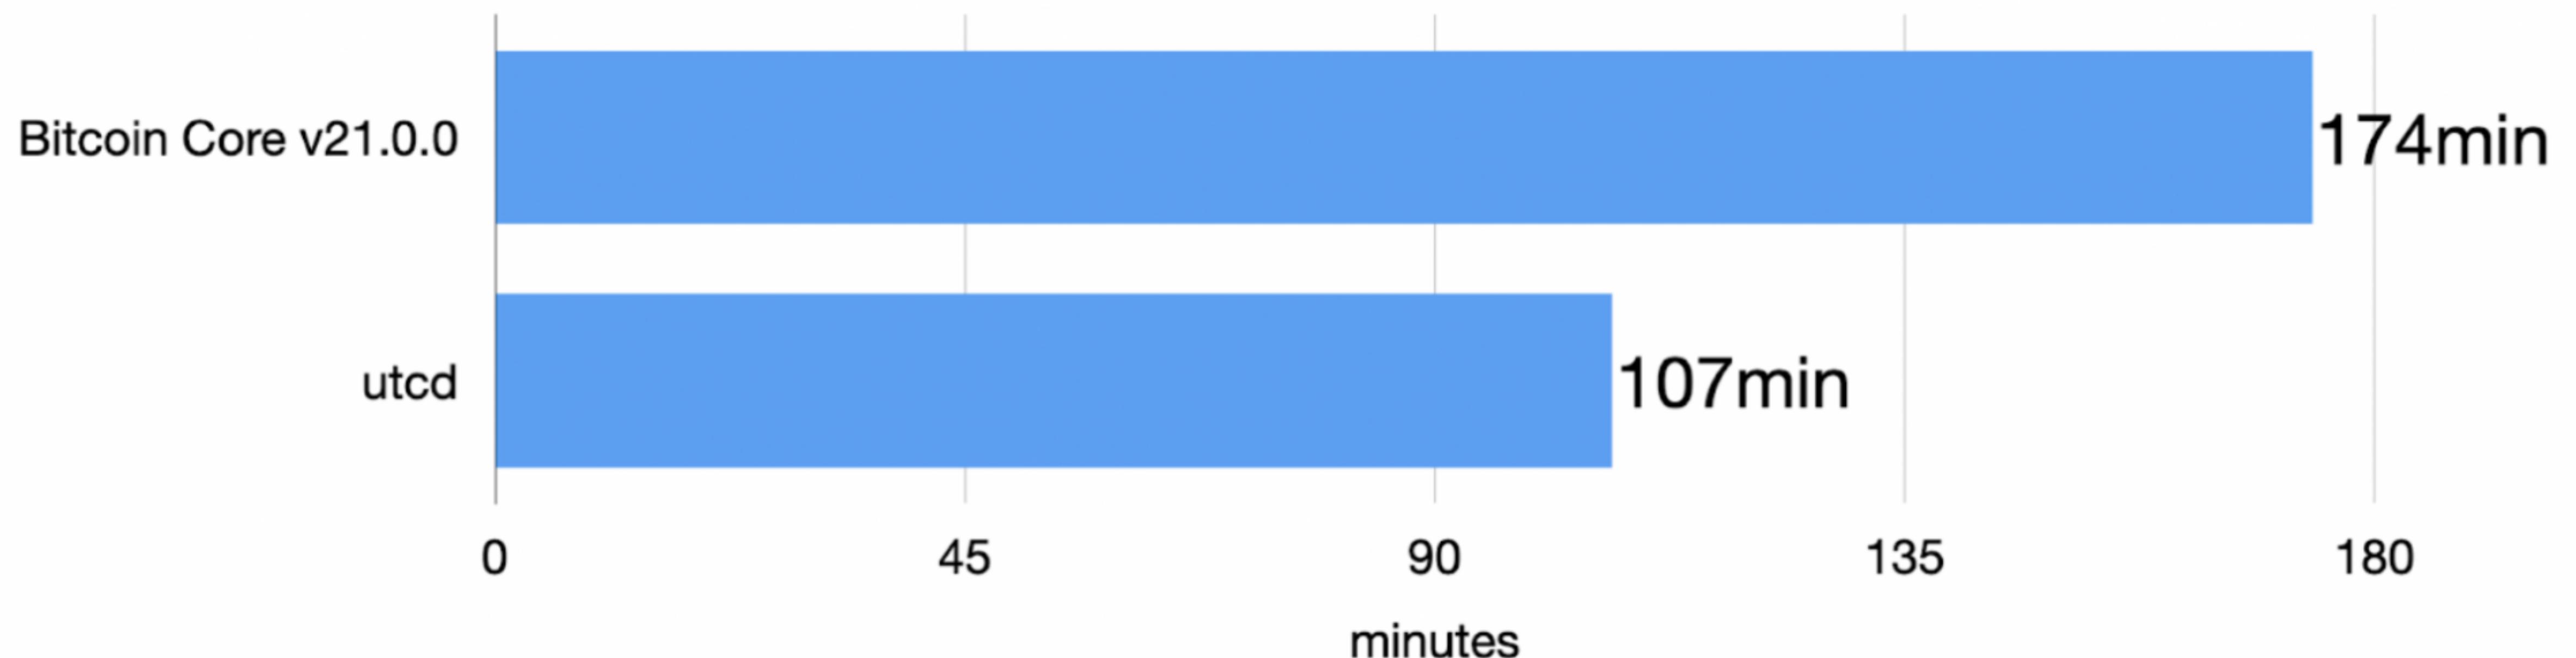
## Efficient parallel validation



# *Efficient* Parallel Validation

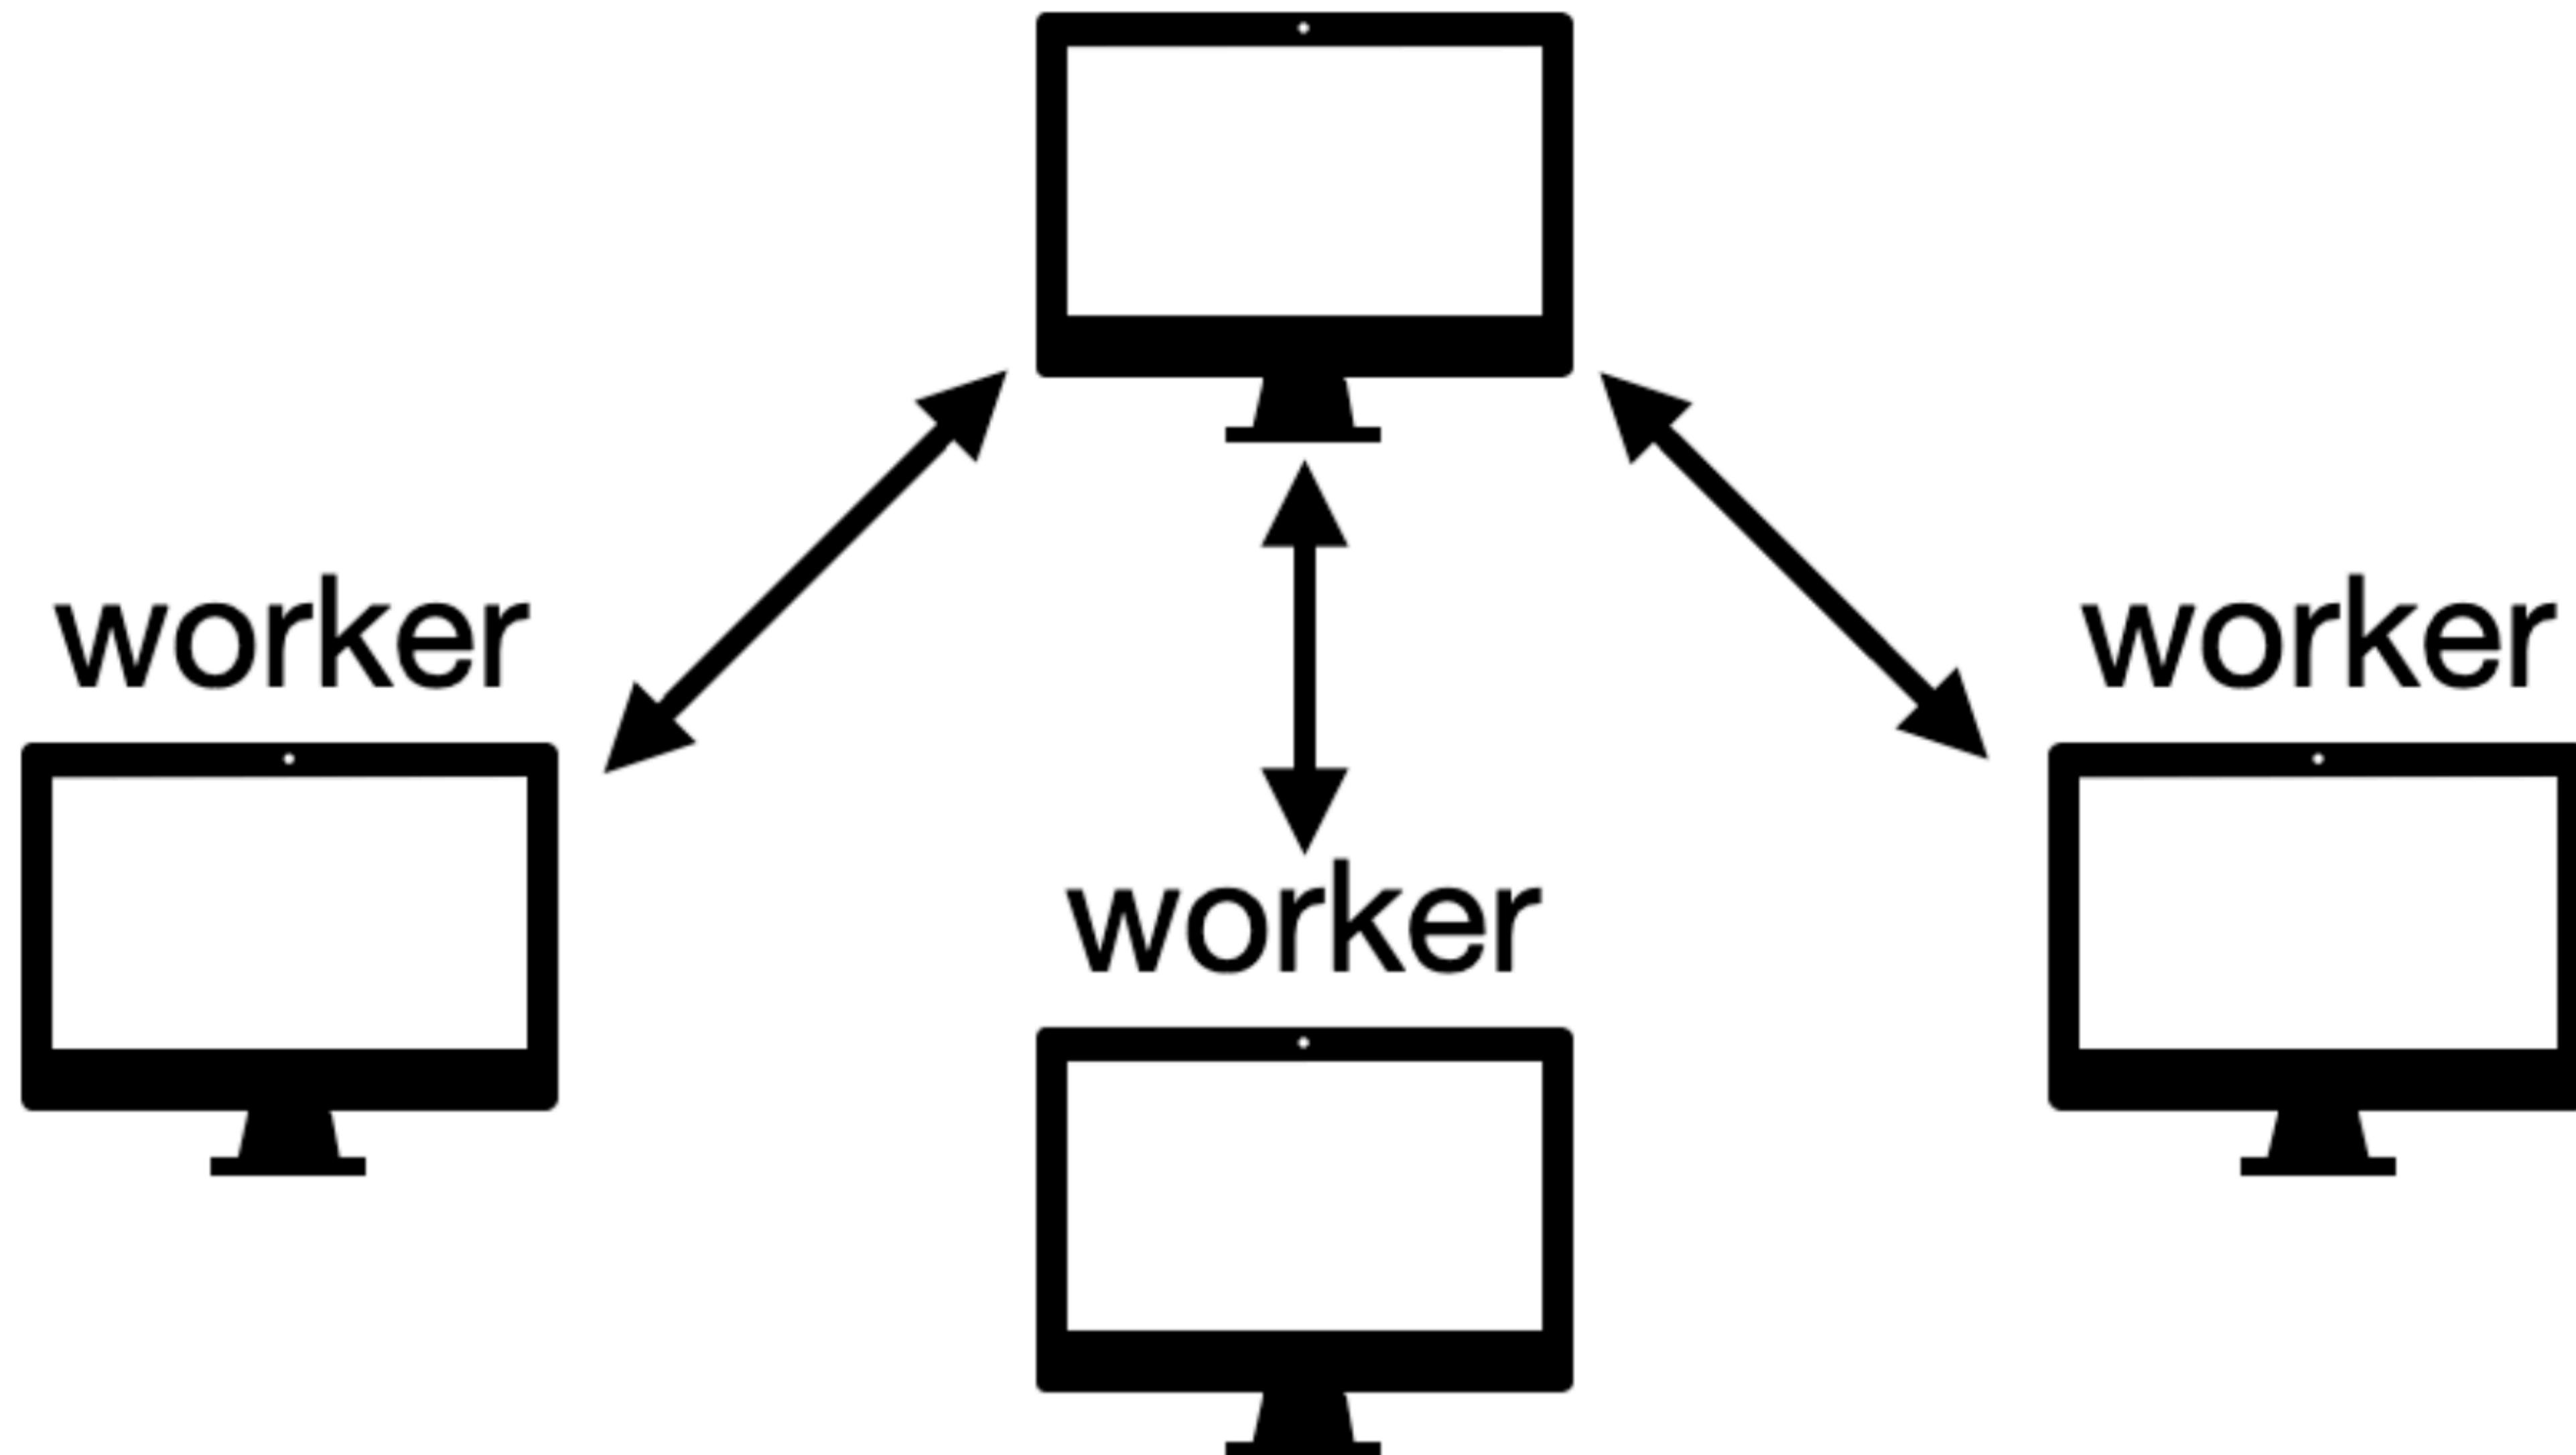
**SHA256**  **> Map access** 

## Initial Block Download Speeds - local nodes (default mode, no signature checks until block 654,683)

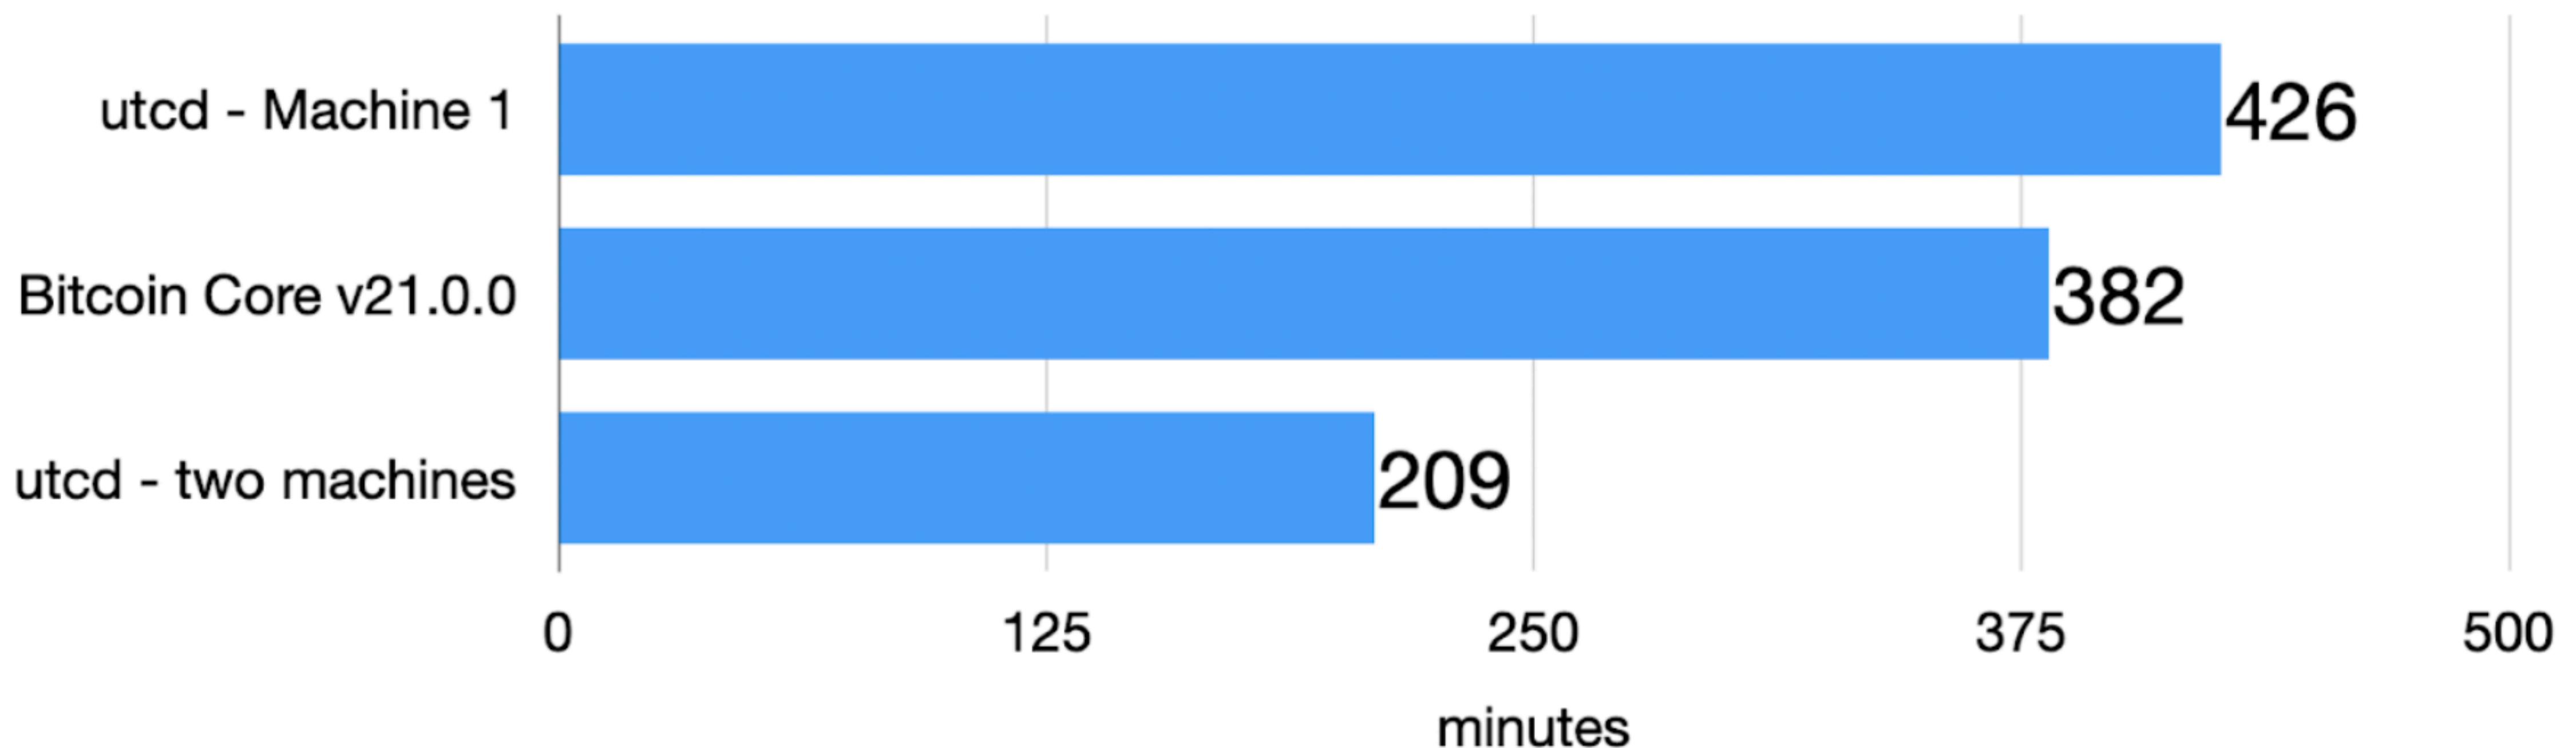


**~62% faster  
blockchain  
validation**

# Coordinator/worker



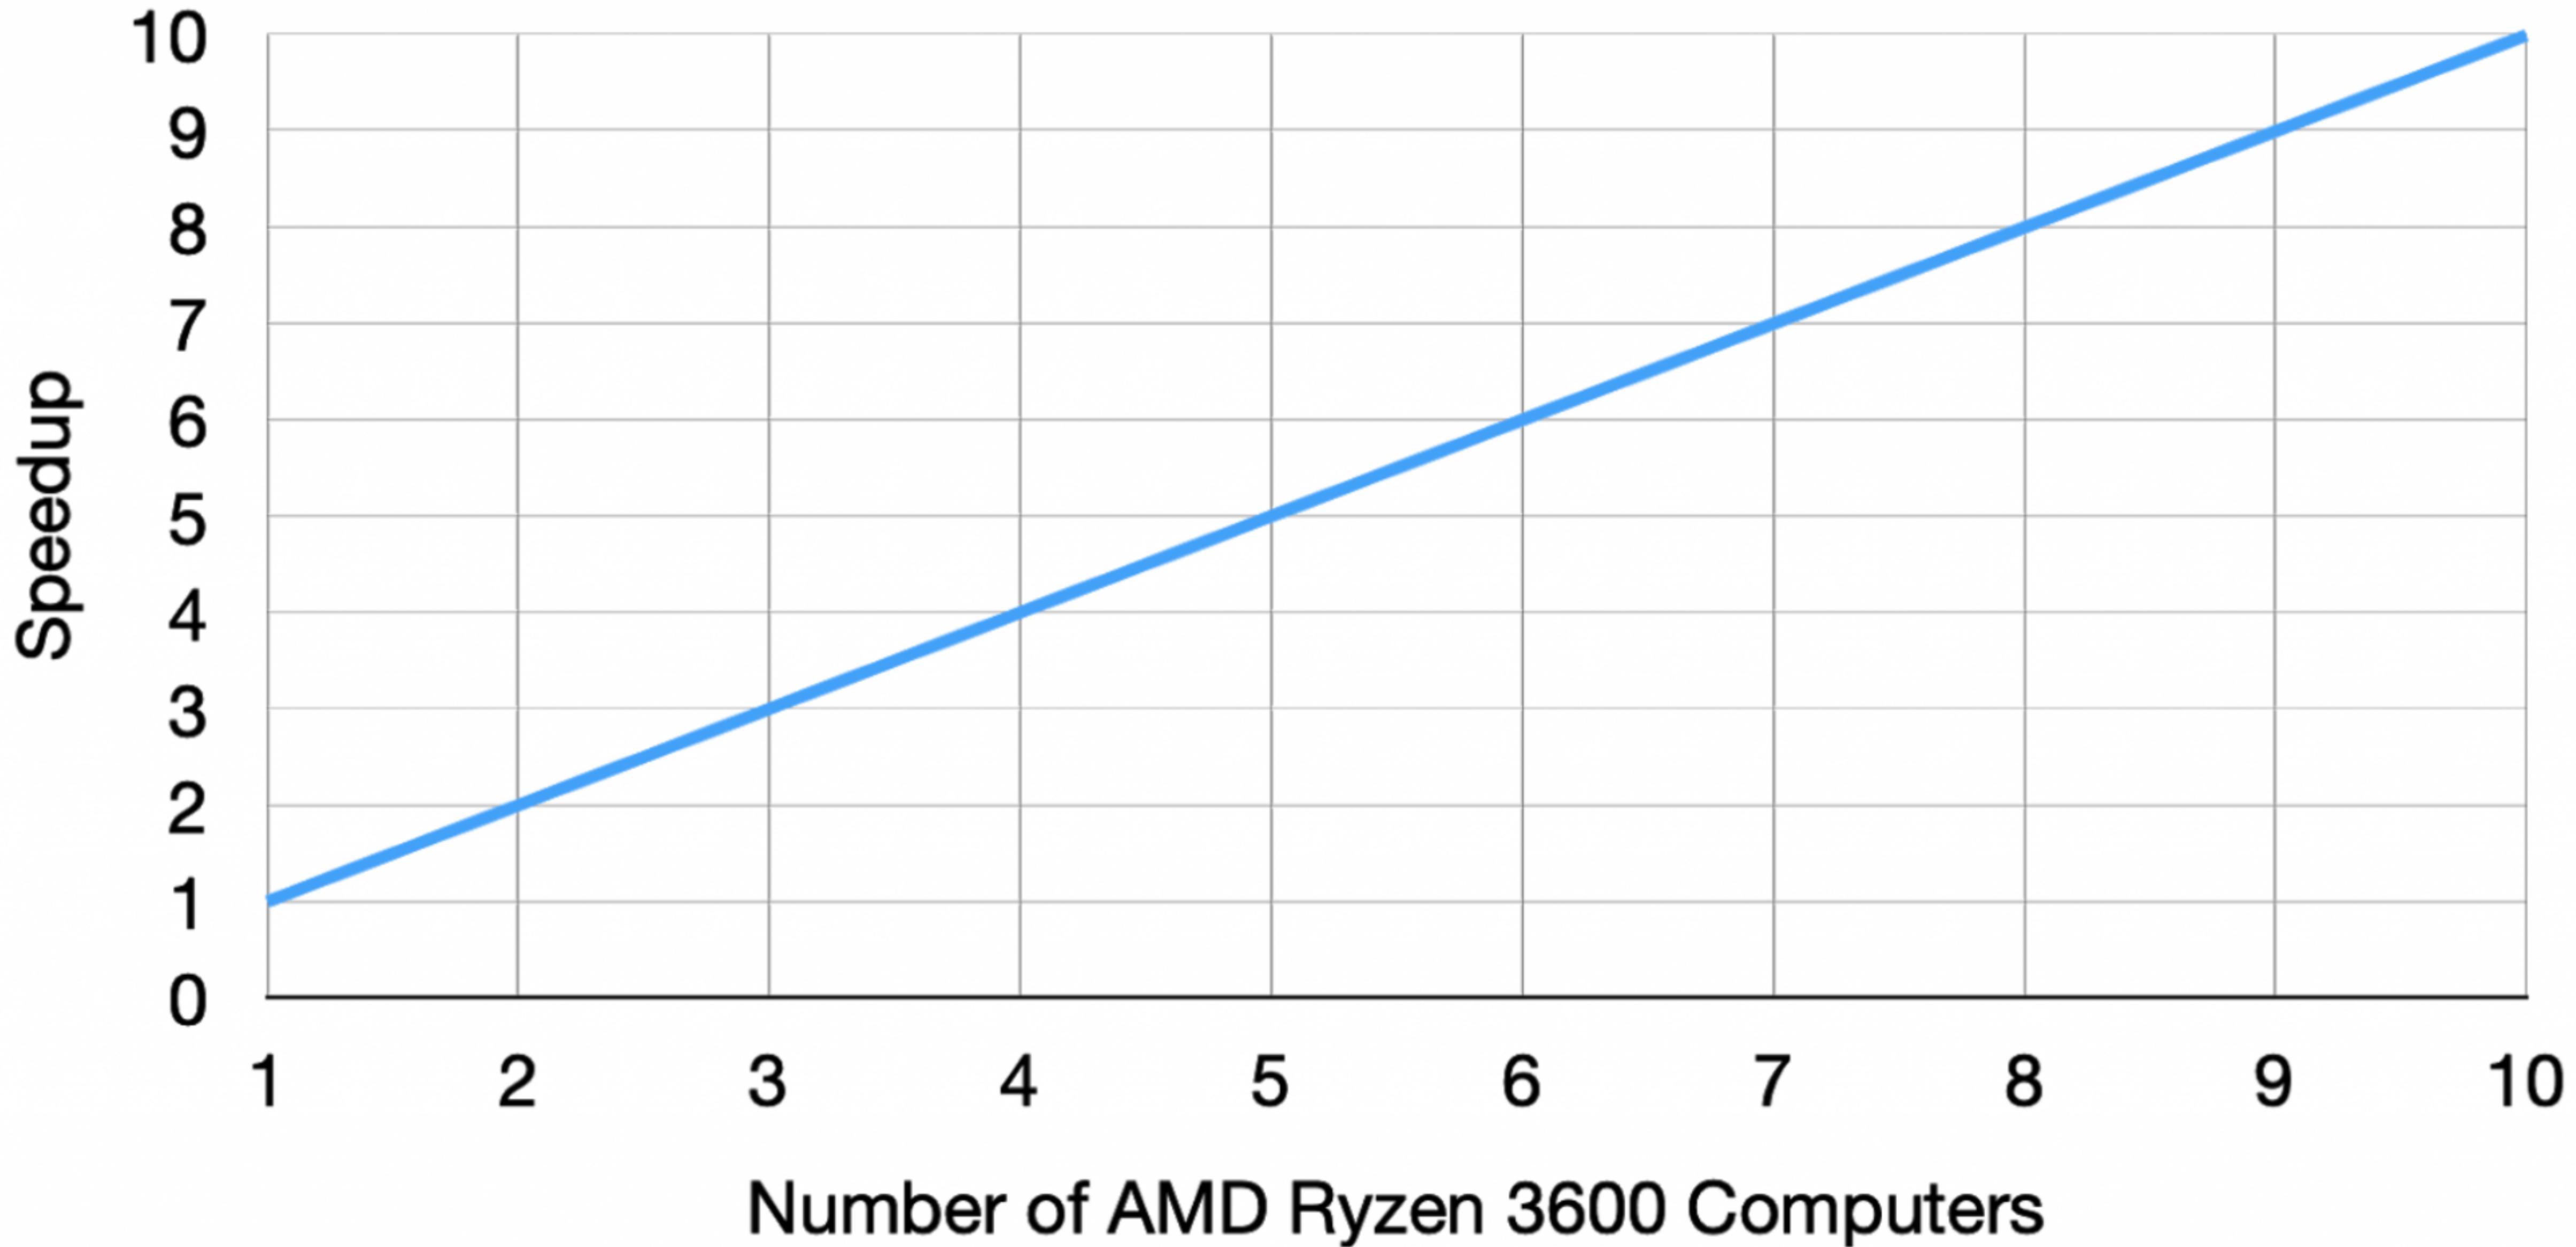
## Initial Block Download Speeds to block 671,000 (Full signature check)



# **2x faster blockchain validation**

**When using 2 computers**

# Utreexo Parallel Node Speedups (based on Amdahl's Law)



# **Summary**

## **What I want you to take away**

- We want to preserve trustlessness while having the most amount of Bitcoin users

# **Summary**

## **What I want you to take away**

- We want to preserve trustlessness while having the most amount of Bitcoin users
- Current bottleneck is with block validation

# Summary

## What I want you to take away

- We want to preserve trustlessness while having the most amount of Bitcoin users
- Current bottleneck is with block validation
- Utreexo helps eliminate this bottleneck

# Questions?